

Processing
“*BIG-DATA*”
In Real Time

Yanai Franchi , Tikal



Vacation to Barcelona



NON-END-RER-MU ONLY/DTE CHG M YR150/RFD MYR300				DEPARTURE DATE AND INVOICE DATE 08SEP07	ARRIVAL DATE 18/LXGUZO EXCHANGED FOR	MSIAN HARMONY TOUR N TRAVEL - 1 KUALA LUMPUR /MY 20300626 /LL/U75B/1			
				IATA BSP					
1	DEPARTURE AIRPORT KUALA LUMPUR	HQ	5024N 20DEC 1645	OKYEE45	FARE BASIS	03FEB20K			
2	SHANGHAI PUDONG	MU	5023N 06JAN 1000	OKYEE45		03FEB20K			
3	KUALA LUMPUR	VOID	VOID						
4	--VOID--	VOID	VOID						
5	--VOID--	VOID	VOID						
MVR 5774KUE MD SHA843. 95MU KUL843. 95NUC1587. 70END ROE3. 4224									
DEPARTURE AIRPORT KUALA LUMPUR									
FAIRFAX CHARGE 51MY									
FAIRFAX CHARGE 42CN									
FAIRFAX CHARGE 314YDINVABY									
MVR 6183									
REF ID: 77867									
CARRIER									
APP CODE: KUL07010									
ORIGINAL ISSUE									
1 781									

After a Long Travel Day



Going to a Salsa Club



Best Salsa Club

NOW



- Good Music
- Crowded –
Now!

Same Problem in “gogobot”

gogobot Destinations Tribes Hotels More Find Places or People ... Sign In Sign Up

Things to do in Barcelona Back to list view

OVERVIEW HOTELS VACATION RENTALS THINGS TO DO RESTAURANTS

696 Results

1 La Sagrada Família Popular with Design Local History

2 Park Güell Popular with Design Local History

3 Las Ramblas Popular with Local Budget Backpackers

4 La Boqueria Popular with Foodies Local Backpackers

5 Camp Nou Popular with History Families Local

6 Casa Milà

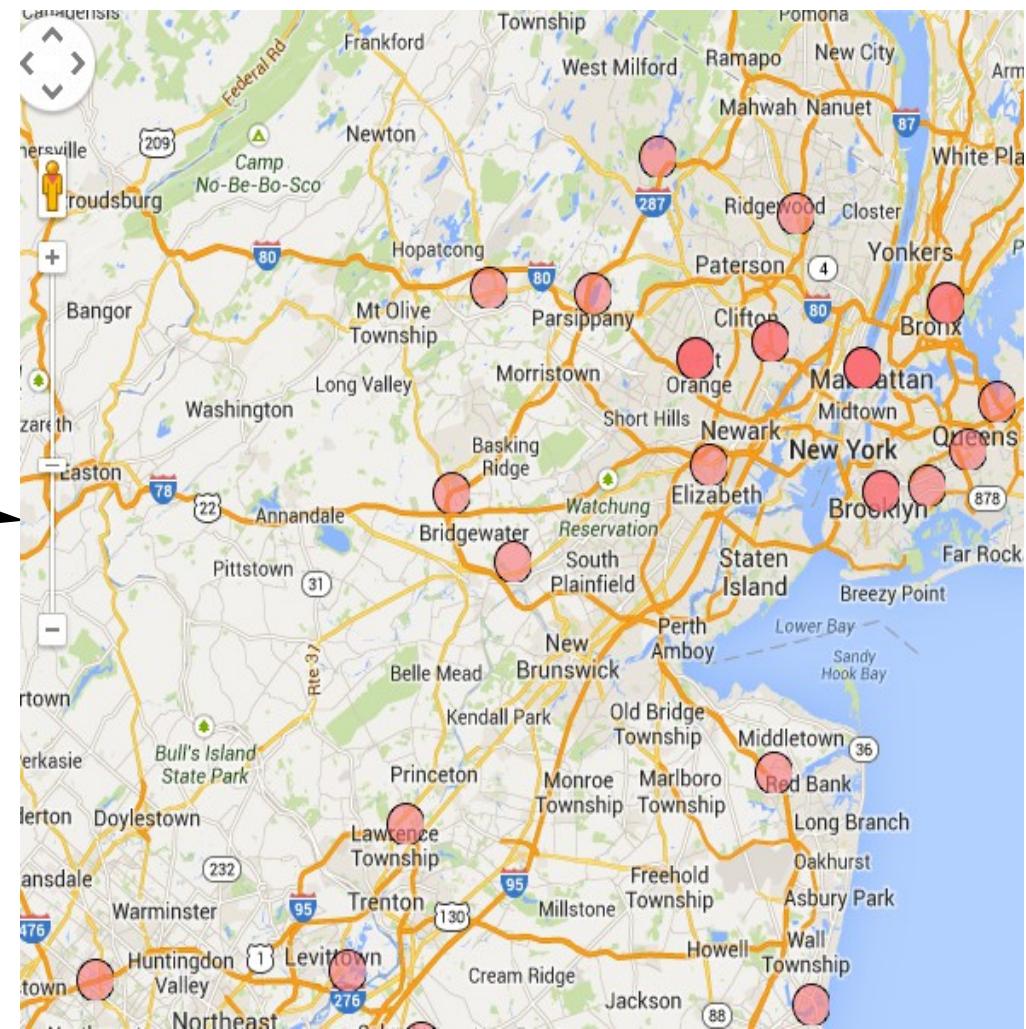
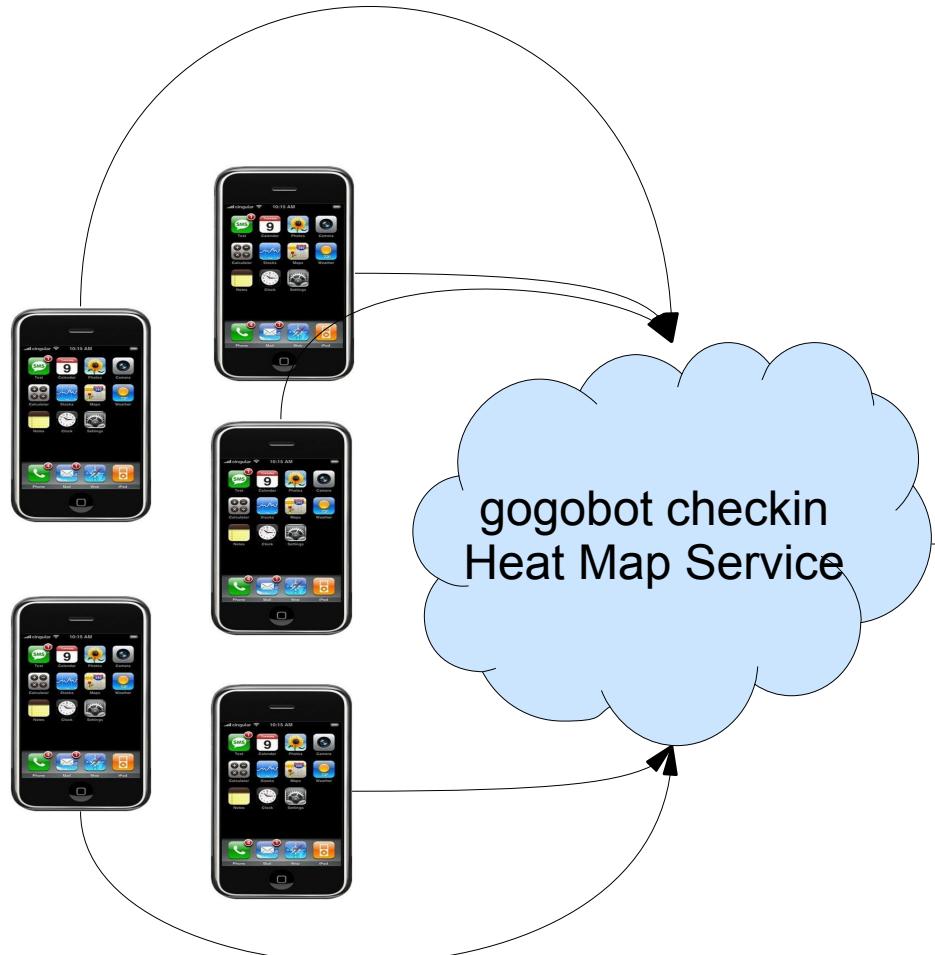
A map of Barcelona, Spain, highlighting several key locations. The map includes labels for 'Torre De Collserola', 'CosmoCaixa', 'B-20', 'LA SALUT', 'GRÀCIA', 'VILA DE GRÀCIA', 'SARRIÀ-SANT GERVASI', 'ST. GERVASI - LA BONANOVA', 'SANT GERVASI', 'ASTURIÉS', 'LES CORTS', 'VILLARROEL', 'L'ANTIGA ESCUERRA DE L'EIXAMPLE', 'LES TRES TORRES', 'VIA AUGUSTA', 'Passeig de la Bonanova', 'Fundació Eina', 'IESE Business School', 'Monastery of Pedralbes', 'Escola Universitària Salesians de Sarrià - UAB', 'Hospital de Barcelona', 'Finca Güell', 'Campus Nord', 'PEDRALBES', and 'LA MATERNITAT'. Numbered pins (1 through 6) mark specific points of interest corresponding to the list on the left.

← 1 2 3 ... 50 →



Do it
Realtime!

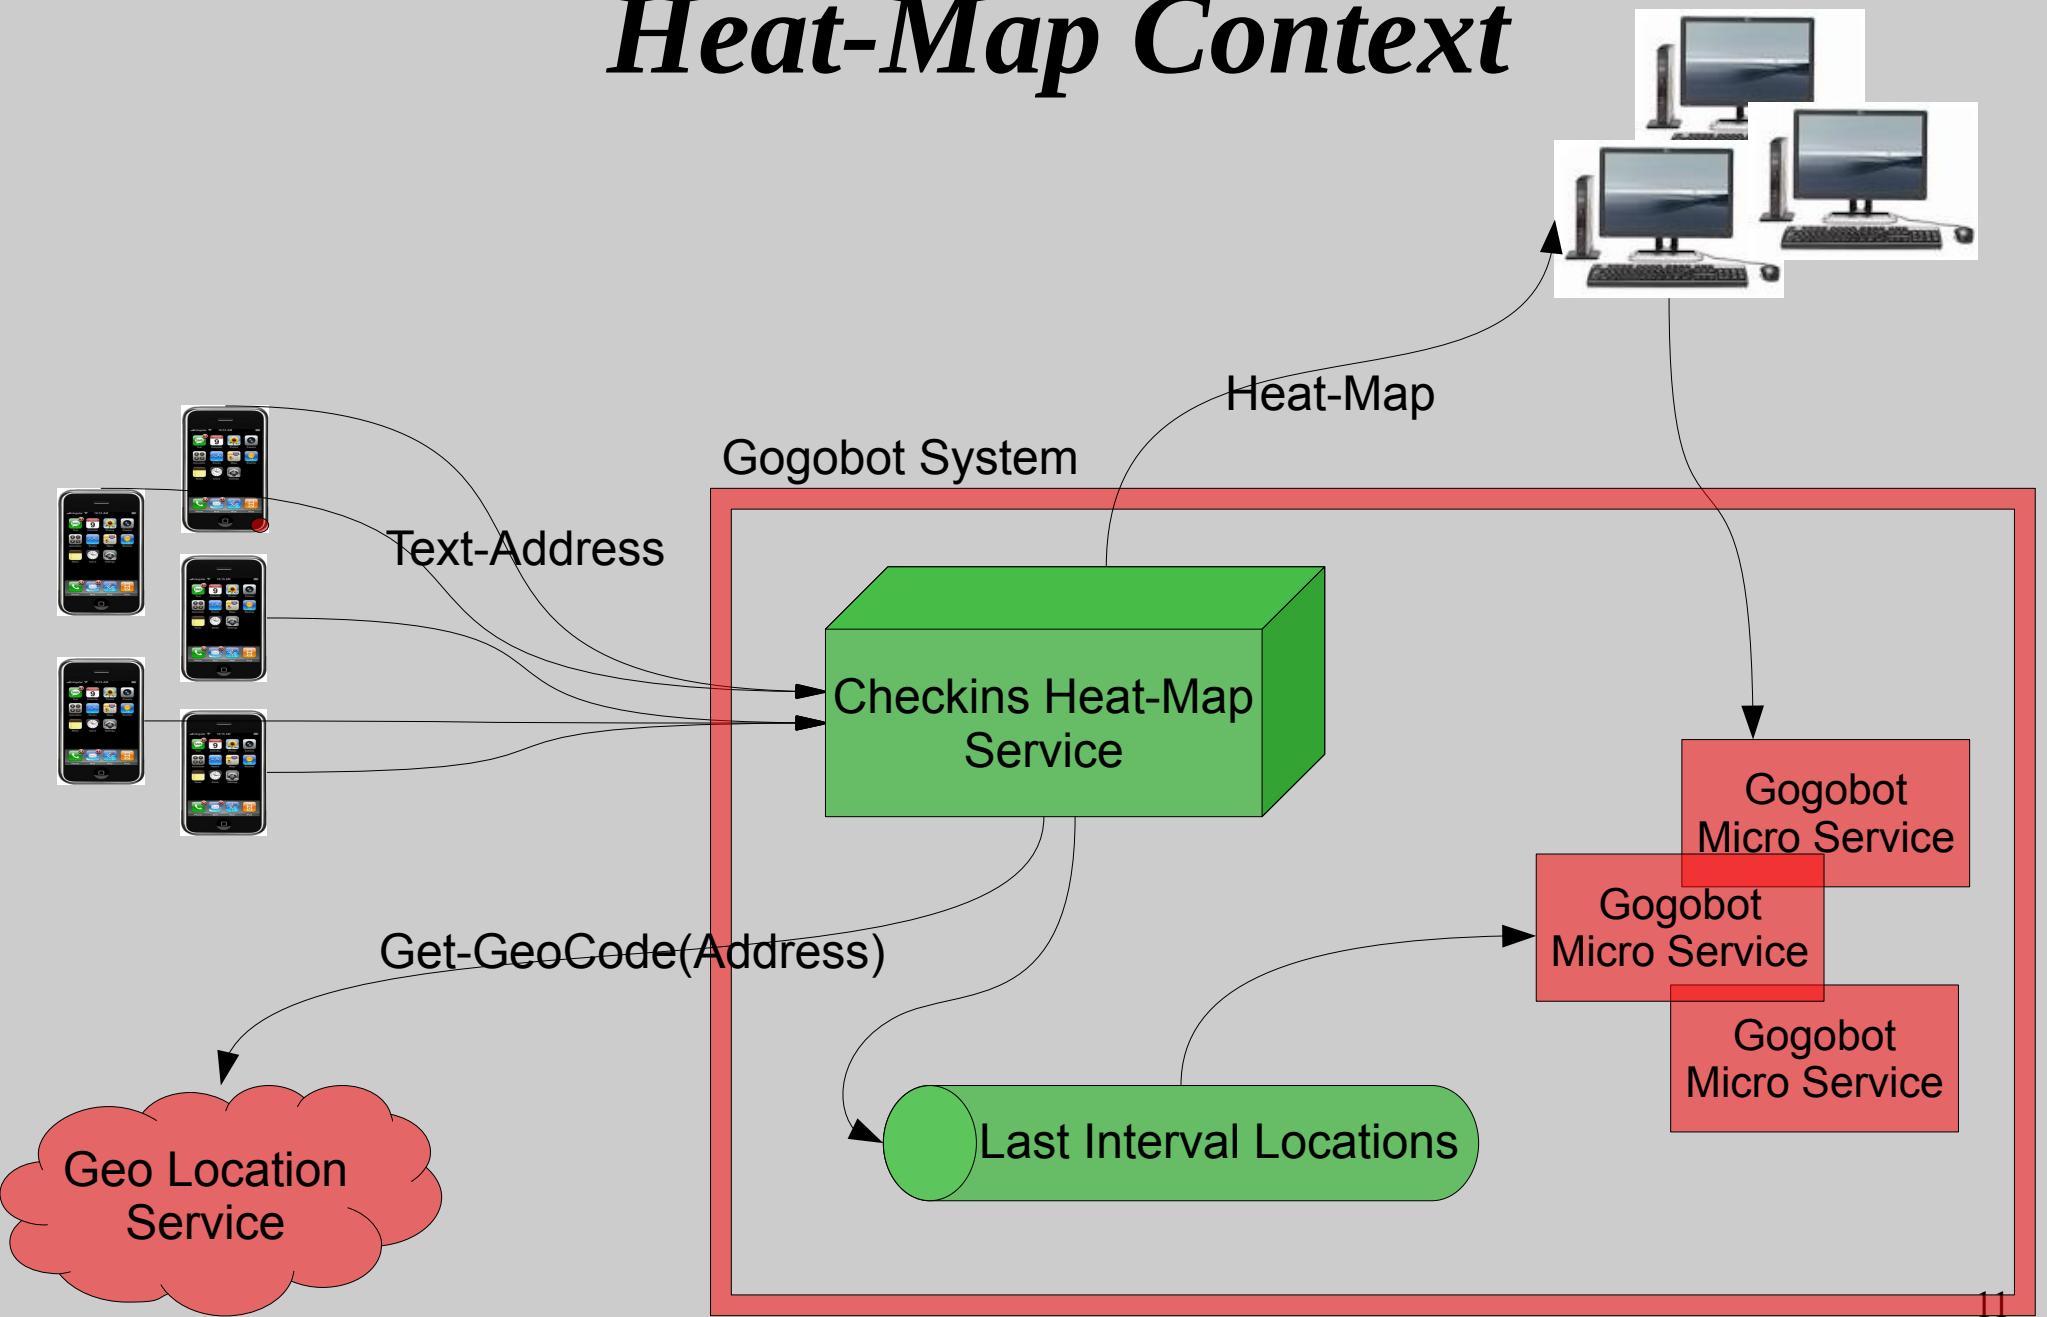
Lets' Develop “Gogobot Checkins Heat-Map”



Key Notes

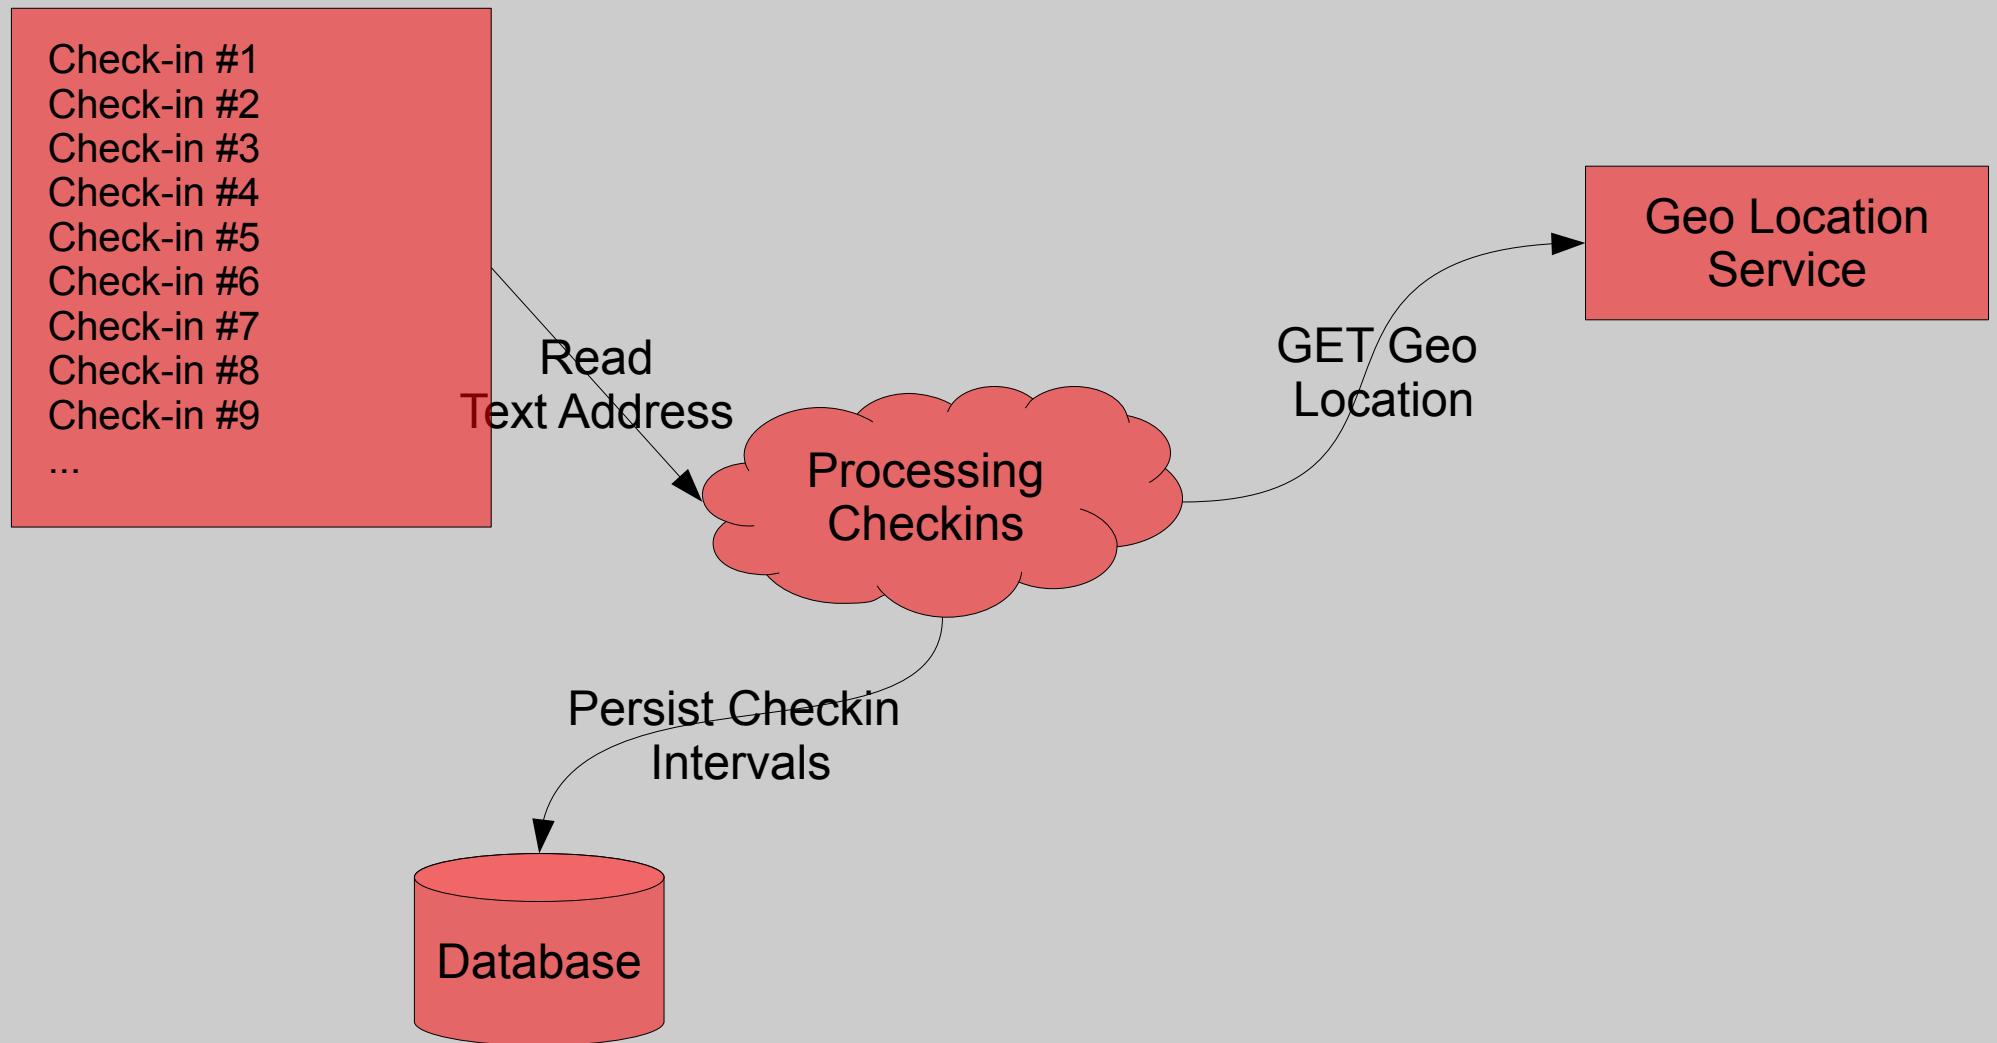
- Collector Service - Collects checkins as text addresses
 - We need to use GeoLocation Service
- Upon elapsed interval, the last locations list will be displayed as Heat-Map in GUI.
- Web Scale service – 10Ks checkins/seconds all over the world (imaginary, but lets do it for the exercise).
- Accuracy – Sample data, NOT critical data.
 - Proportionately representative
 - Data volume is large enough to compensate for data loss.

Heat-Map Context



Plan A

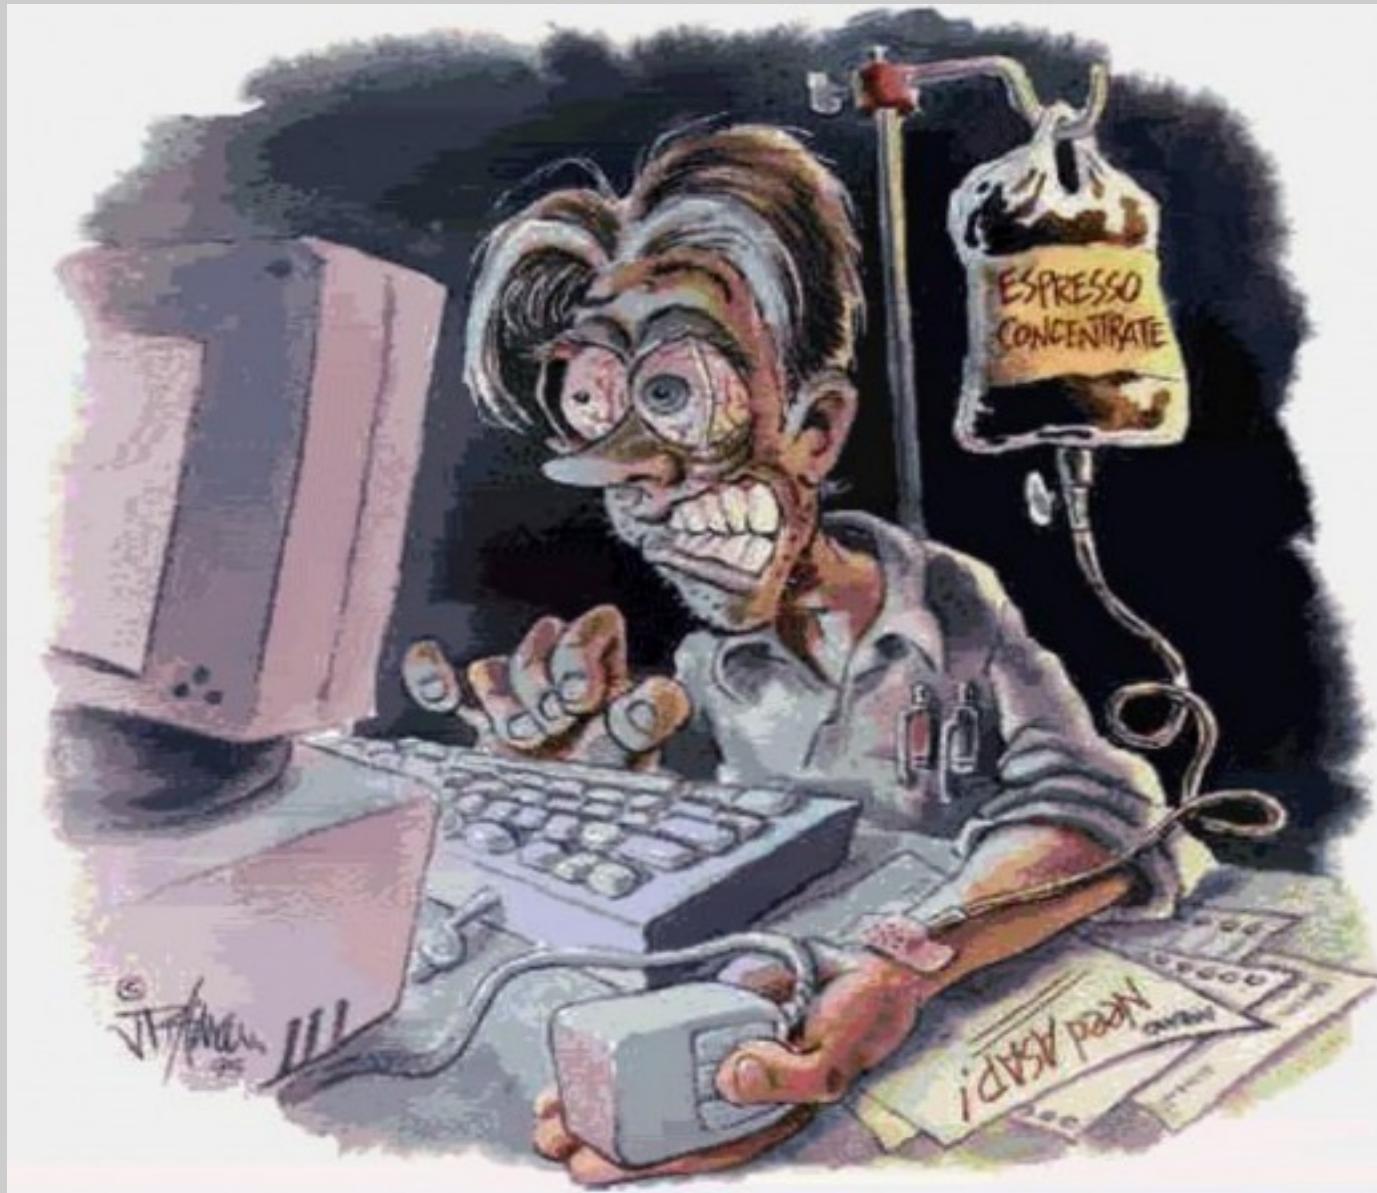
Simulate Checkins with a File



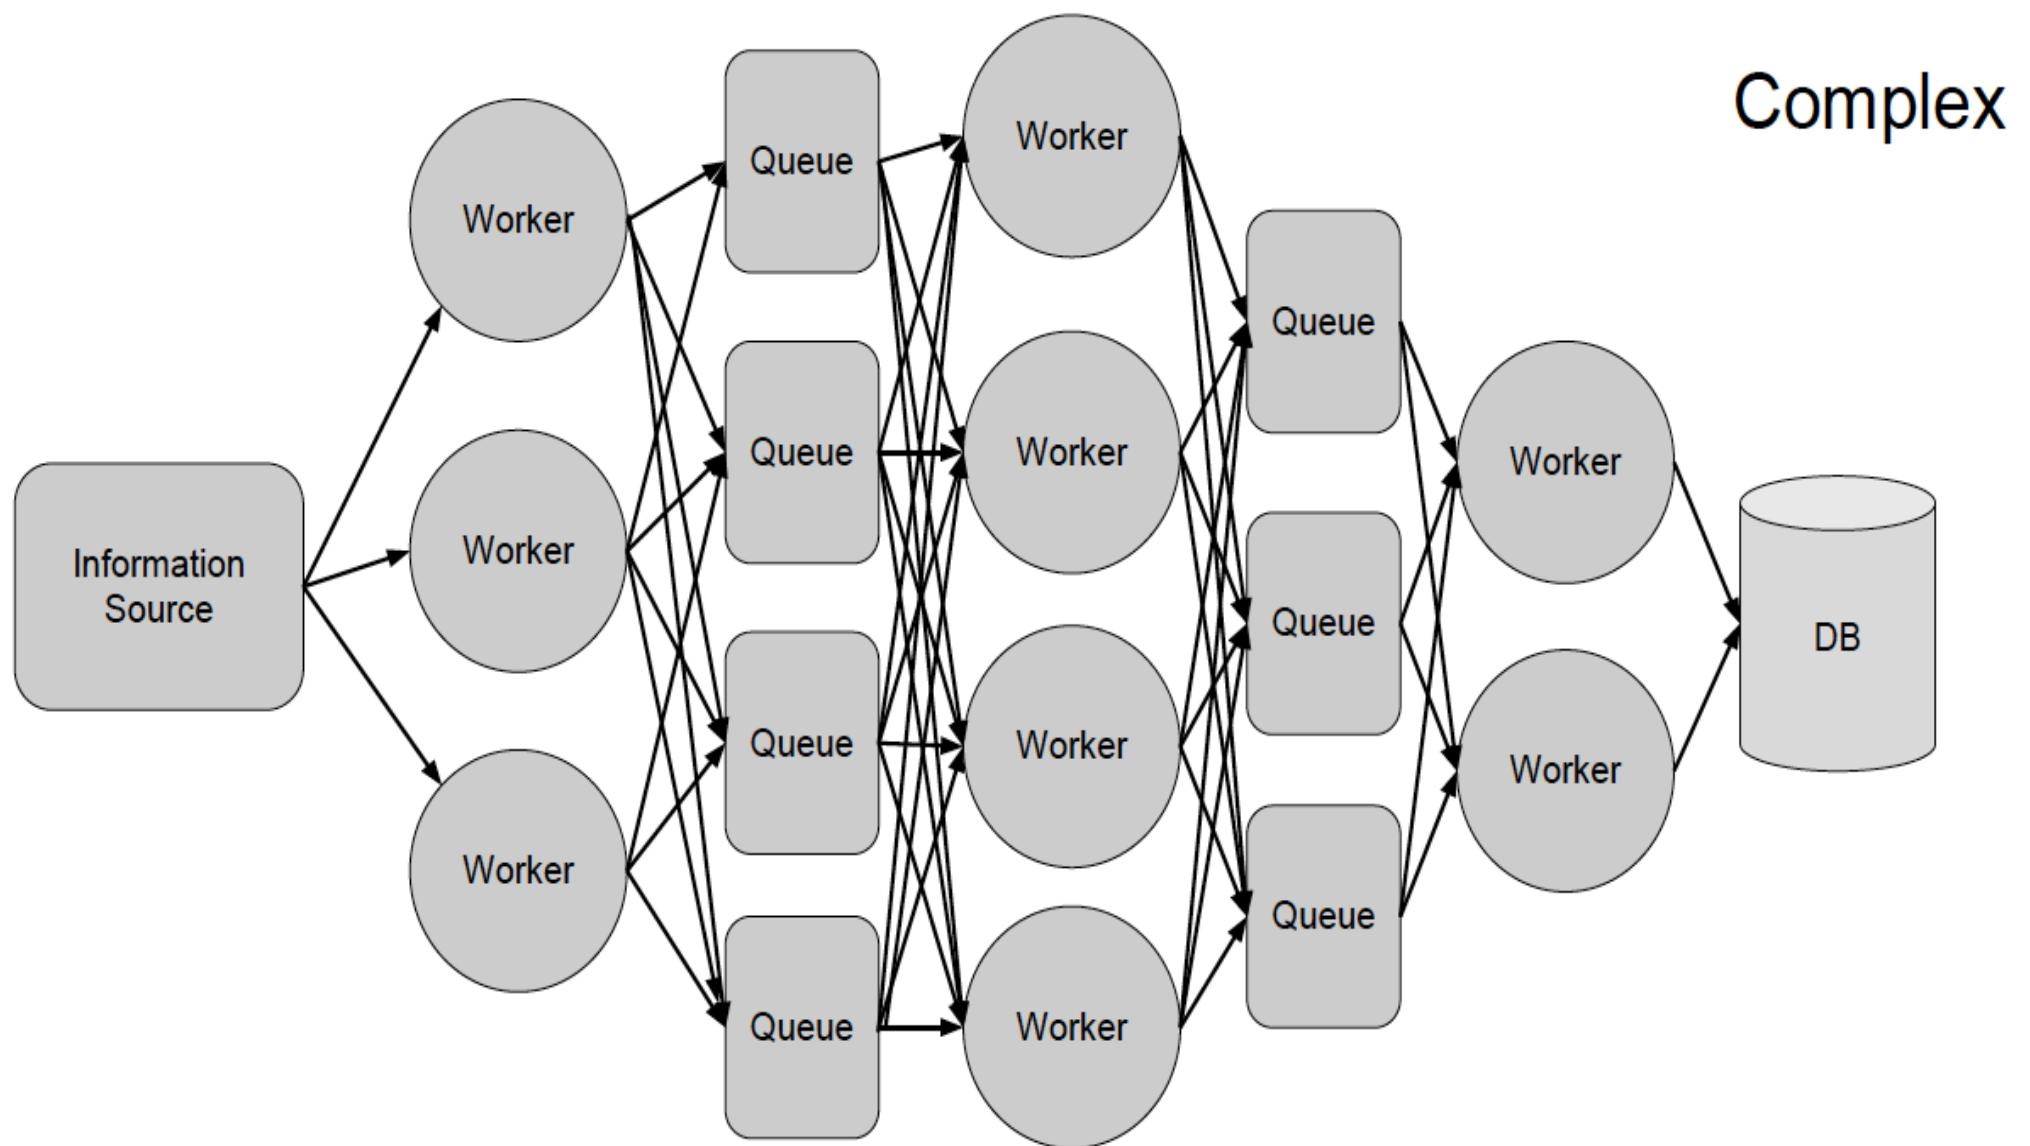
A wide-angle photograph of a massive waterfall, likely Niagara Falls, showing the Horseshoe Falls section. The water flows powerfully over dark, rocky ledges, creating white foam and spray. In the background, a dense forest lines the opposite bank under a dramatic, overcast sky filled with heavy clouds.

*Tons of Addresses
Arriving Every Second*

Architect - First Reaction...



Second Reaction...



Developer First Reaction



Second Reaction



Problems ?

- Tedious: Spend time configuring where to send messages, deploying workers, and deploying intermediate queues.
- Brittle: There's little fault-tolerance.
- Painful to scale: Partition of running worker/s is complicated.

What We Want ?

- Horizontal scalability
- Fault-tolerance
- No intermediate message brokers!
- Higher level abstraction than message passing
- “Just works”
- Guaranteed data processing (not in this case)

Apache Storm

- ✓ Horizontal scalability
- ✓ Fault-tolerance
- ✓ No intermediate message brokers!
- ✓ Higher level abstraction than message passing
- ✓ “Just works”
- ✓ Guaranteed data processing

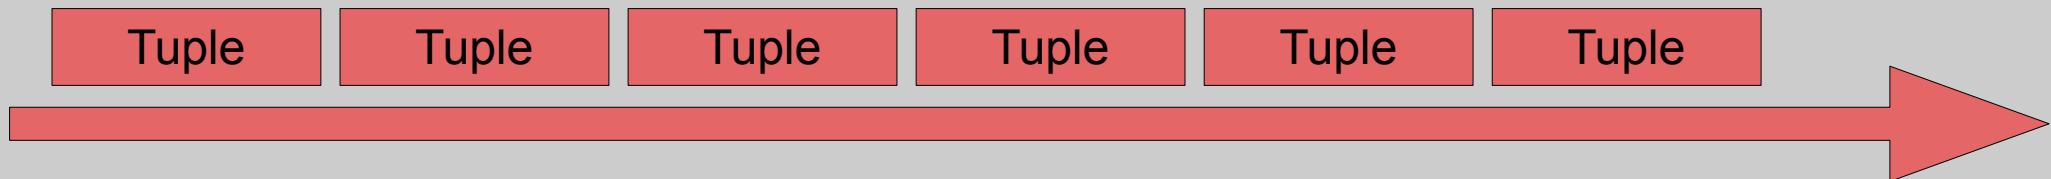
Anatomy of Storm



What is Storm ?

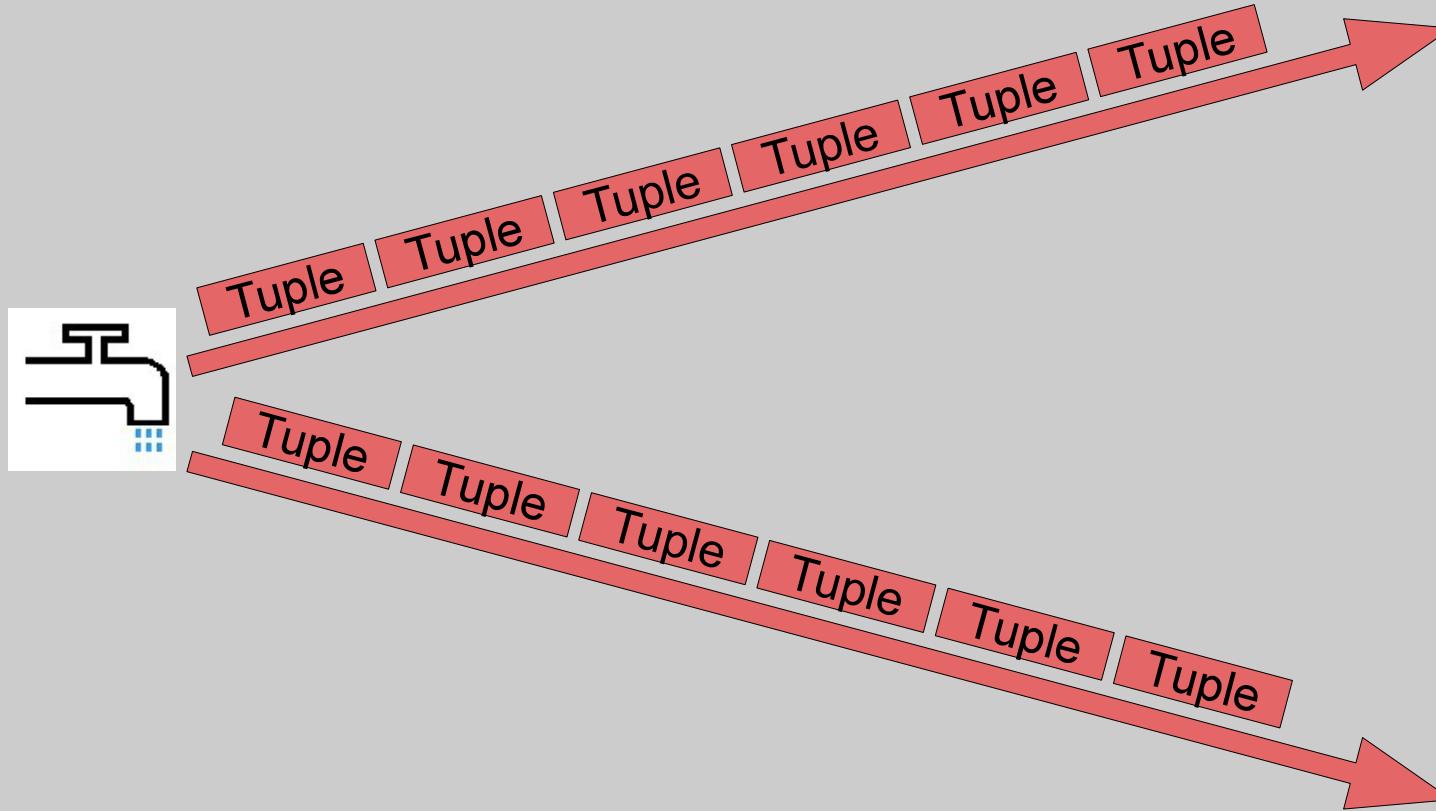
- CEP - Open source and distributed realtime computation system.
 - Makes it easy to reliably process unbounded streams of tuples
 - Doing for realtime processing what Hadoop did for batch processing.
- Fast - 1M Tuples/sec per node.
 - It is scalable,fault-tolerant, guarantees your data will be processed, and is easy to set up and operate.

Streams



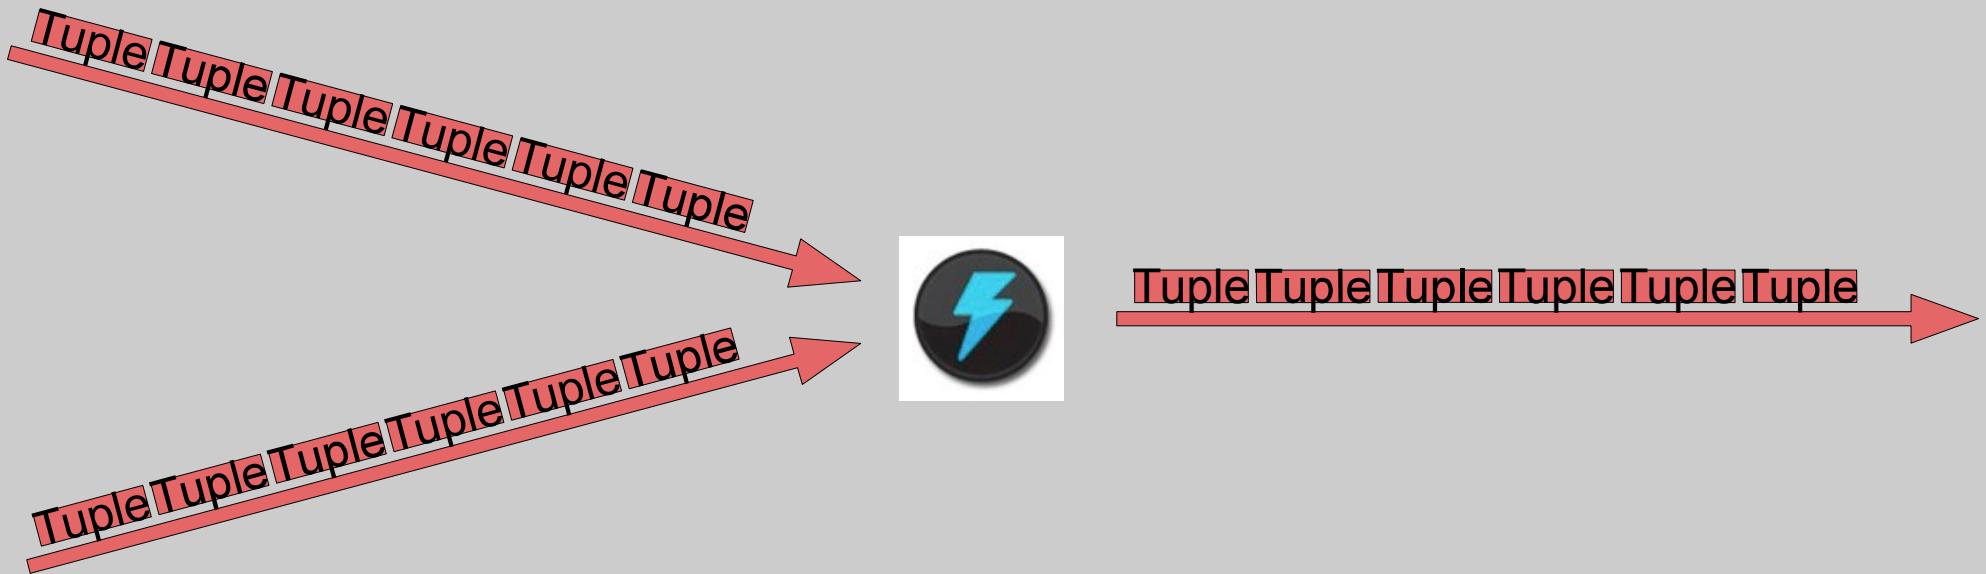
Unbounded sequence of tuples

Spouts



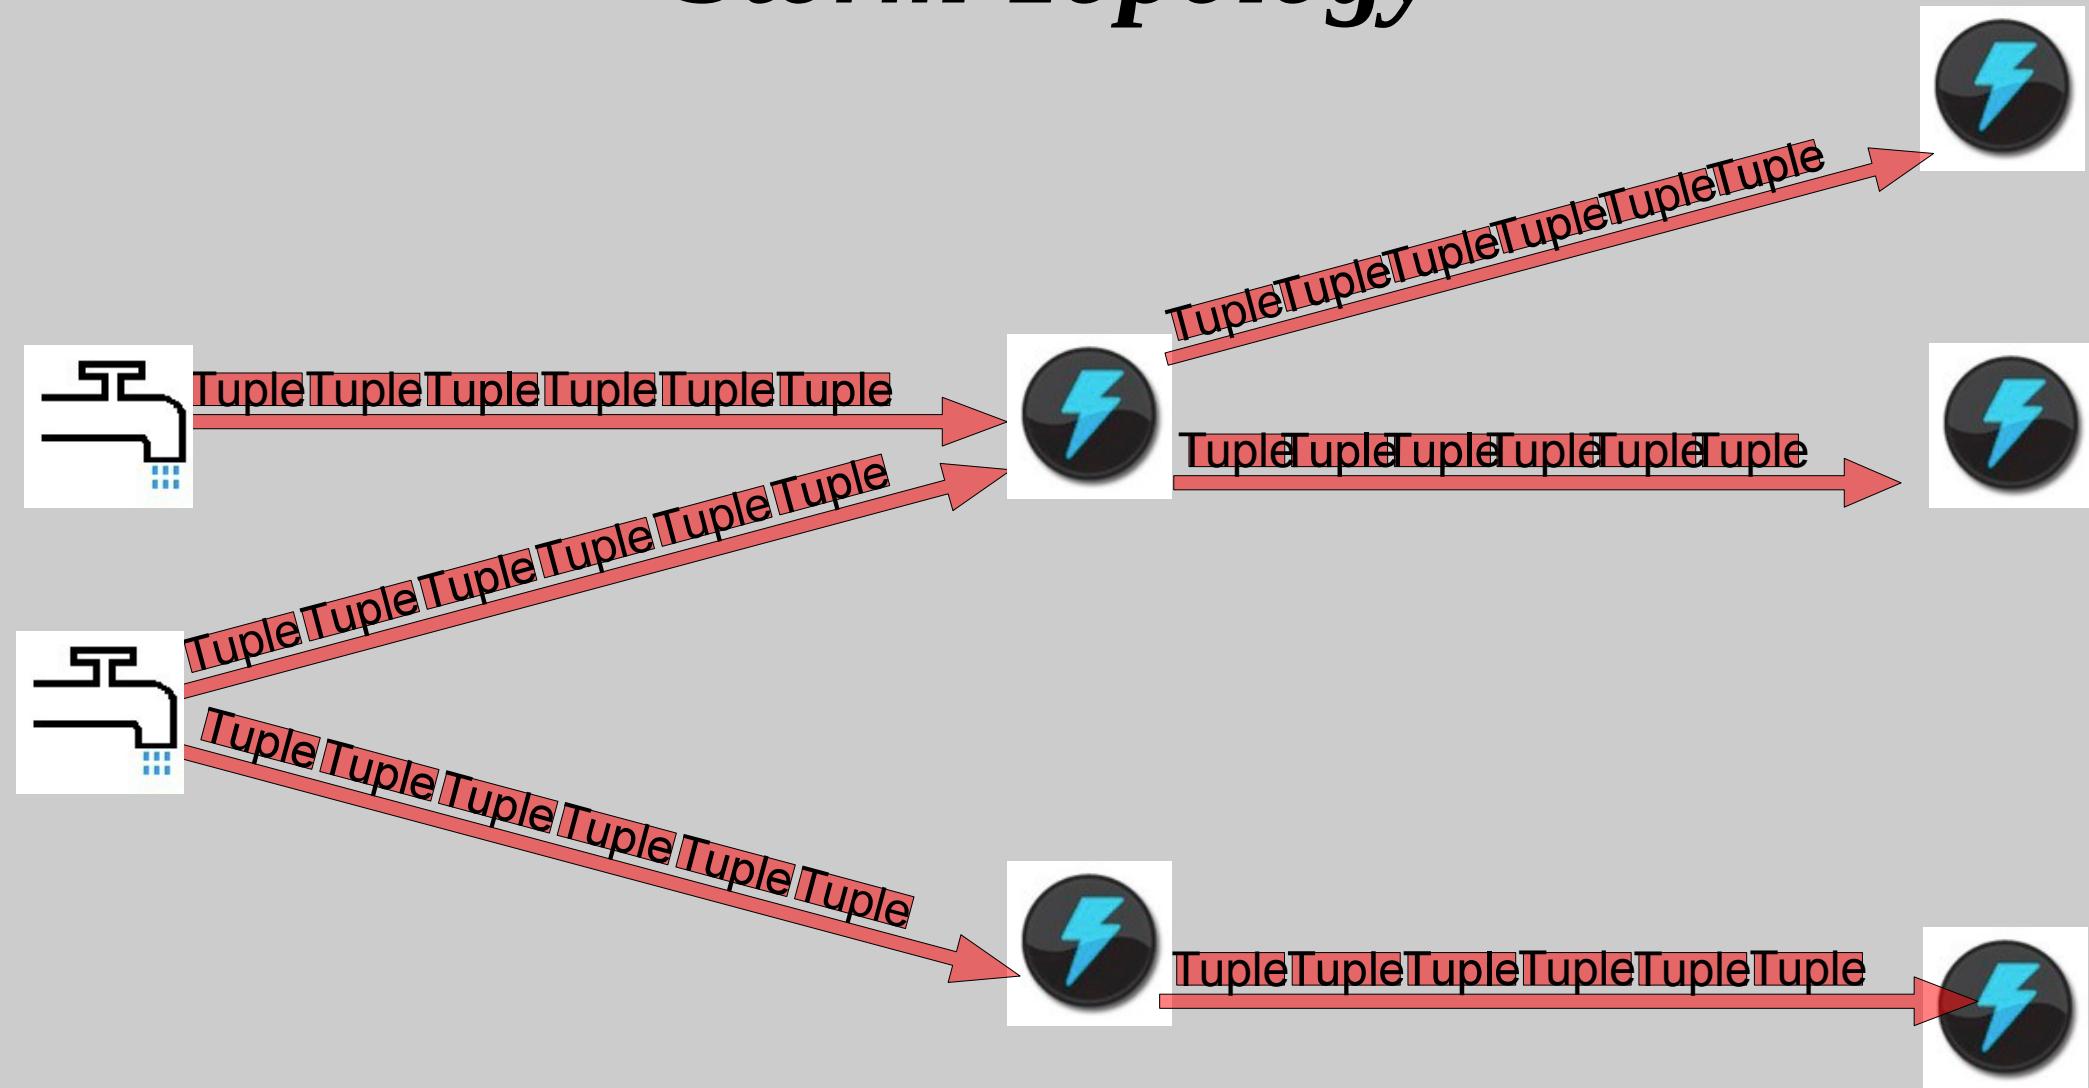
Sources of Streams

Bolts



Processes input streams and produces new streams

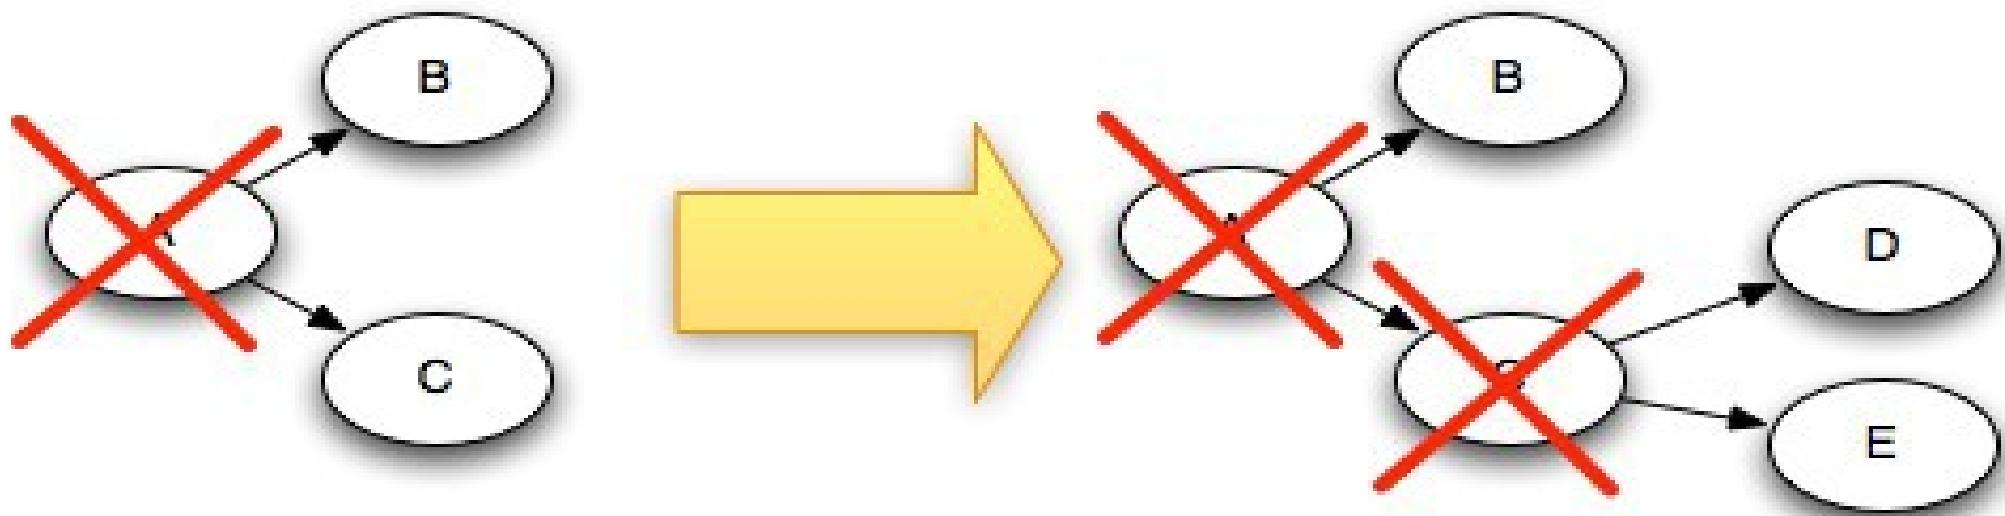
Storm Topology



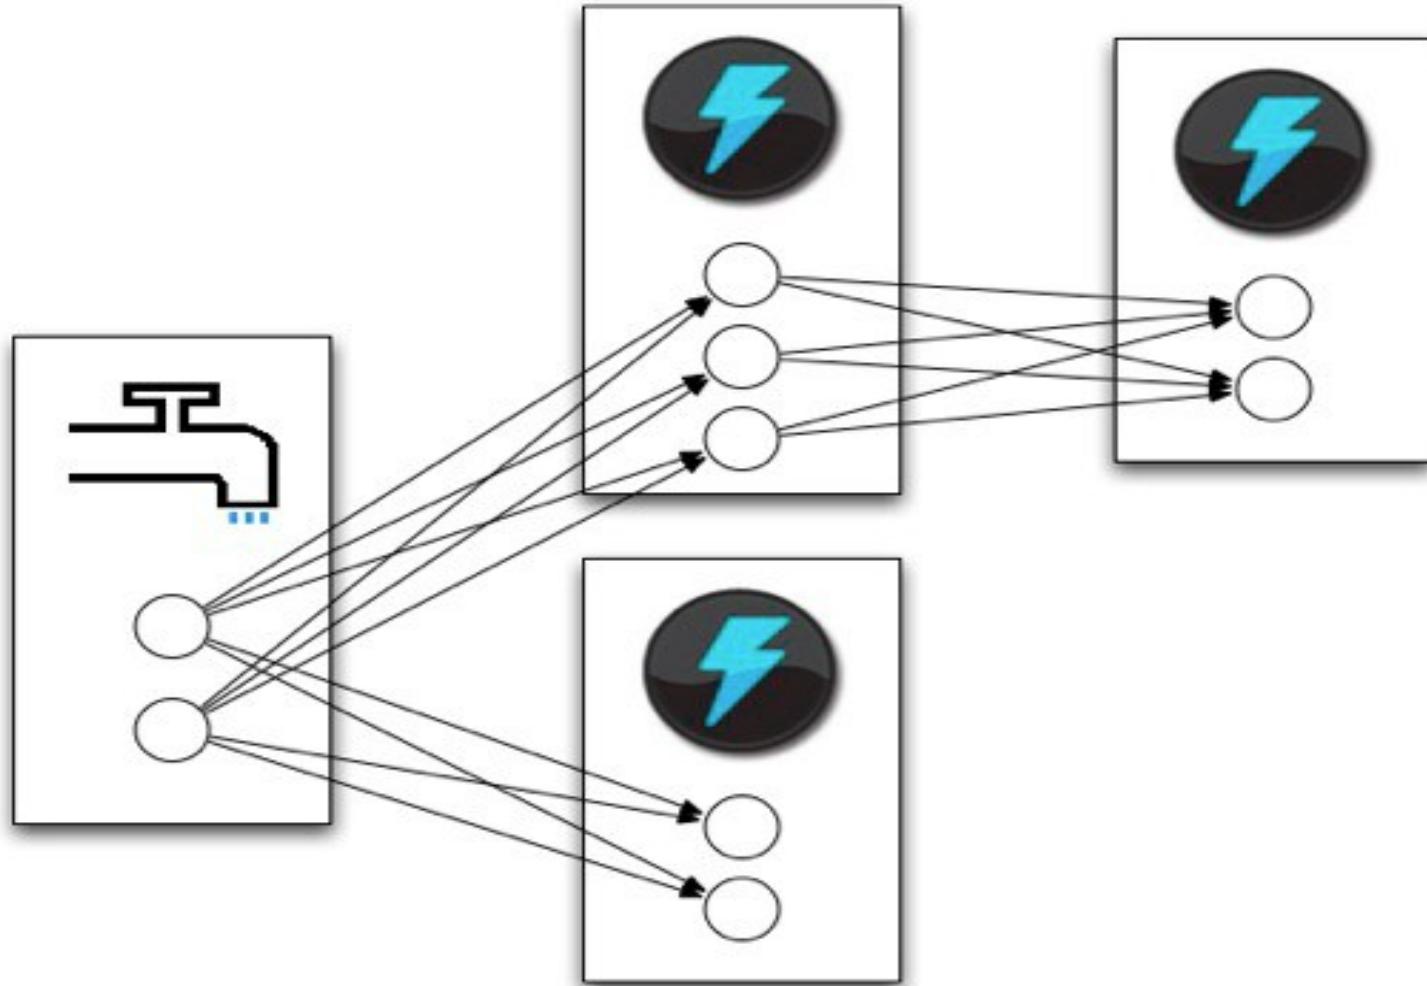
Network of spouts and bolts

Guarantee for Processing

- Storm guarantees the full processing of a tuple by tracking its state
- In case of failure, Storm can re-process it.
- Source tuples with full “acked” trees are removed from the system

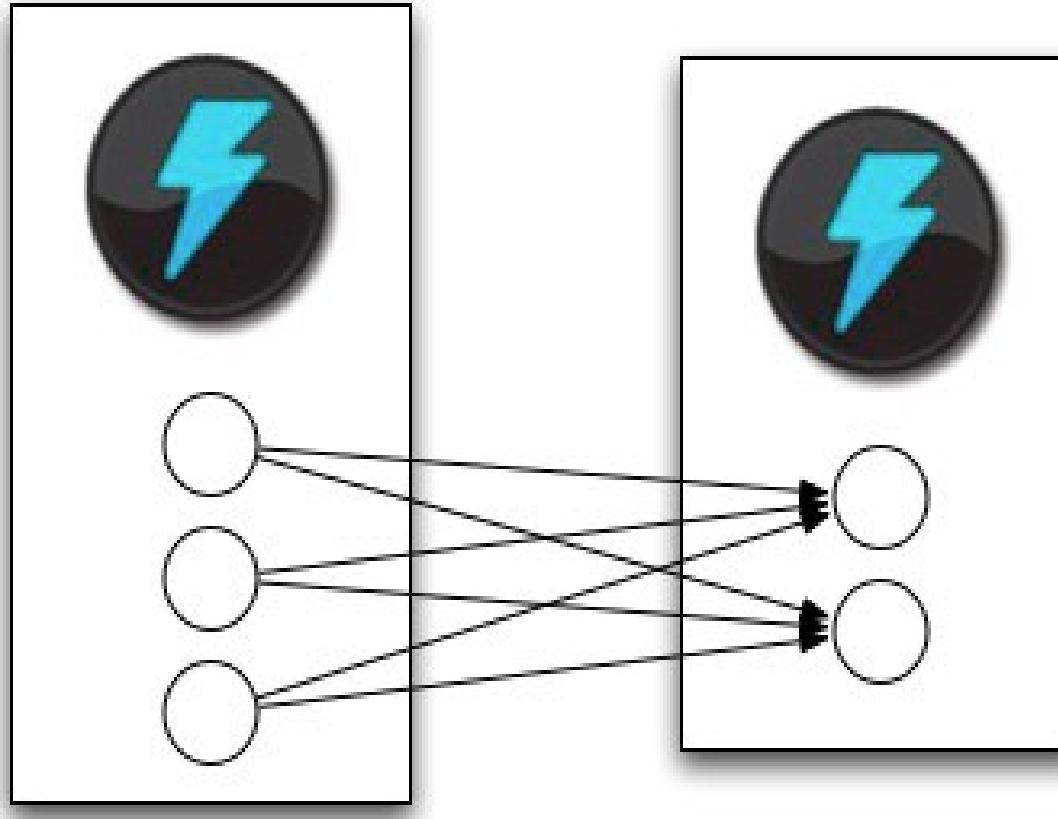


Tasks (Bolt/Spout Instance)



Spouts and bolts execute as many tasks across the cluster

Stream Grouping

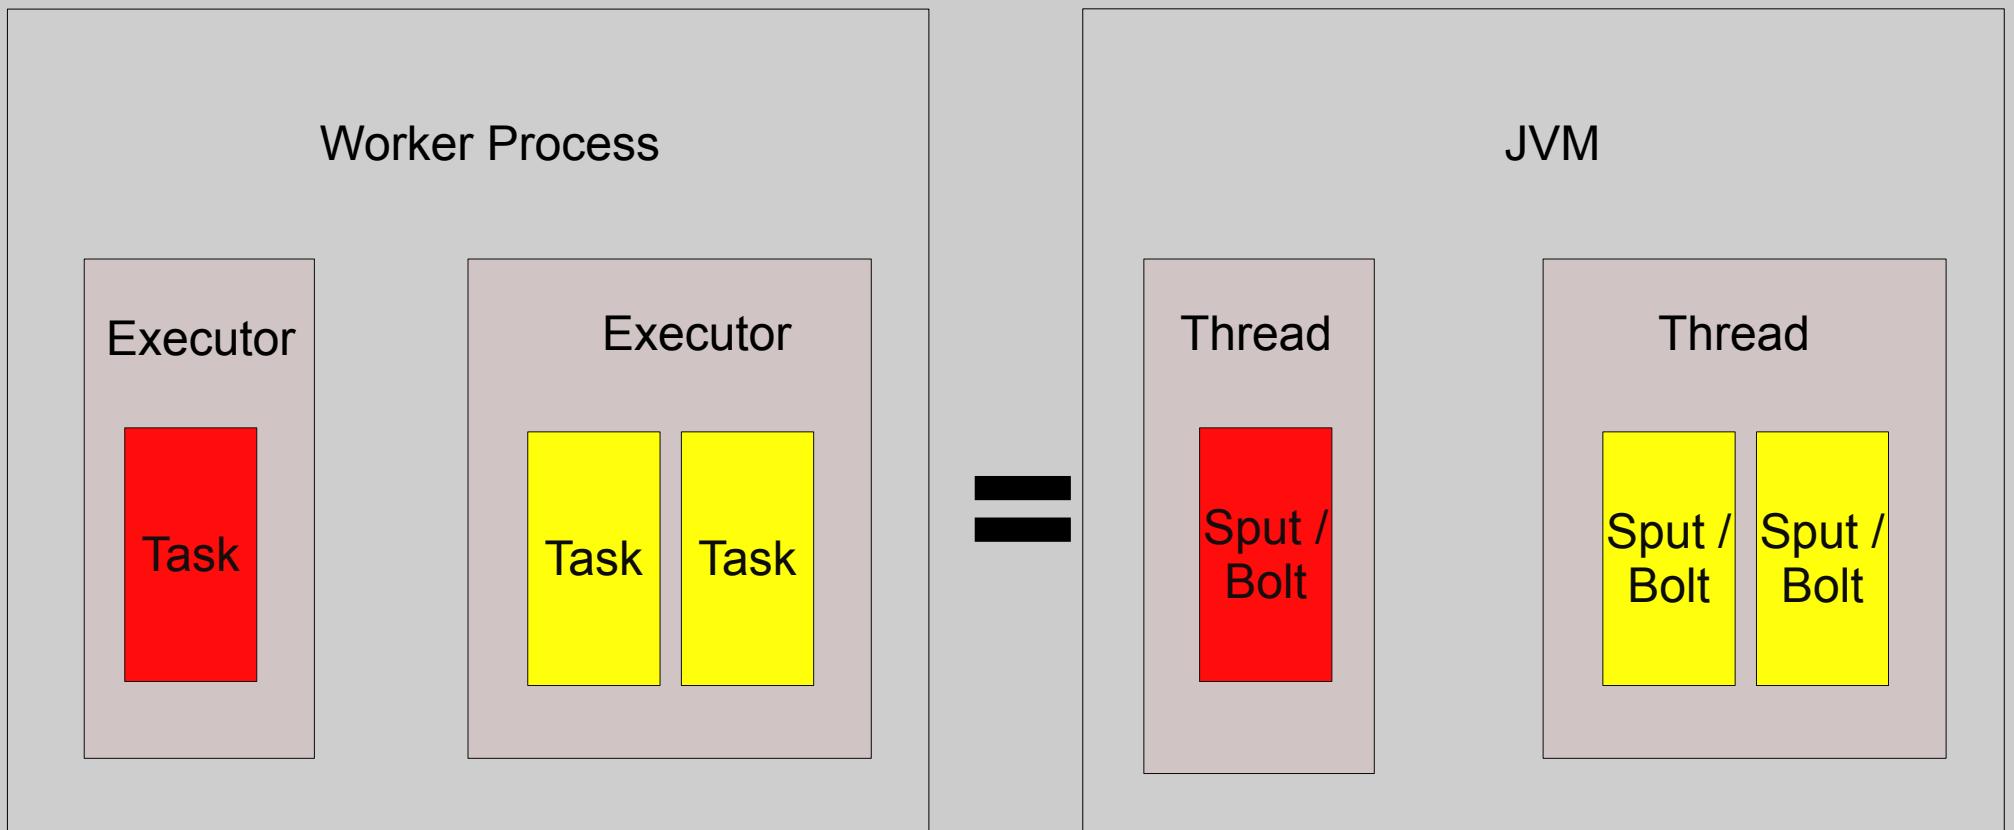


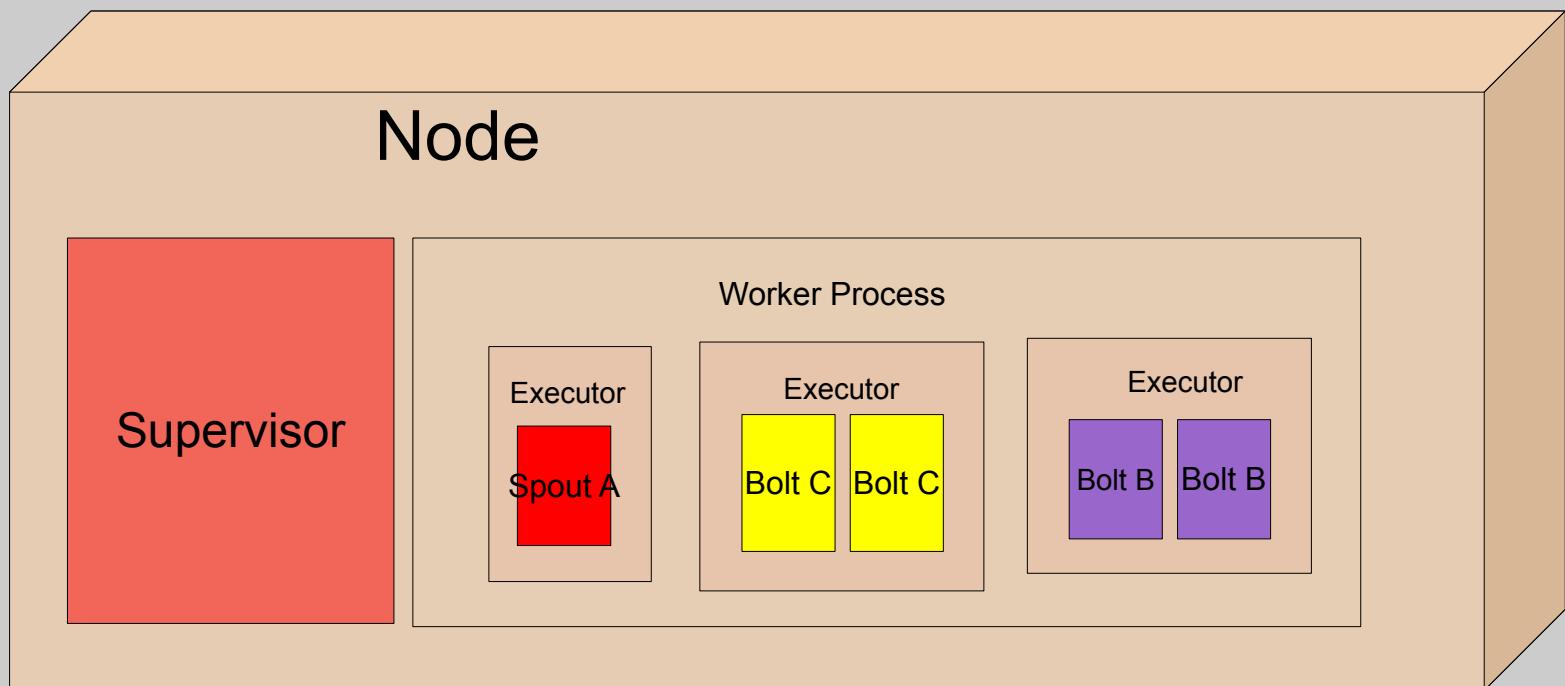
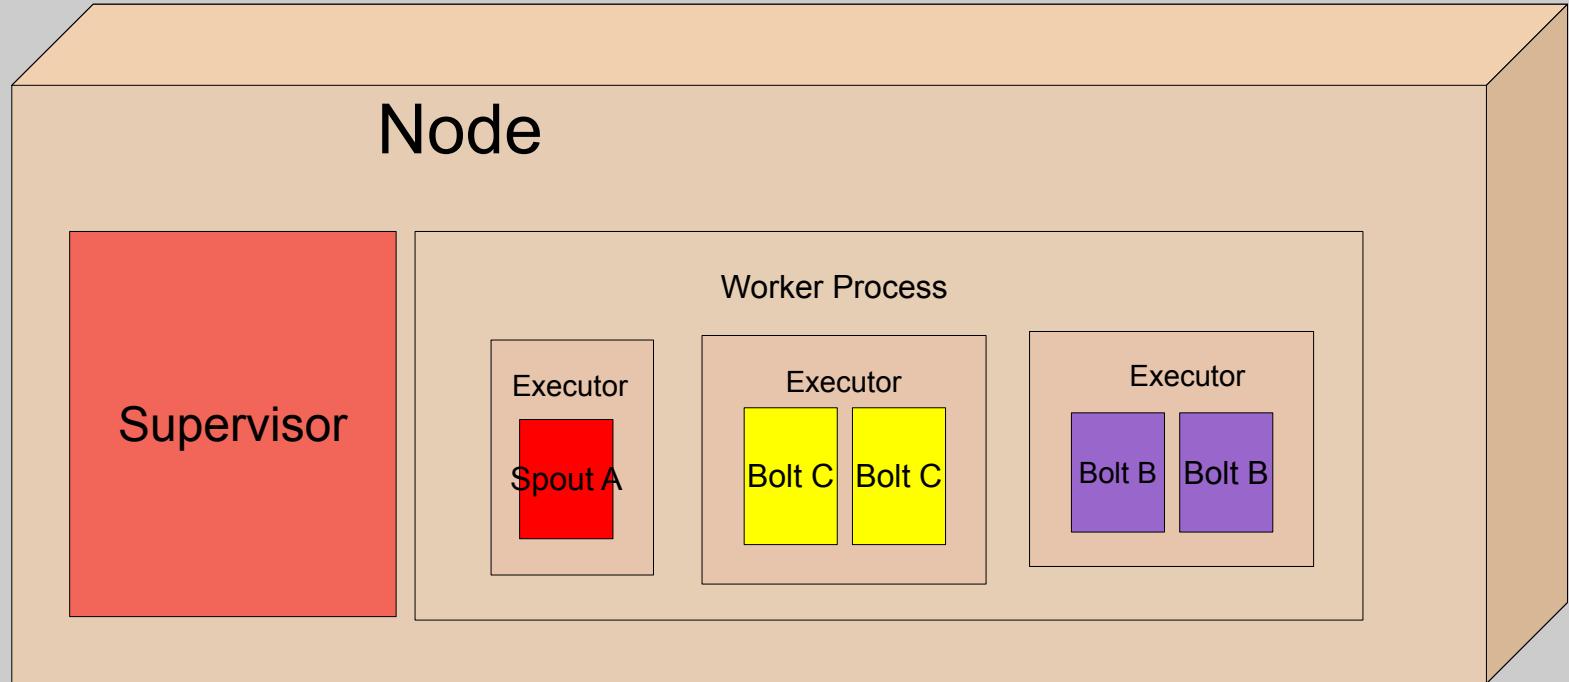
When a tuple is emitted, which task (instance) does it go to?

Stream Grouping

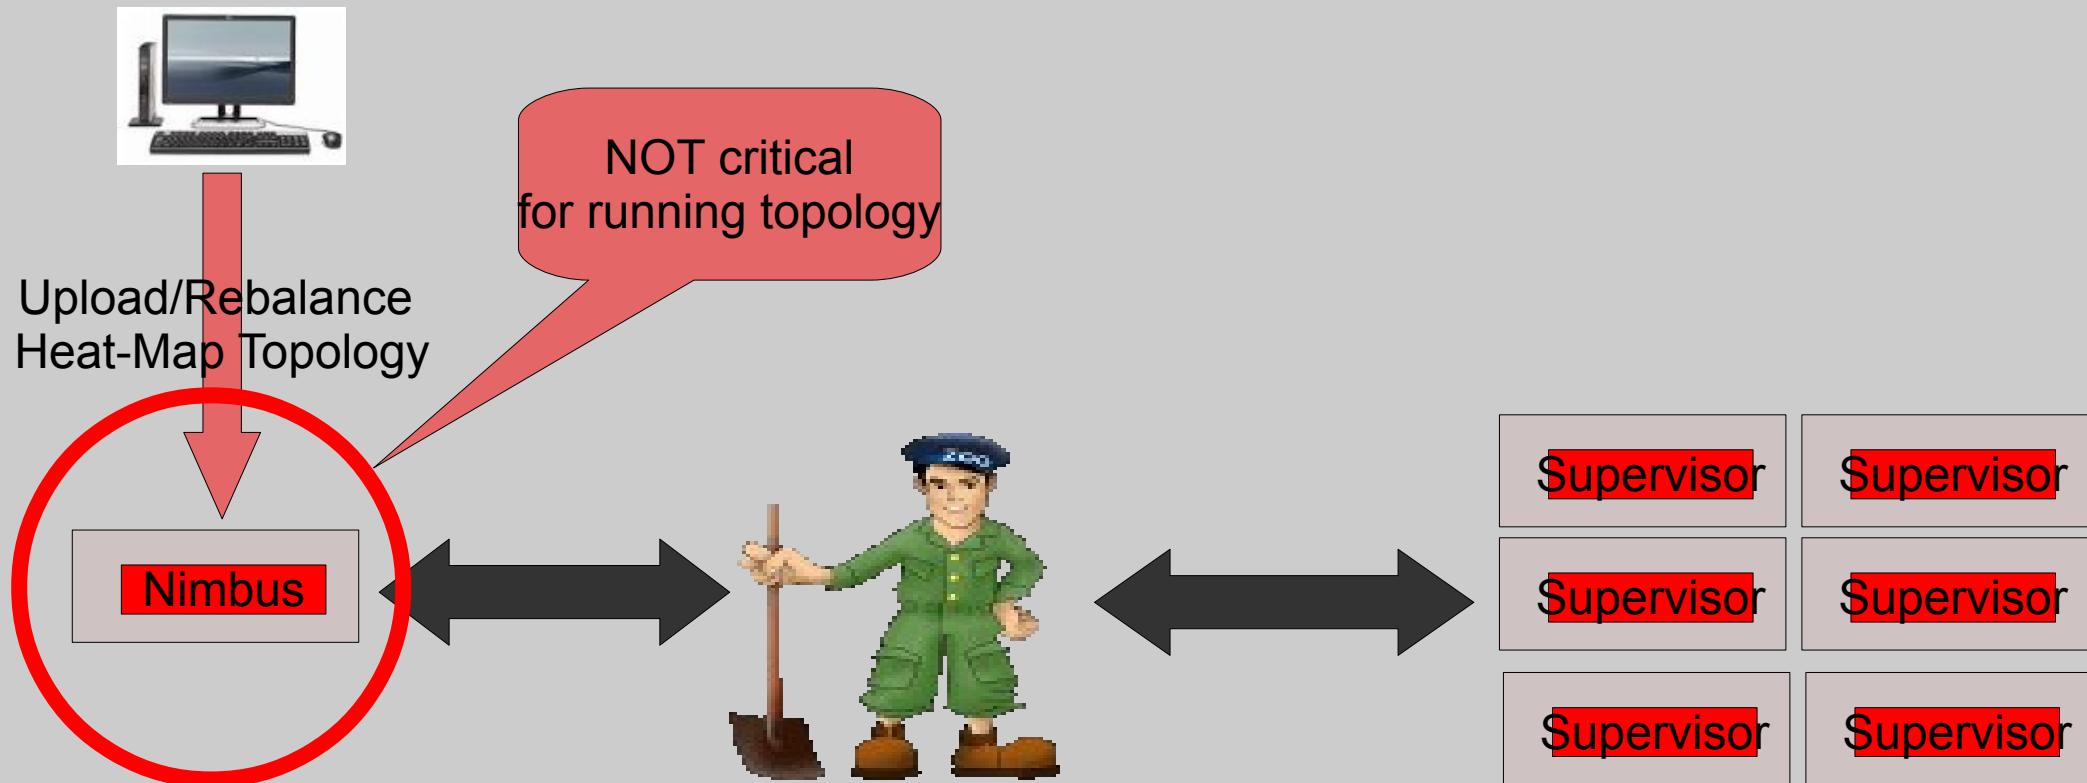
- Shuffle grouping: pick a random task
- Fields grouping: consistent hashing on a subset of tuple fields
- All grouping: send to all tasks
- Global grouping: pick task with lowest id

Tasks , Executors , Workers



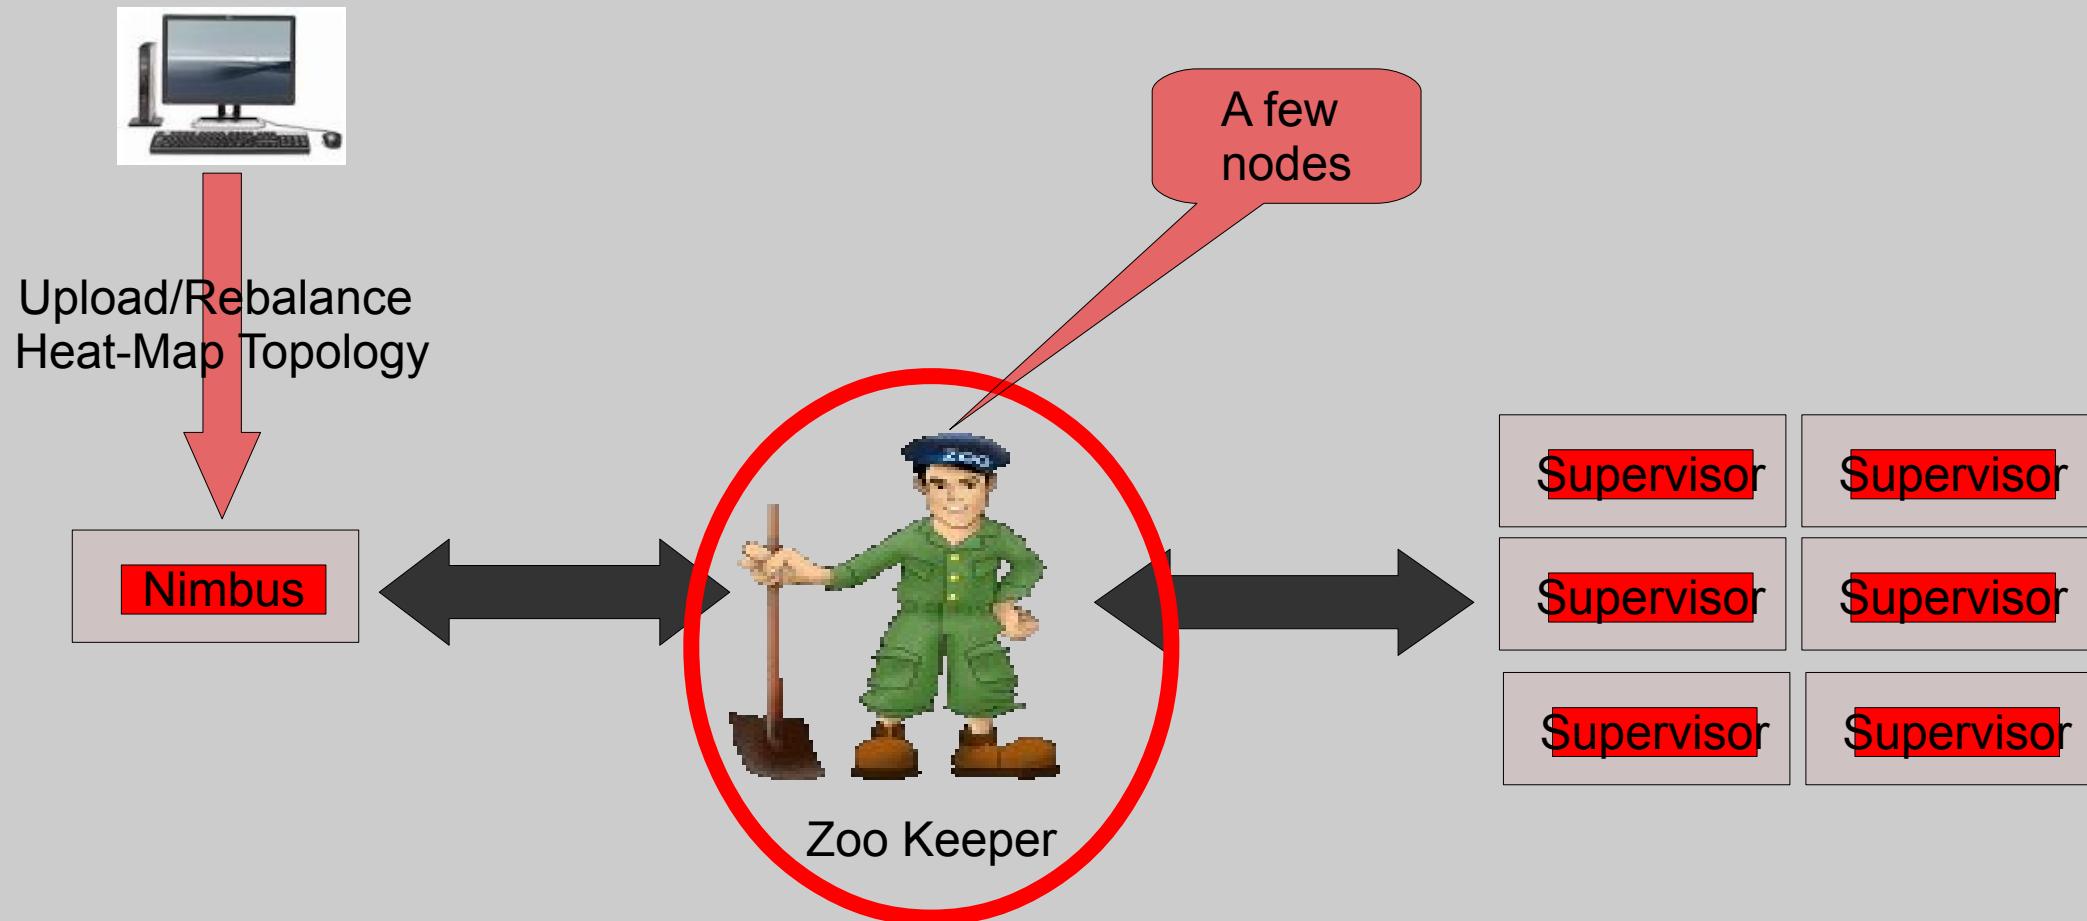


Storm Architecture



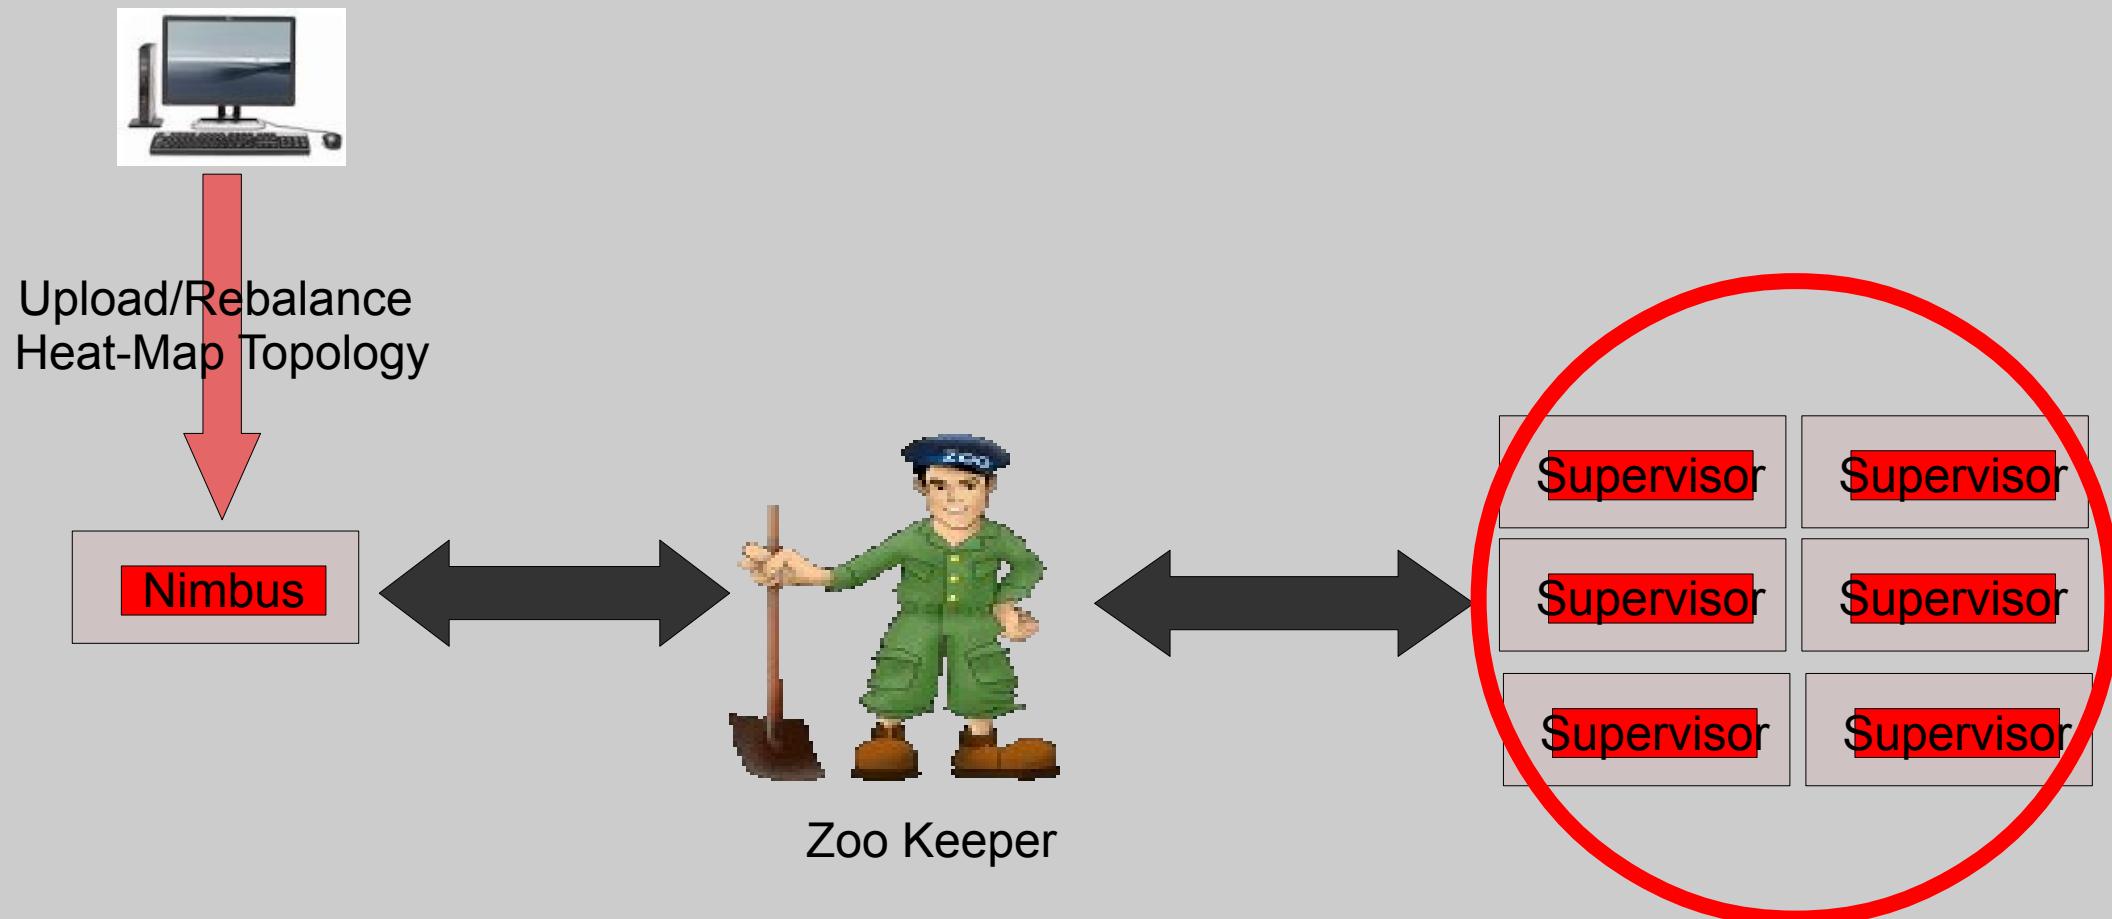
*Master Node
(similar to Hadoop JobTracker)*

Storm Architecture



Used For Cluster Coordination

Storm Architecture



Run Worker Processes

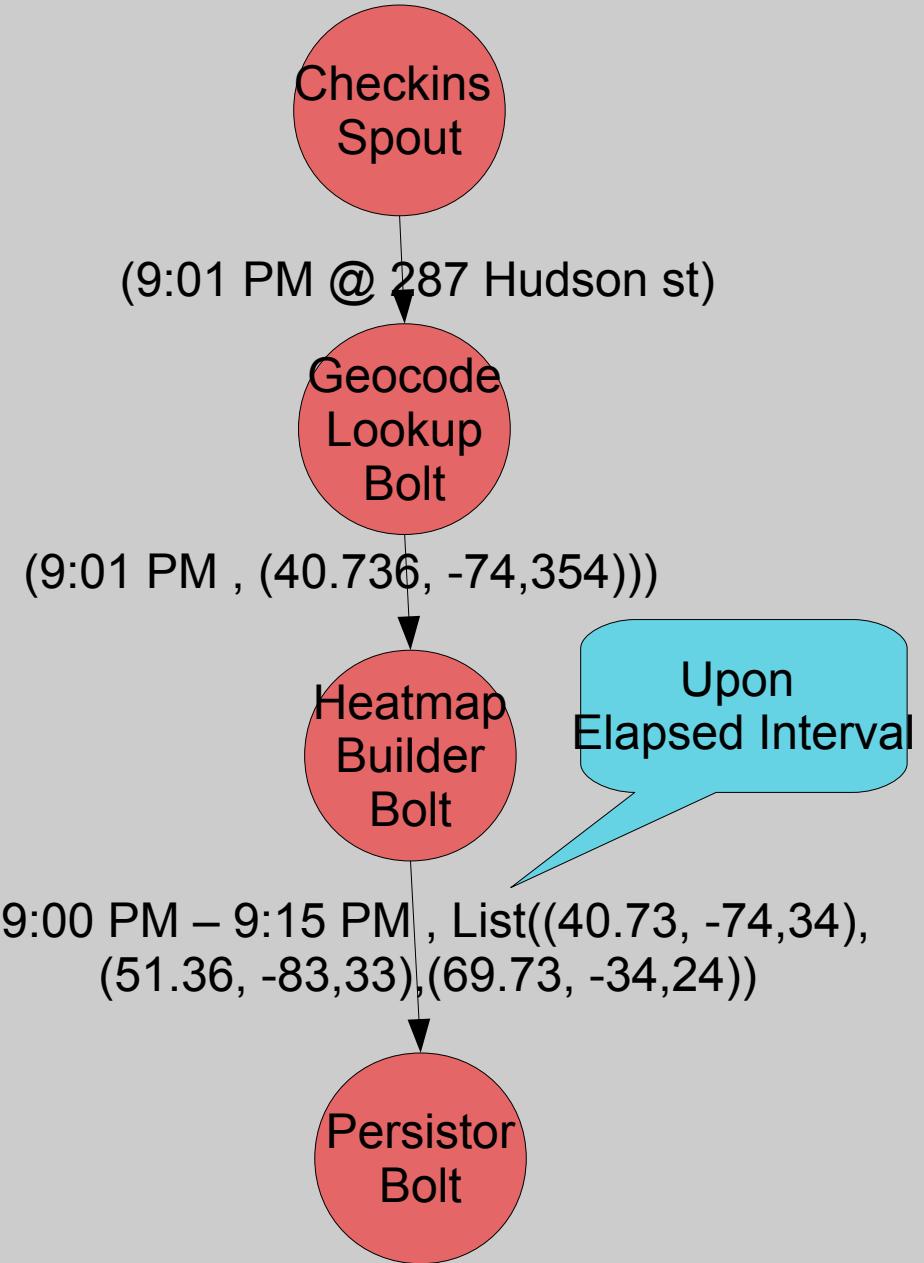
Assembling Heatmap Topology



HeatMap Input/Output Tuples

- Input Tuples: Timestamp and Text Address :
 - (9:00:07 PM , “287 Hudson St New York NY 10013”)
- Output Tuple: Time interval, and a list of points for it:
 - (9:00:00 PM to 9:00:15 PM,
List((40.719,-73.987),(40.726,-74.001),(40.719,-73.987)))

Heat Map Storm Topology



Checkins Spout

```
public class CheckinsSpout extends BaseRichSpout {  
    private List<String> sampleLocations;  
    private int nextEmitIndex;  
    private SpoutOutputCollector outputCollector;  
  
    @Override  
    public void open(Map map, TopologyContext topologyContext,  
                     SpoutOutputCollector spoutOutputCollector) {  
        this.outputCollector = spoutOutputCollector;  
        this.nextEmitIndex = 0;  
        sampleLocations = IOUtils.readLines(  
            ClassLoader.getSystemResourceAsStream("sample-locations.txt"));  
    }  
  
    @Override  
    public void nextTuple() {  
        String address = checkins.get(nextEmitIndex);  
        String checkin = new Date().getTime() + "@ADDRESS:" + address;  
        outputCollector.emit(new Values(checkin));  
        nextEmitIndex = (nextEmitIndex + 1) % sampleLocations.size();  
    }  
  
    @Override  
    public void declareOutputFields(OutputFieldsDeclarer declarer) {  
        declarer.declare(new Fields("str"));  
    }  
}
```

We hold state
No need for thread safety

Been called iteratively by Storm

Declare output fields

Geocode Lookup Bolt

```
public class GeocodeLookupBolt extends BaseBasicBolt {  
    private LocatorService locatorService;  
  
    @Override  
    public void prepare(Map stormConf, TopologyContext context) {  
        locatorService = new GoogleLocatorService();  
    }  
  
    @Override  
    public void execute(Tuple tuple, BasicOutputCollector outputCollector) {  
        String str = tuple.getStringByField("str");  
        String[] parts = str.split("@");  
        Long time = Long.valueOf(parts[0]);  
        String address = parts[1];  
  
        LocationDTO locationDTO = locatorService.getLocation(address);  
        if(locationDTO!=null)  
            outputCollector.emit(new Values(time,locationDTO) );  
    }  
  
    @Override  
    public void declareOutputFields(OutputFieldsDeclarer fieldsDeclarer) {  
        fieldsDeclarer.declare(new Fields("time", "location"));  
    }  
}
```



Get Geocode,
Create DTO

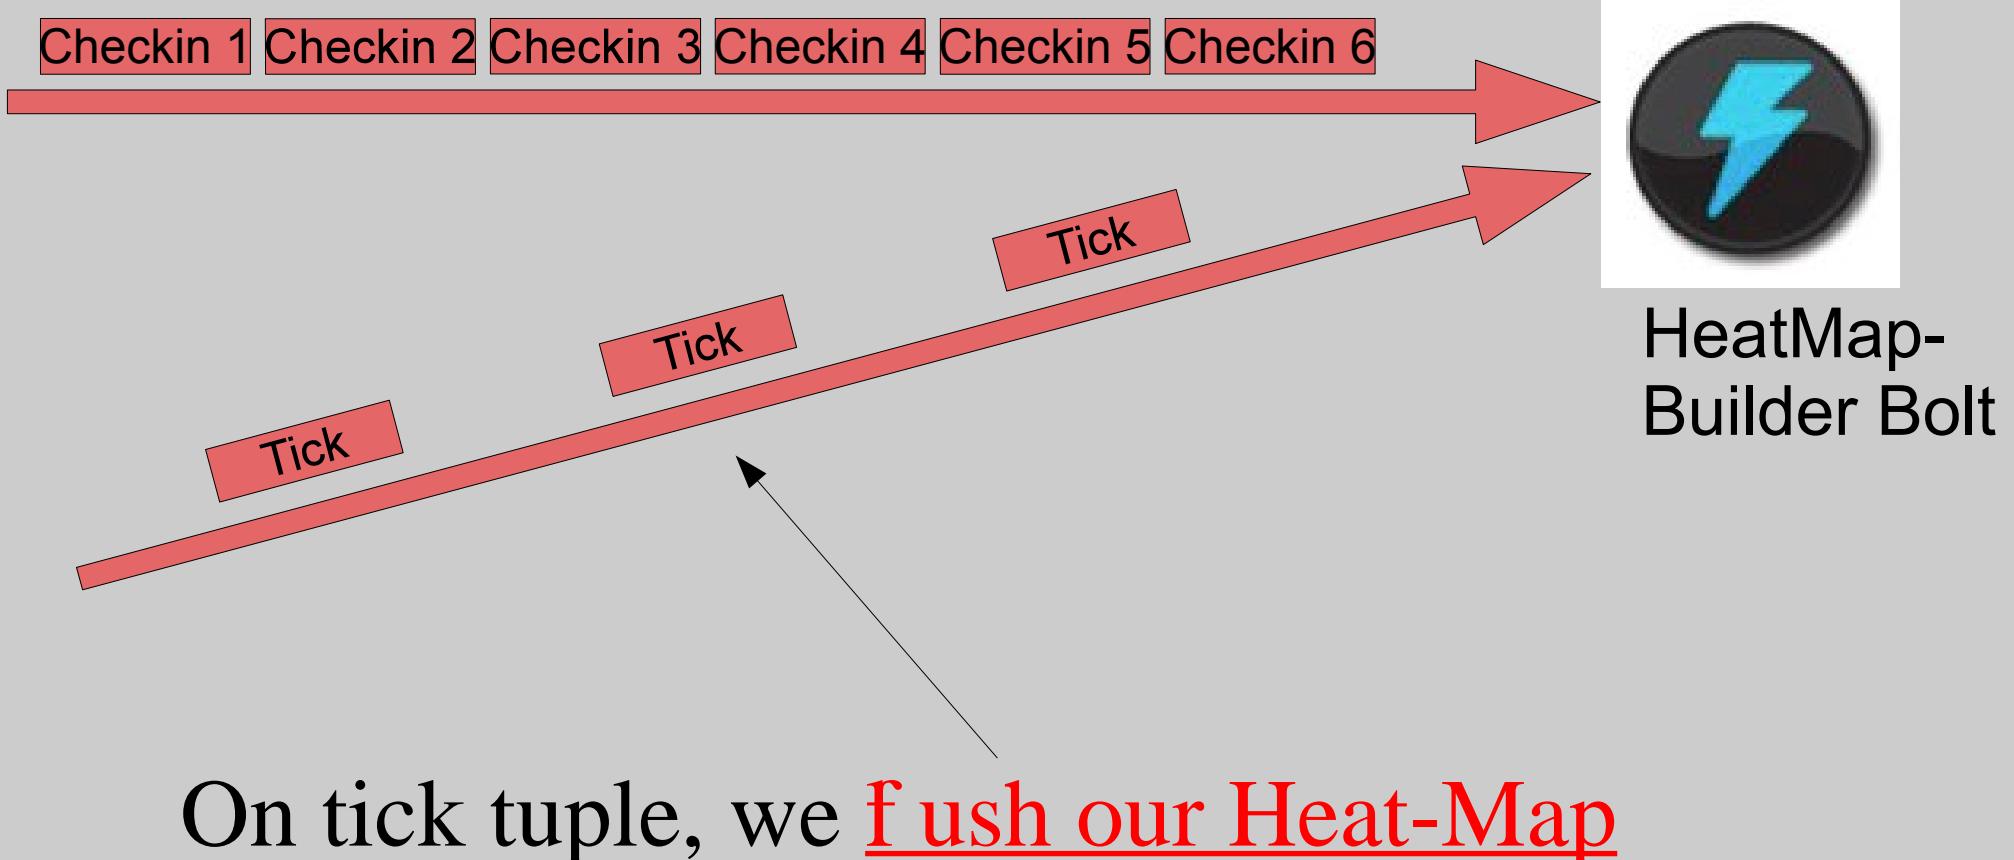
Tick Tuple – Repeating Mantra

remember diet

remember diet

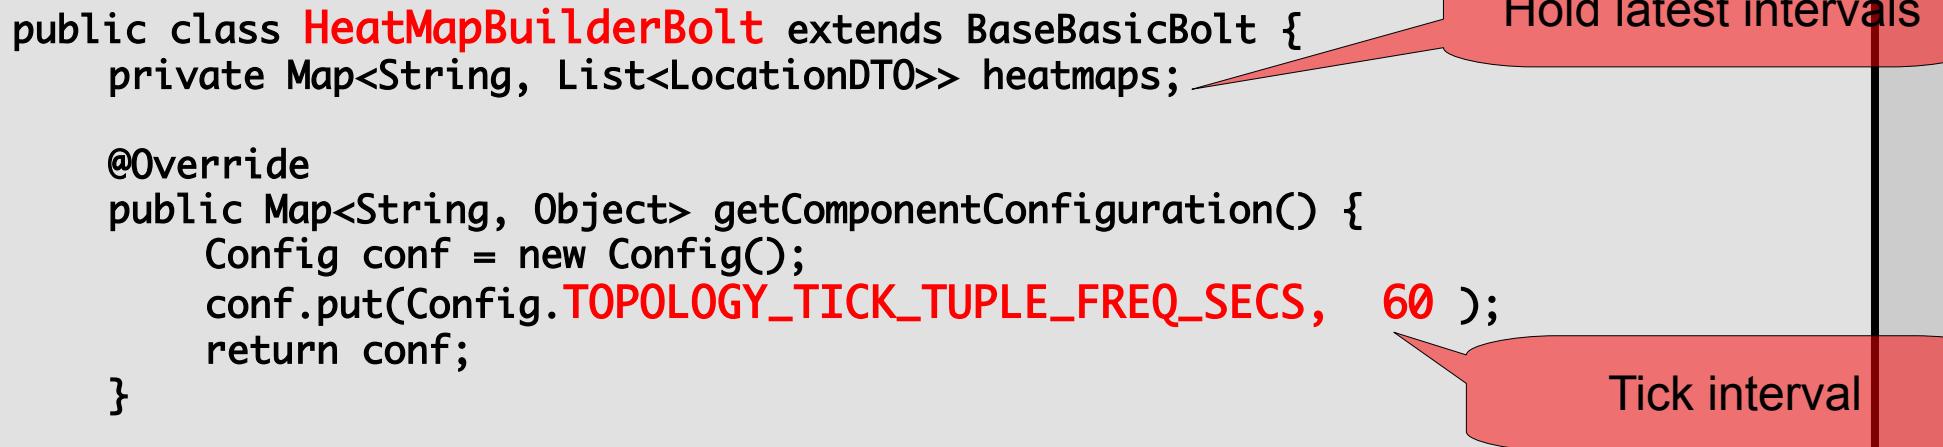
remember diet

Two Streams to Heat-Map Builder



Tick Tuple in Action

```
public class HeatMapBuilderBolt extends BaseBasicBolt {  
    private Map<String, List<LocationDTO>> heatmaps;  
  
    @Override  
    public Map<String, Object> getComponentConfiguration() {  
        Config conf = new Config();  
        conf.put(Config.TOPOLOGY_TICK_TUPLE_FREQ_SECS, 60);  
        return conf;  
    }  
  
    @Override  
    public void execute(Tuple tuple, BasicOutputCollector outputCollector) {  
        if (isTickTuple(tuple)) {  
            // Emit accumulated intervals  
        } else {  
            // Add check-in info to the current interval in the Map  
        }  
    }  
  
    private boolean isTickTuple(Tuple tuple) {  
        return tuple.getSourceComponent().equals(Constants.SYSTEM_COMPONENT_ID)  
            && tuple.getSourceStreamId().equals(Constants.SYSTEM_TICK_STREAM_ID);  
    }  
}
```



Persister Bolt

```
public class PersistorBolt extends BaseBasicBolt {  
    private Jedis jedis;  
  
    @Override  
    public void execute(Tuple tuple, BasicOutputCollector outputCollector) {  
  
        Long timeInterval = tuple.getLongByField("time-interval");  
        String city = tuple.getStringByField("city");  
        String locationsList = objectMapper.writeValueAsString  
            (tuple.getValueByField("locationsList"));  
  
        String dbKey = "checkins-" + timeInterval+"@"+city; Persist in Redis  
for 24h  
        jedis.setex(dbKey, 3600*24 ,locationsList);  
  
        jedis.publish("location-key", dbKey);  
    } Publish in  
Redis channel  
for debugging
```

Transforming the Tuples

Sample Checkins File

```
Check-in #1  
Check-in #2  
Check-in #3  
Check-in #4  
Check-in #5  
Check-in #6  
Check-in #7  
Check-in #8  
Check-in #9
```

...

Read
Text Addresses

Checkins
Spout

Shuffle Grouping

Group by city

Get Geo
Location

Field Grouping(city)

Heatmap
Builder
Bolt

Shuffle Grouping

Persistor
Bolt

Geo Location
Service

Database

Heat Map Topology

```
public class LocalTopologyRunner {  
    public static void main(String[] args) {  
        TopologyBuilder builder = buildTopology();  
        StormSubmitter.submitTopology(  
            "local-heatmap", new Config(), builder.createTopology());  
  
    }  
  
    private static TopologyBuilder buildTopology() {  
        topologyBuilder builder = new TopologyBuilder();  
  
        builder.setSpout("checkins", new CheckinsSpout());  
  
        builder.setBolt("geocode-lookup", new GeocodeLookupBolt() )  
            .shuffleGrouping("checkins");  
  
        builder.setBolt("heatmap-builder", new HeatMapBuilderBolt() )  
            .fieldsGrouping("geocode-lookup", new Fields("city"));  
  
        builder.setBolt("persistor", new PersistorBolt() )  
            .shuffleGrouping("heatmap-builder");  
  
        return builder;  
    }  
}
```

Its NOT Scaled





FAIL.

Scaling the Topology

```
public class LocalTopologyRunner {  
    conf.setNumWorkers(20);  
    public static void main(String[] args) {  
        TopologyBuilder builder = buildTopology();  
        Config conf = new Config();  
        conf.setNumWorkers(2);  
        StormSubmitter.submitTopology(  
            "local-heatmap", conf, builder.createTopology());  
    }  
  
    private static TopologyBuilder buildTopology() {  
        topologyBuilder builder = new TopologyBuilder();  
  
        builder.setSpout("checkins", new CheckinsSpout(), 4);  
  
        builder.setBolt("geocode-lookup", new GeocodeLookupBolt() , 8)  
            .shuffleGrouping("checkins").setNumTasks(64);  
  
        builder.setBolt("heatmap-builder", new HeatMapBuilderBolt() , 4)  
            .fieldsGrouping("geocode-lookup", new Fields("city"));  
  
        builder.setBolt("persistor", new PersistorBolt() , 2 )  
            .shuffleGrouping("heatmap-builder").setNumTasks(4);49  
  
        return builder;  
    }  
}
```

The code is annotated with three callout boxes:

- A red callout points to the line `conf.setNumWorkers(20);` with the text "Set no. of workers".
- A red callout points to the line `conf.setNumWorkers(2);` with the text "Parallelism hint".
- A red callout points to the line `.setNumTasks(64);` with the text "Increase Tasks For Future".

Demo



Recap – Plan A

Sample Checkins File

```
Check-in #1  
Check-in #2  
Check-in #3  
Check-in #4  
Check-in #5  
Check-in #6  
Check-in #7  
Check-in #8  
Check-in #9  
...
```

Read
Text Address

Storm Heat-Map
Topology

GET Geo
Location

Geo Location
Service

Persist Checkin
Intervals

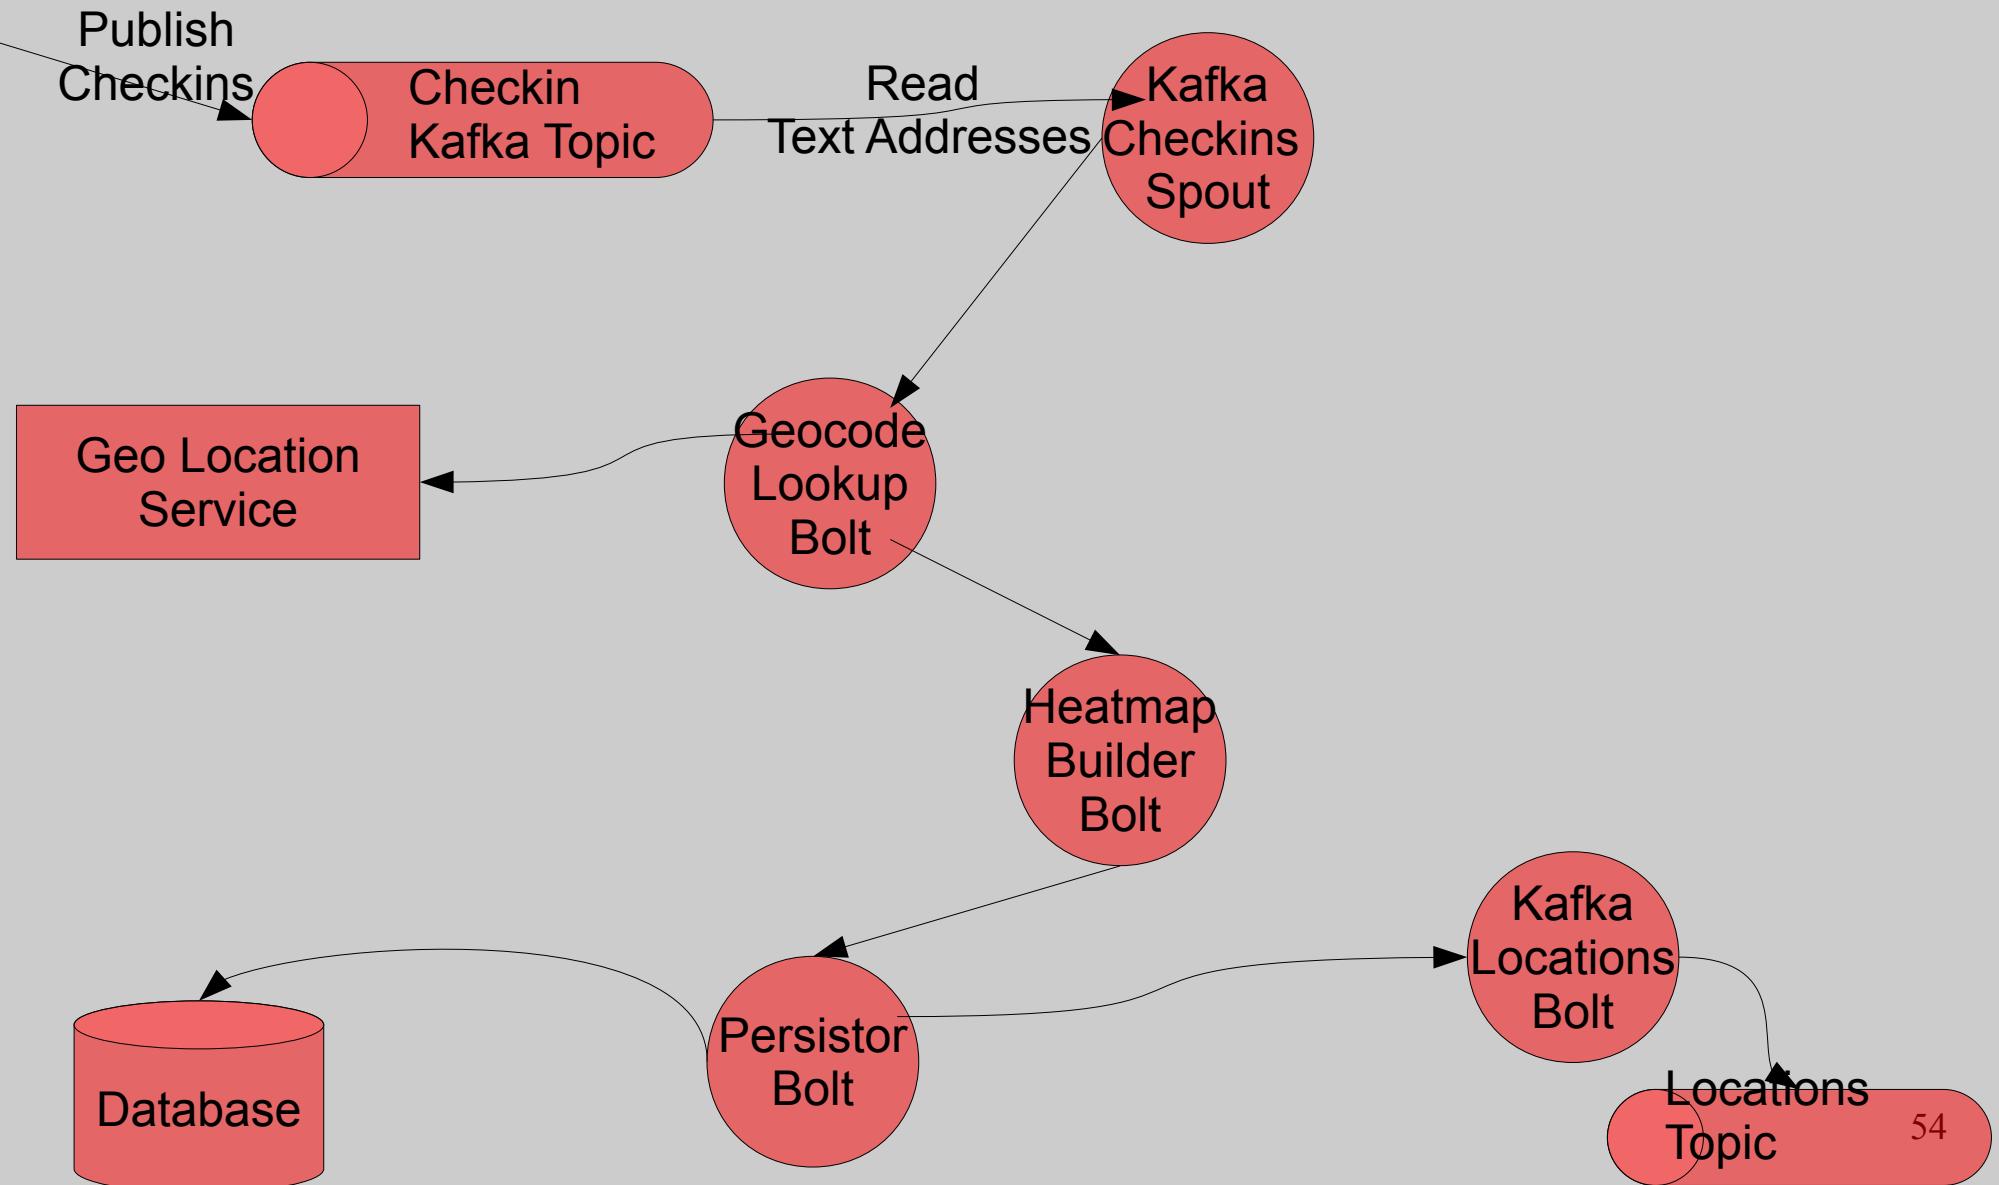


*We have
something working*



Add Kafka Messaging

Plan B - Kafka Spout&Bolt to HeatMap





Why Kafka?

When we have....
Aren't they Good?



They all are Good
But not for all use-cases

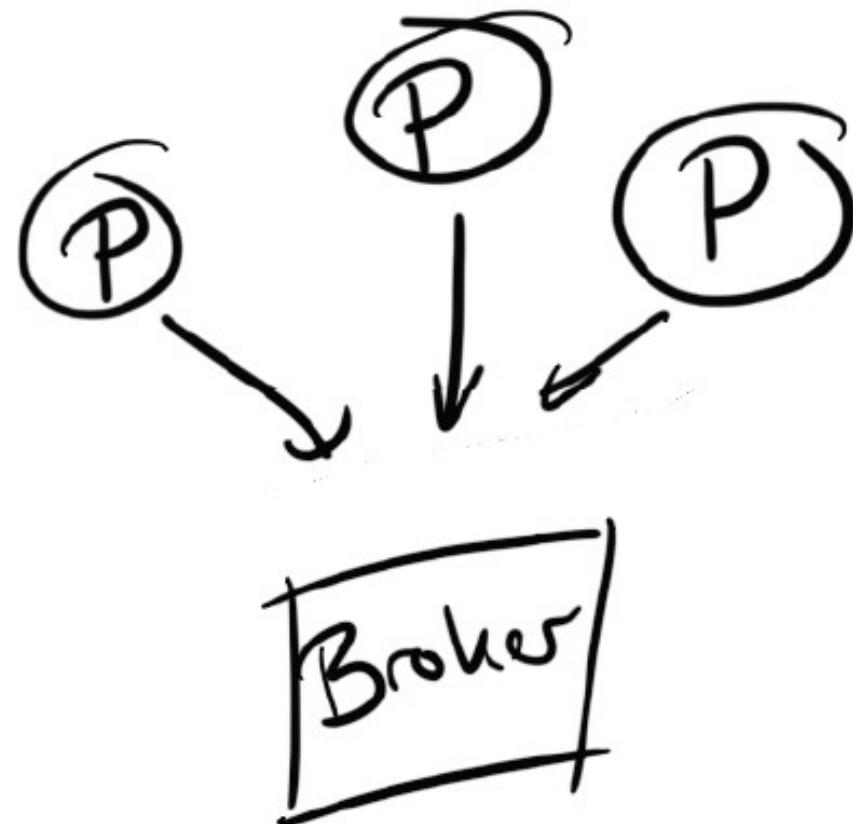
Kafka

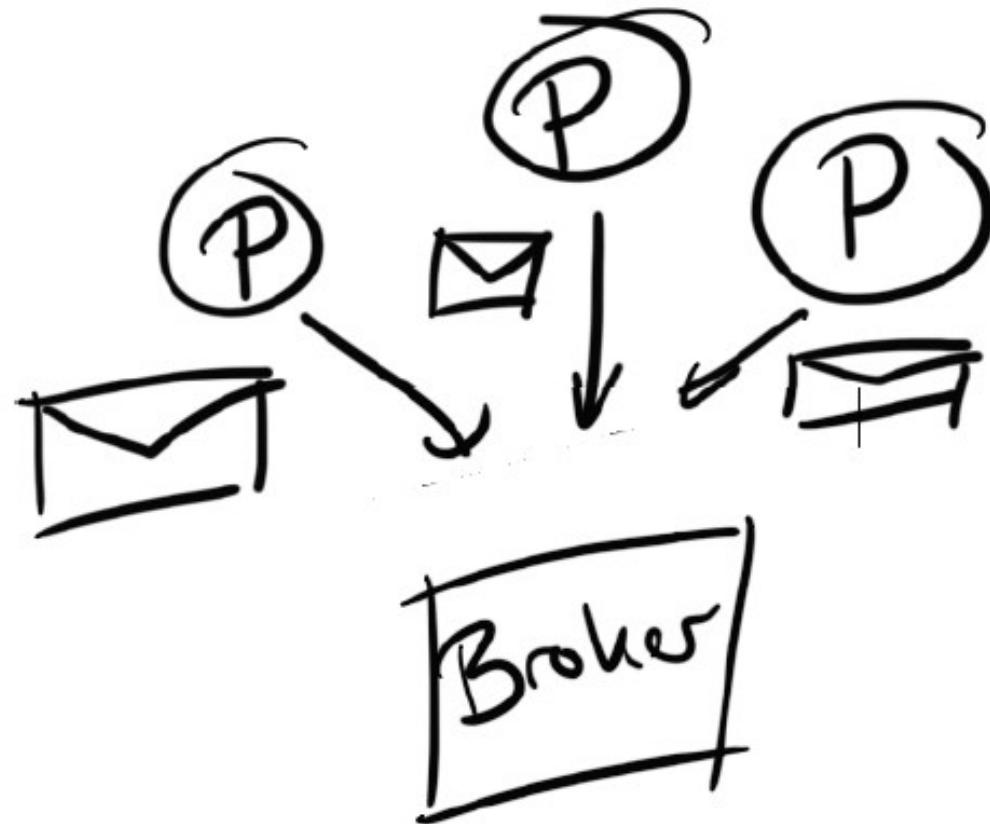
A little introduction

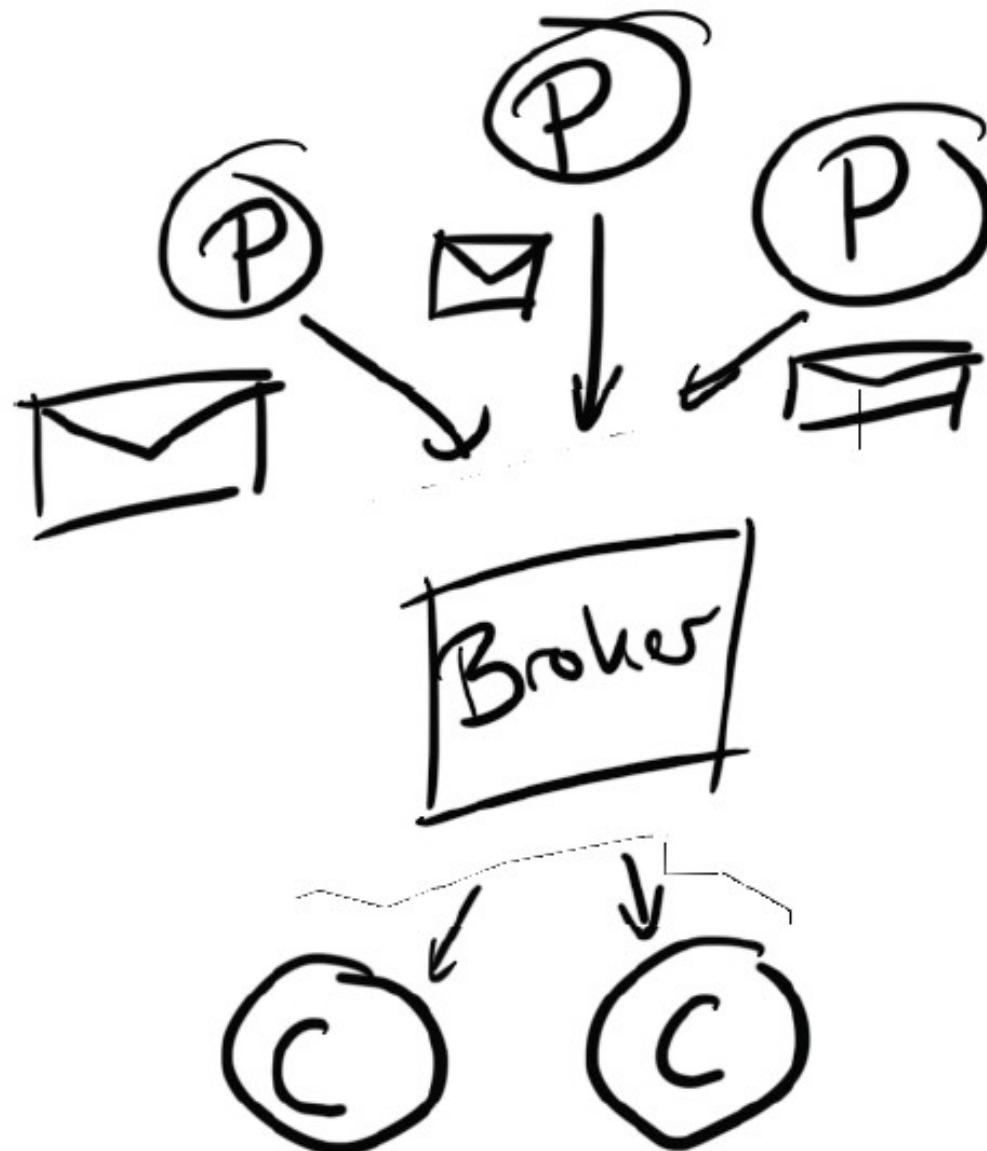


Pub-Sub Messaging System

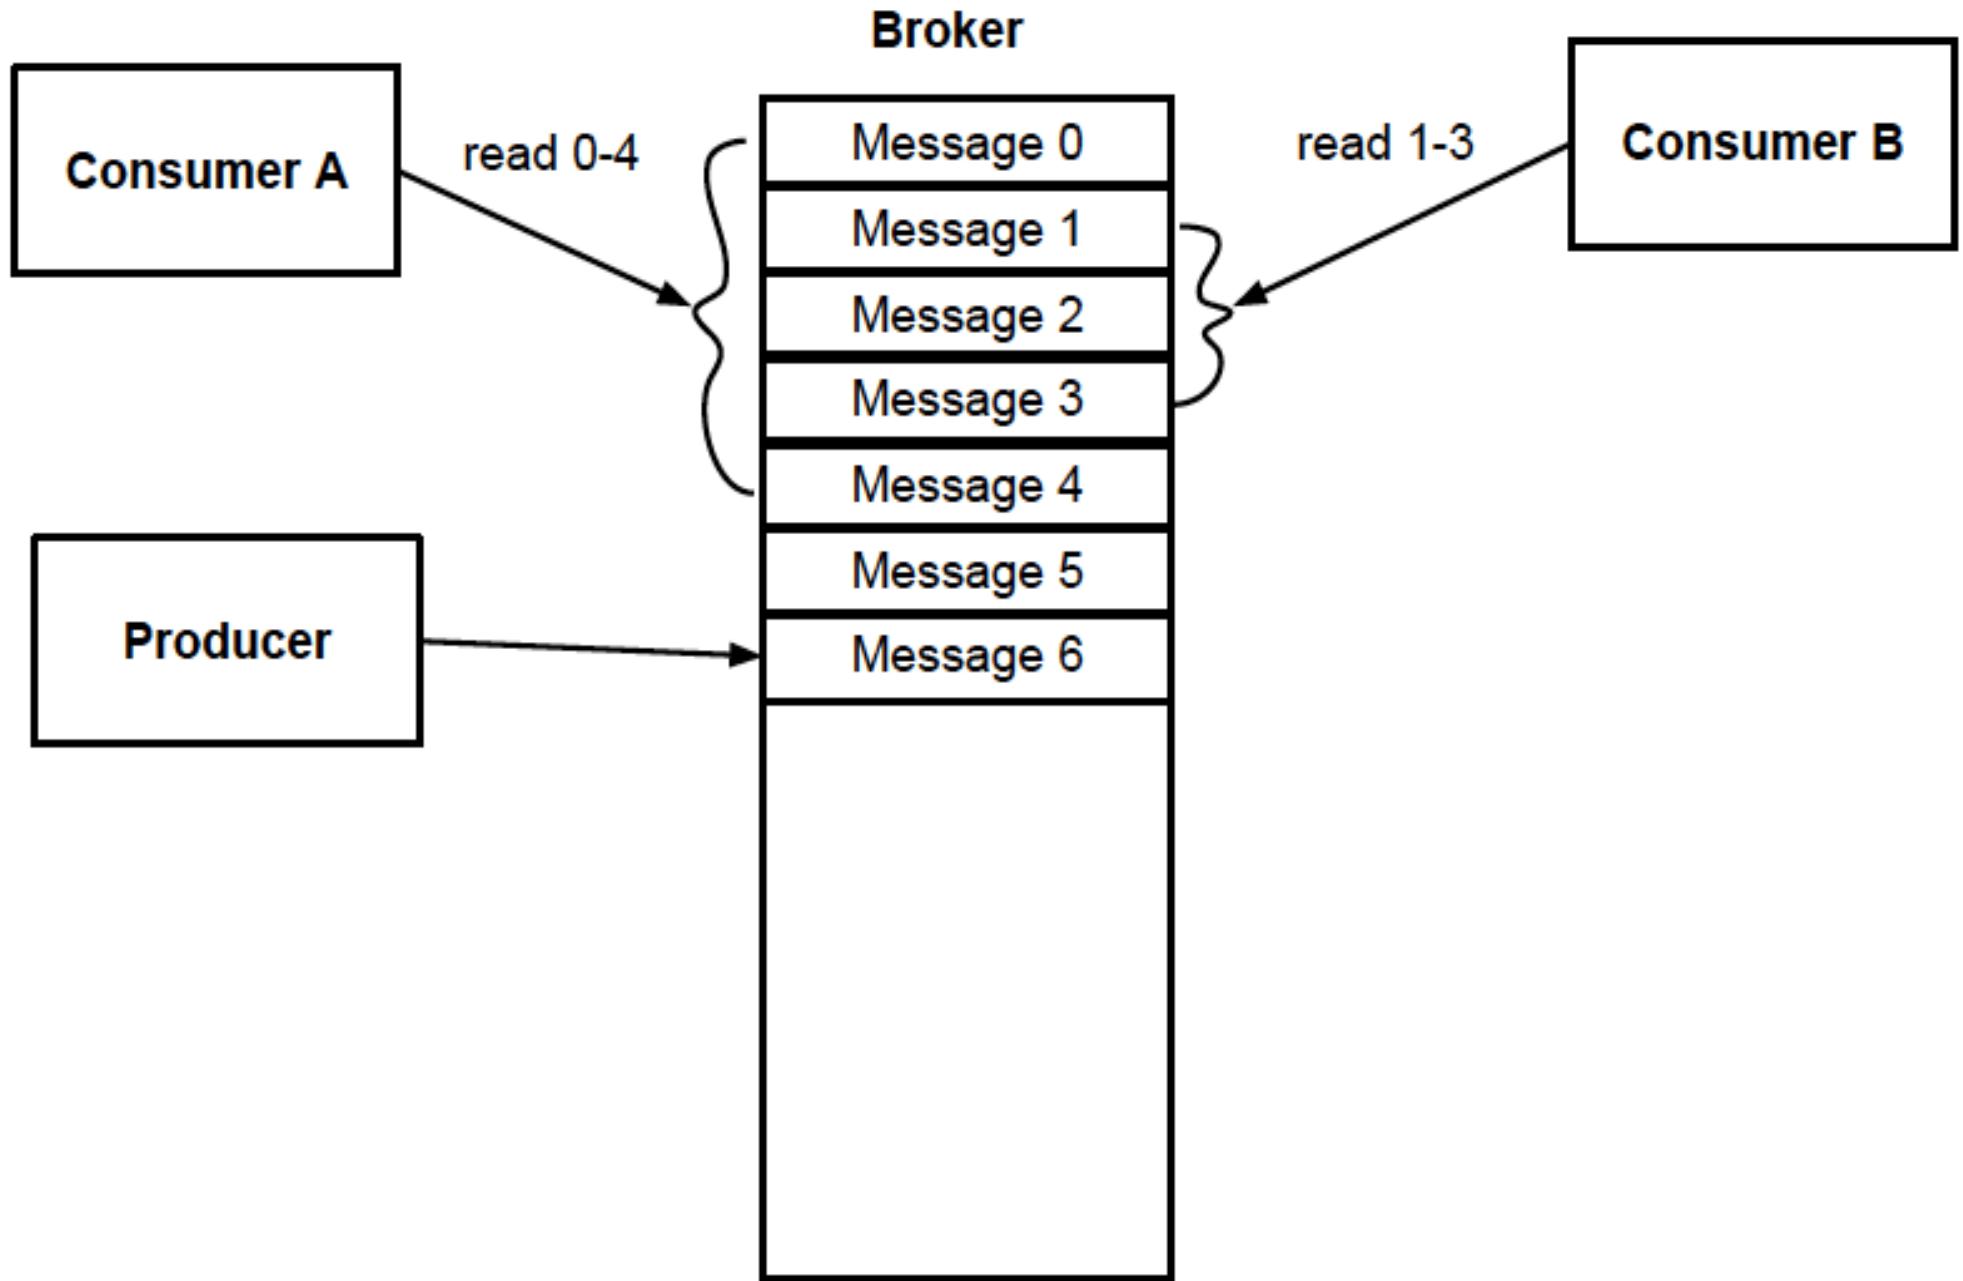
Broker



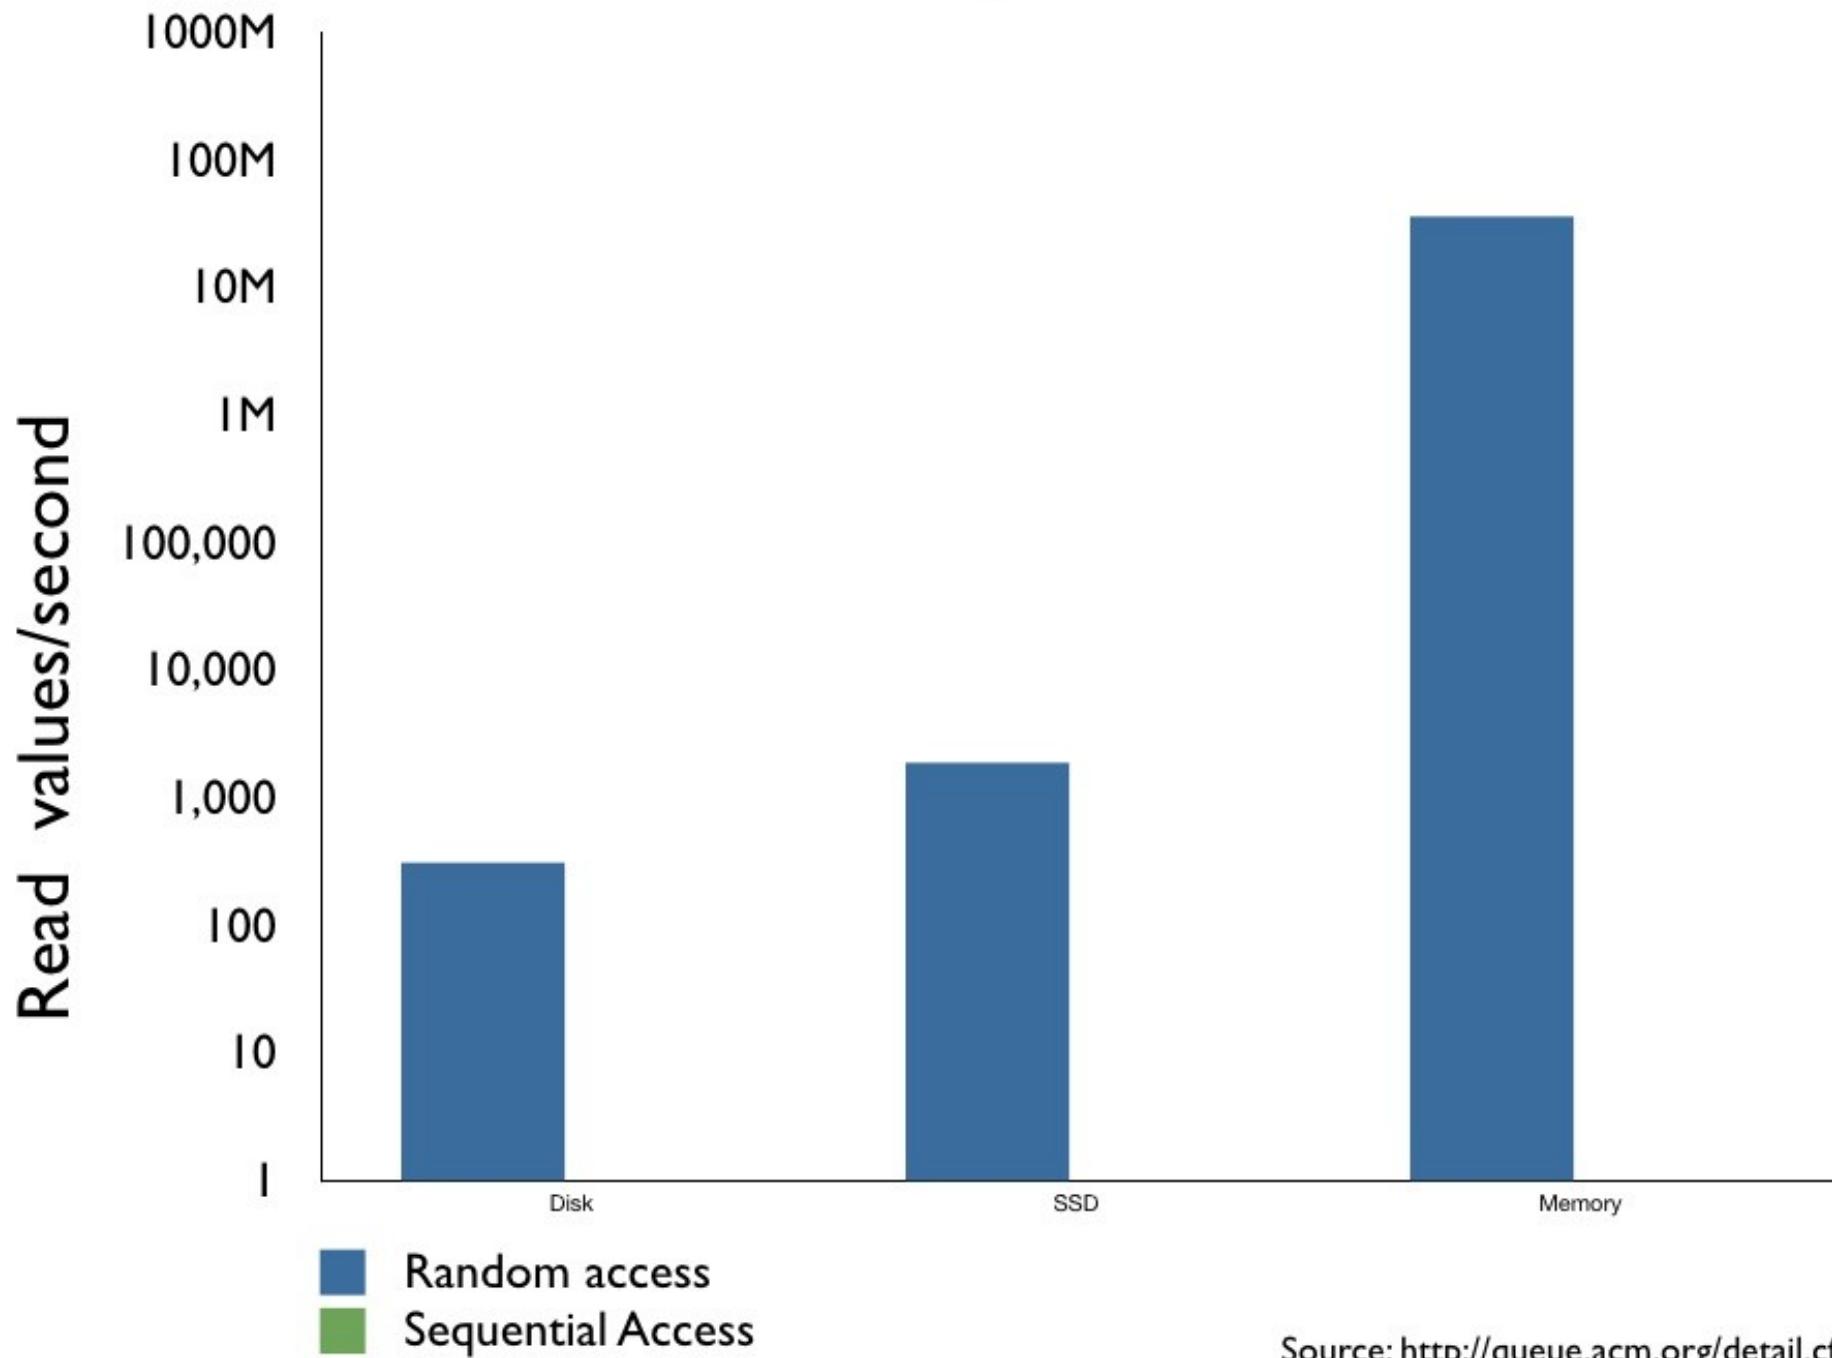




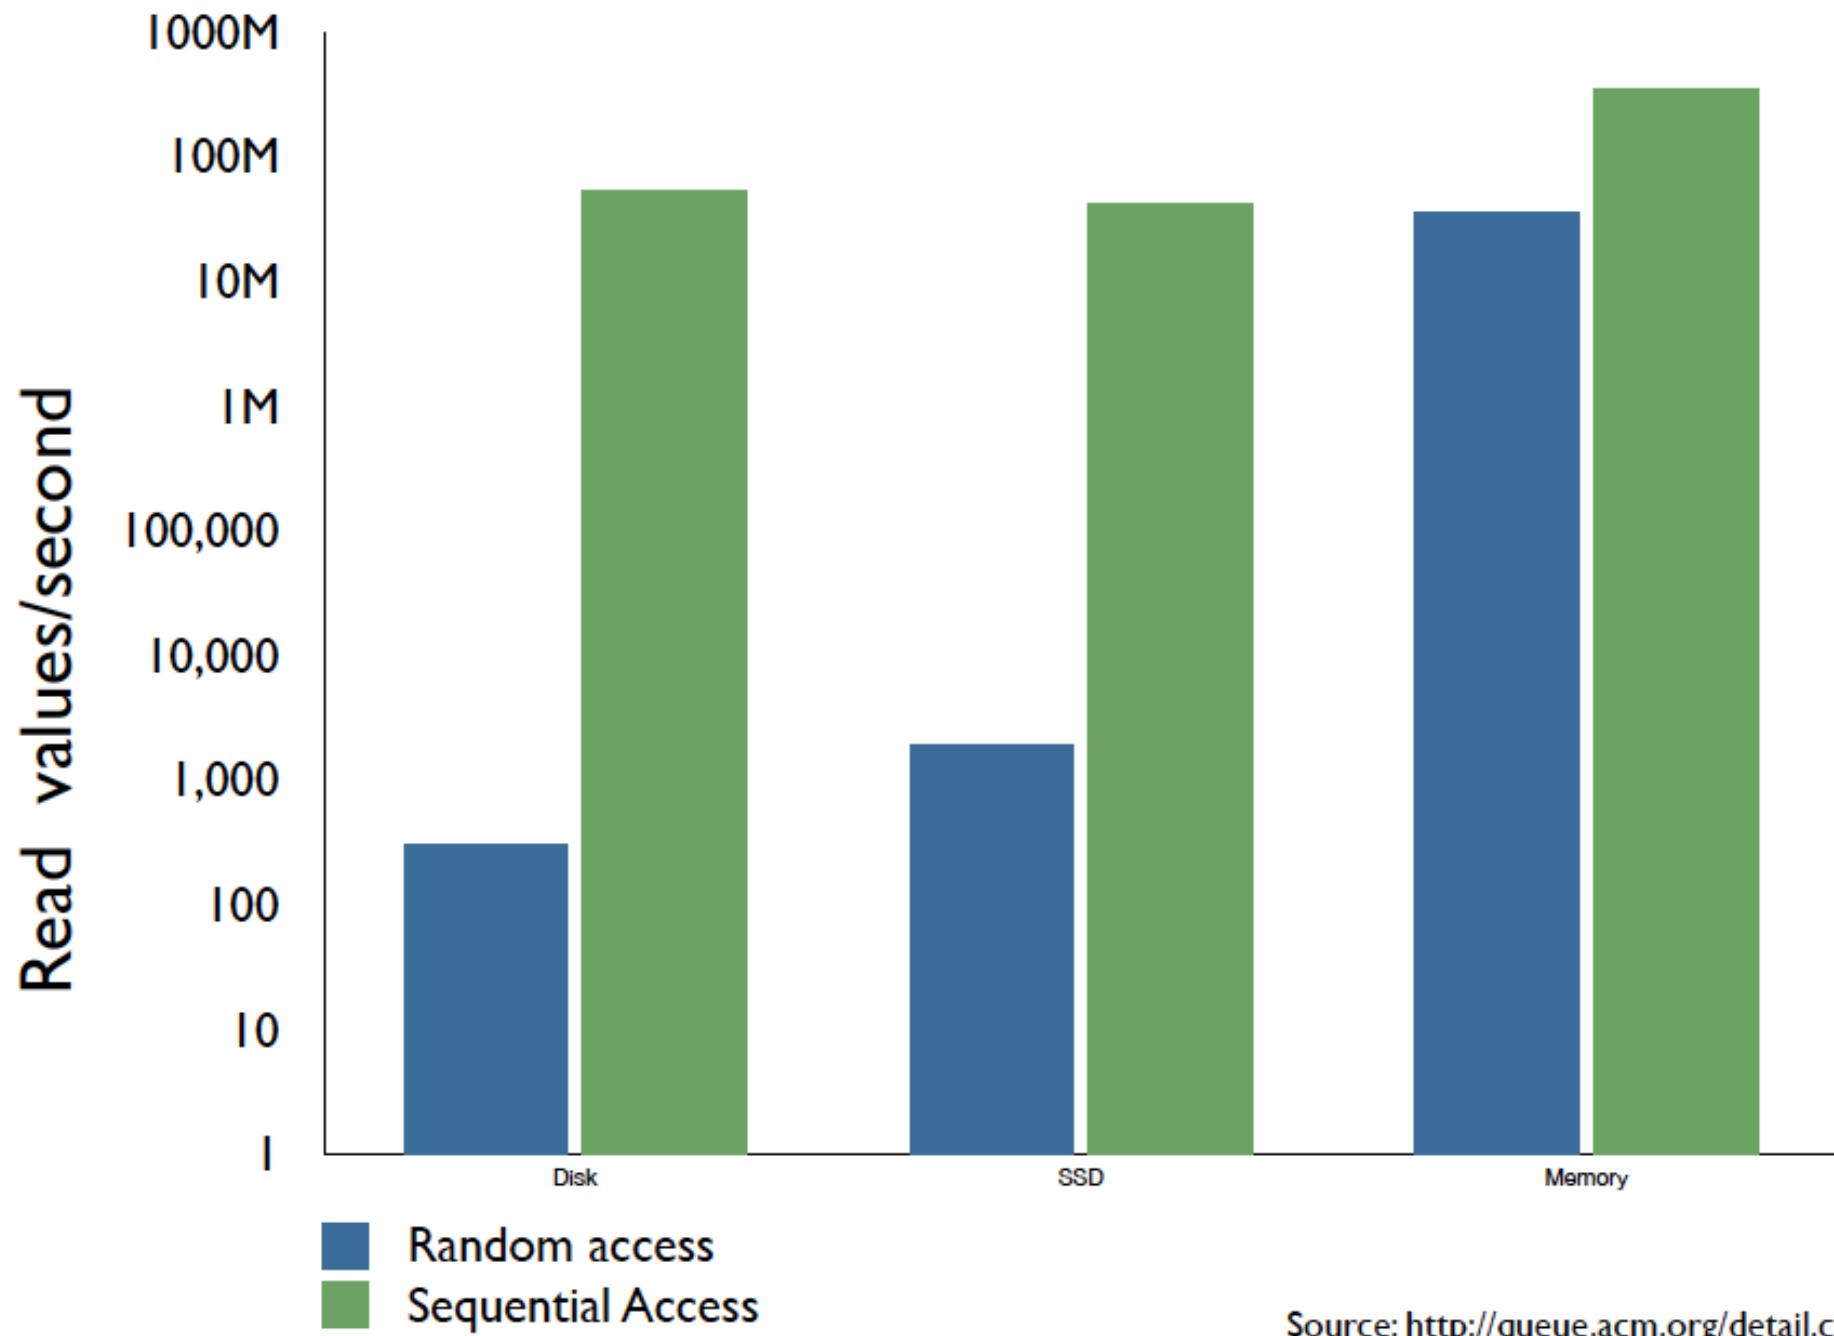
Doesn't Fear the File System



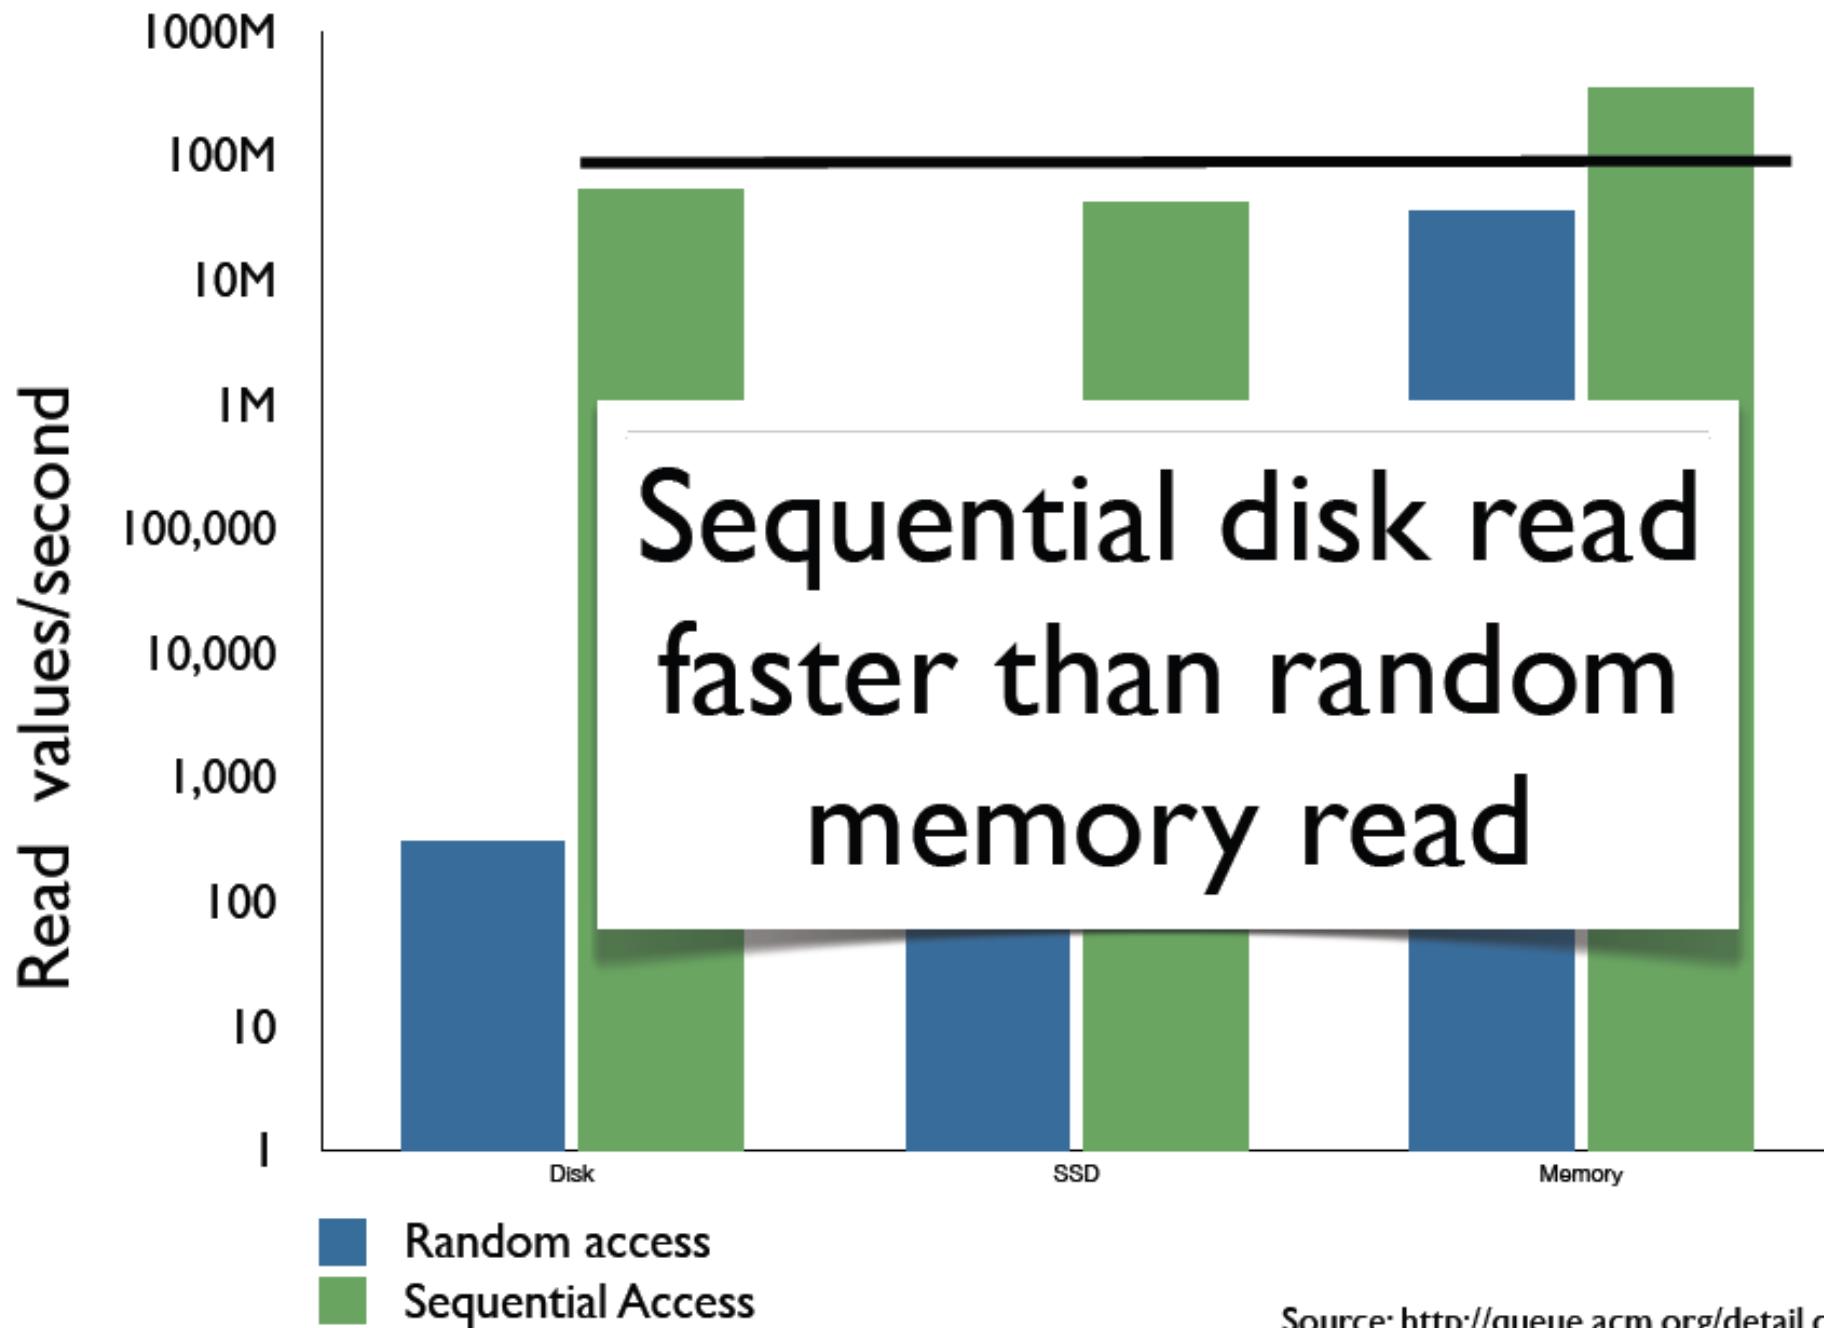
Disk/Memory Performance



Disk/Memory Performance

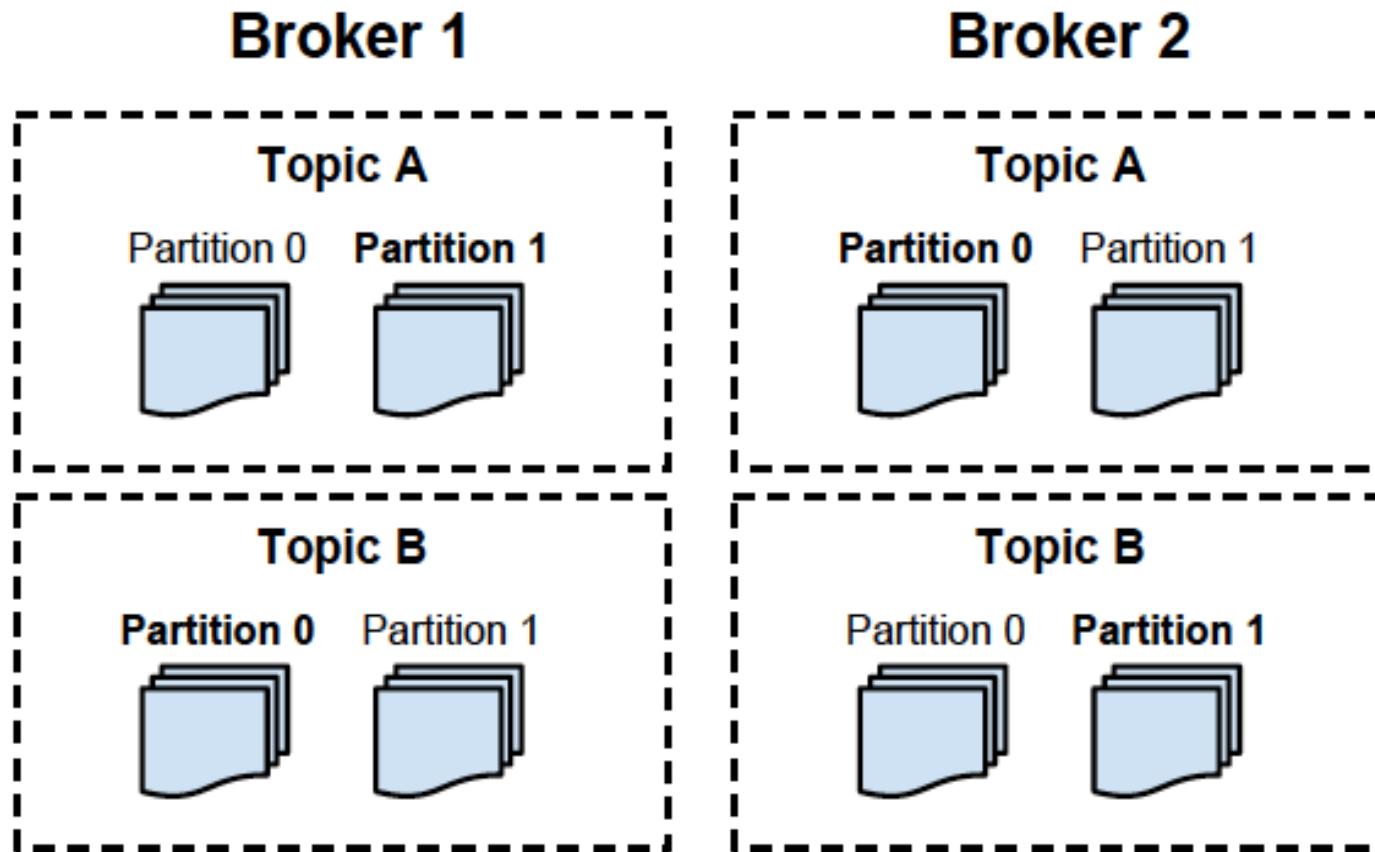


Disk/Memory Performance

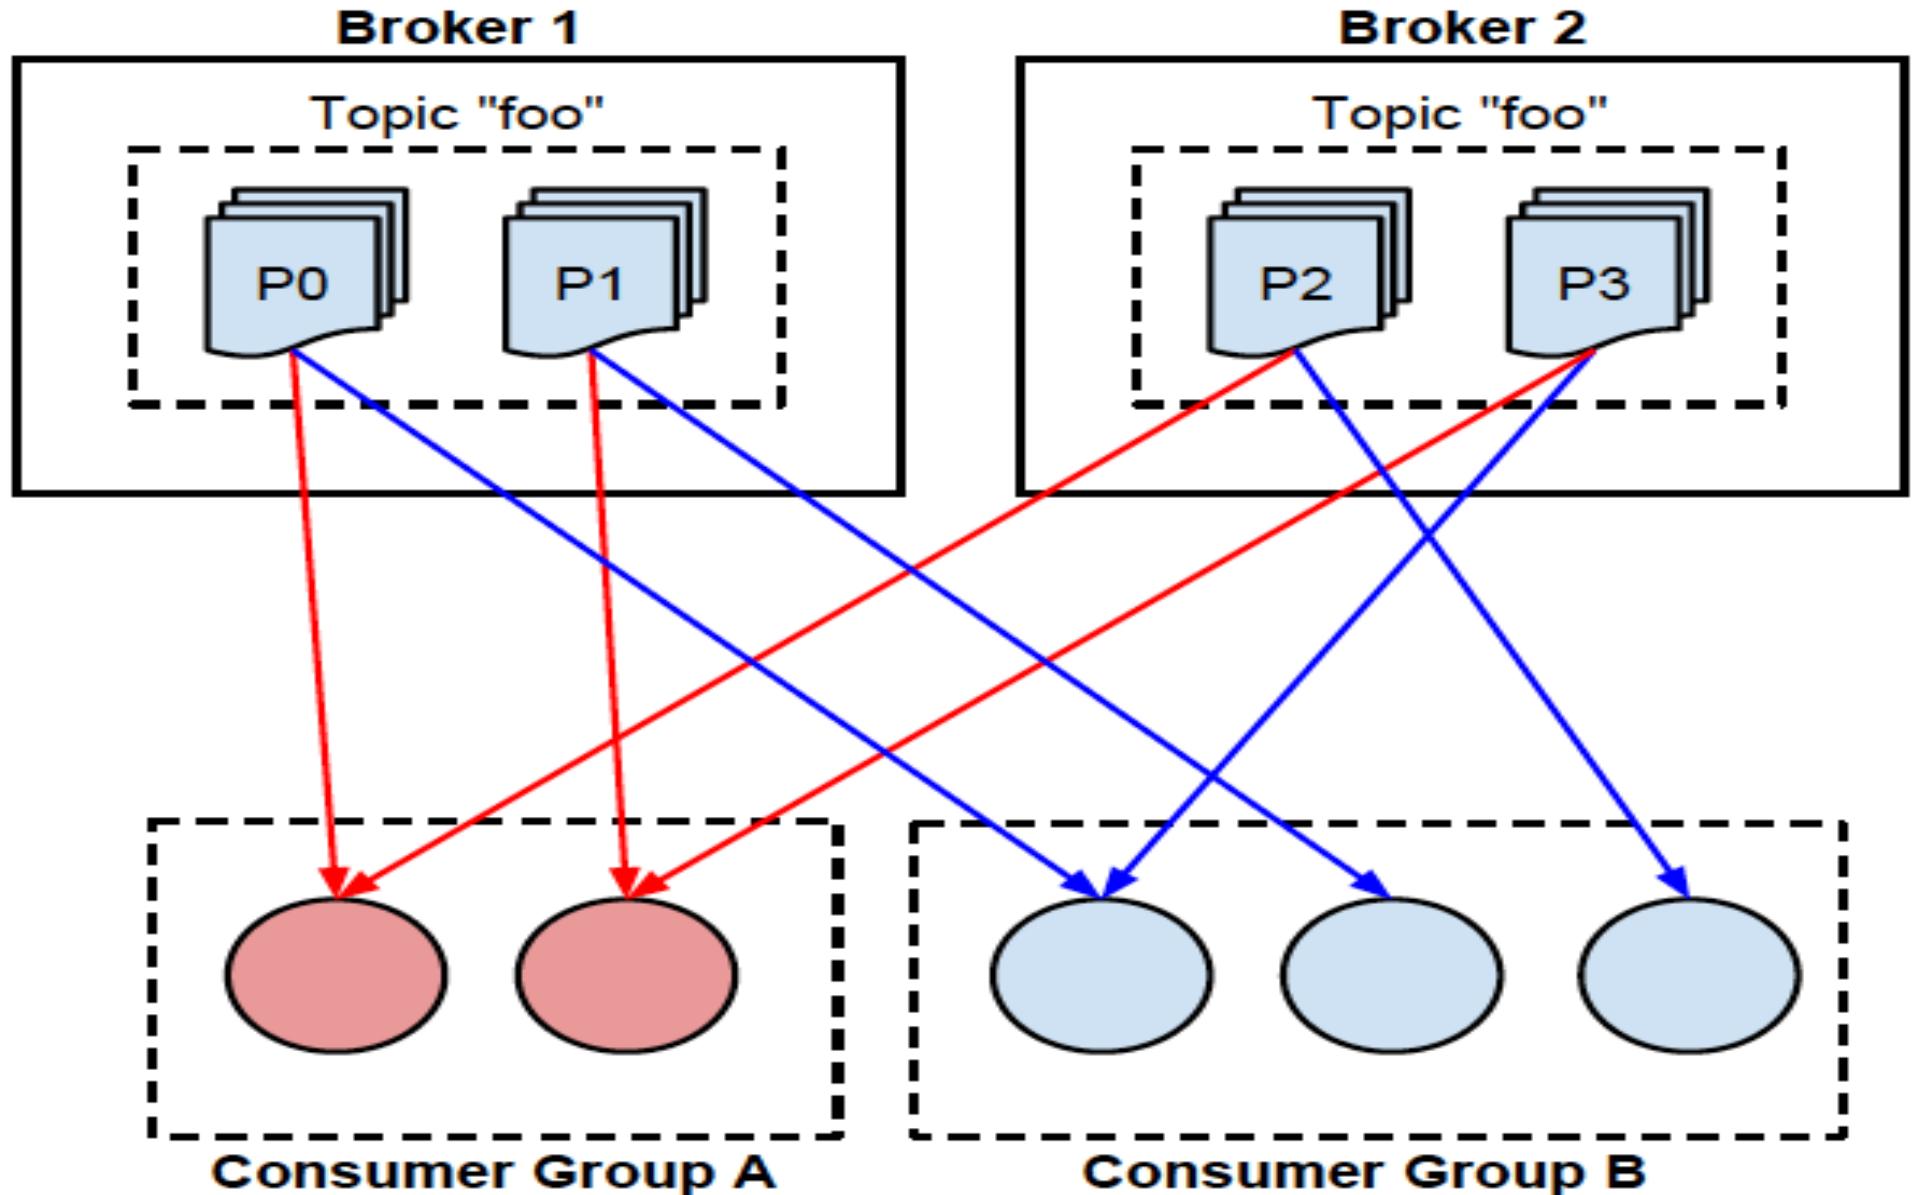


Topics

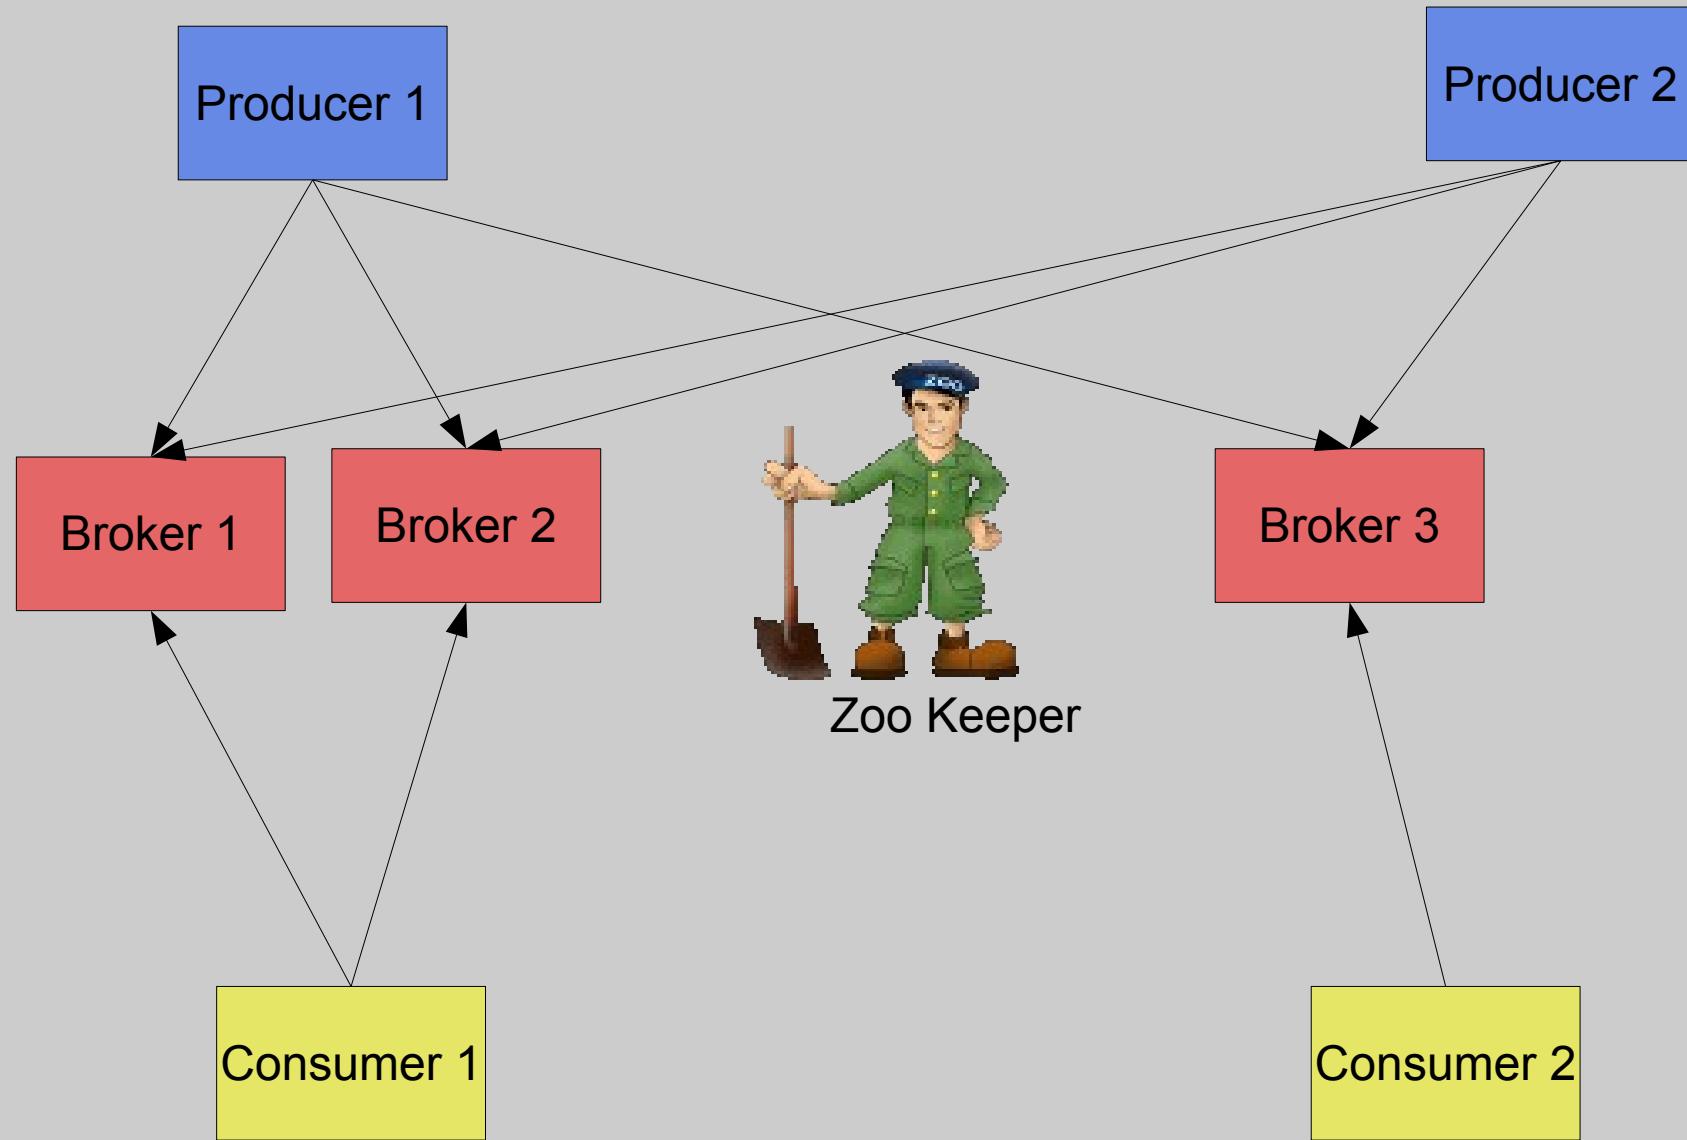
- Logical collections of partitions (the physical files).
- A broker contains some of the partitions for a topic

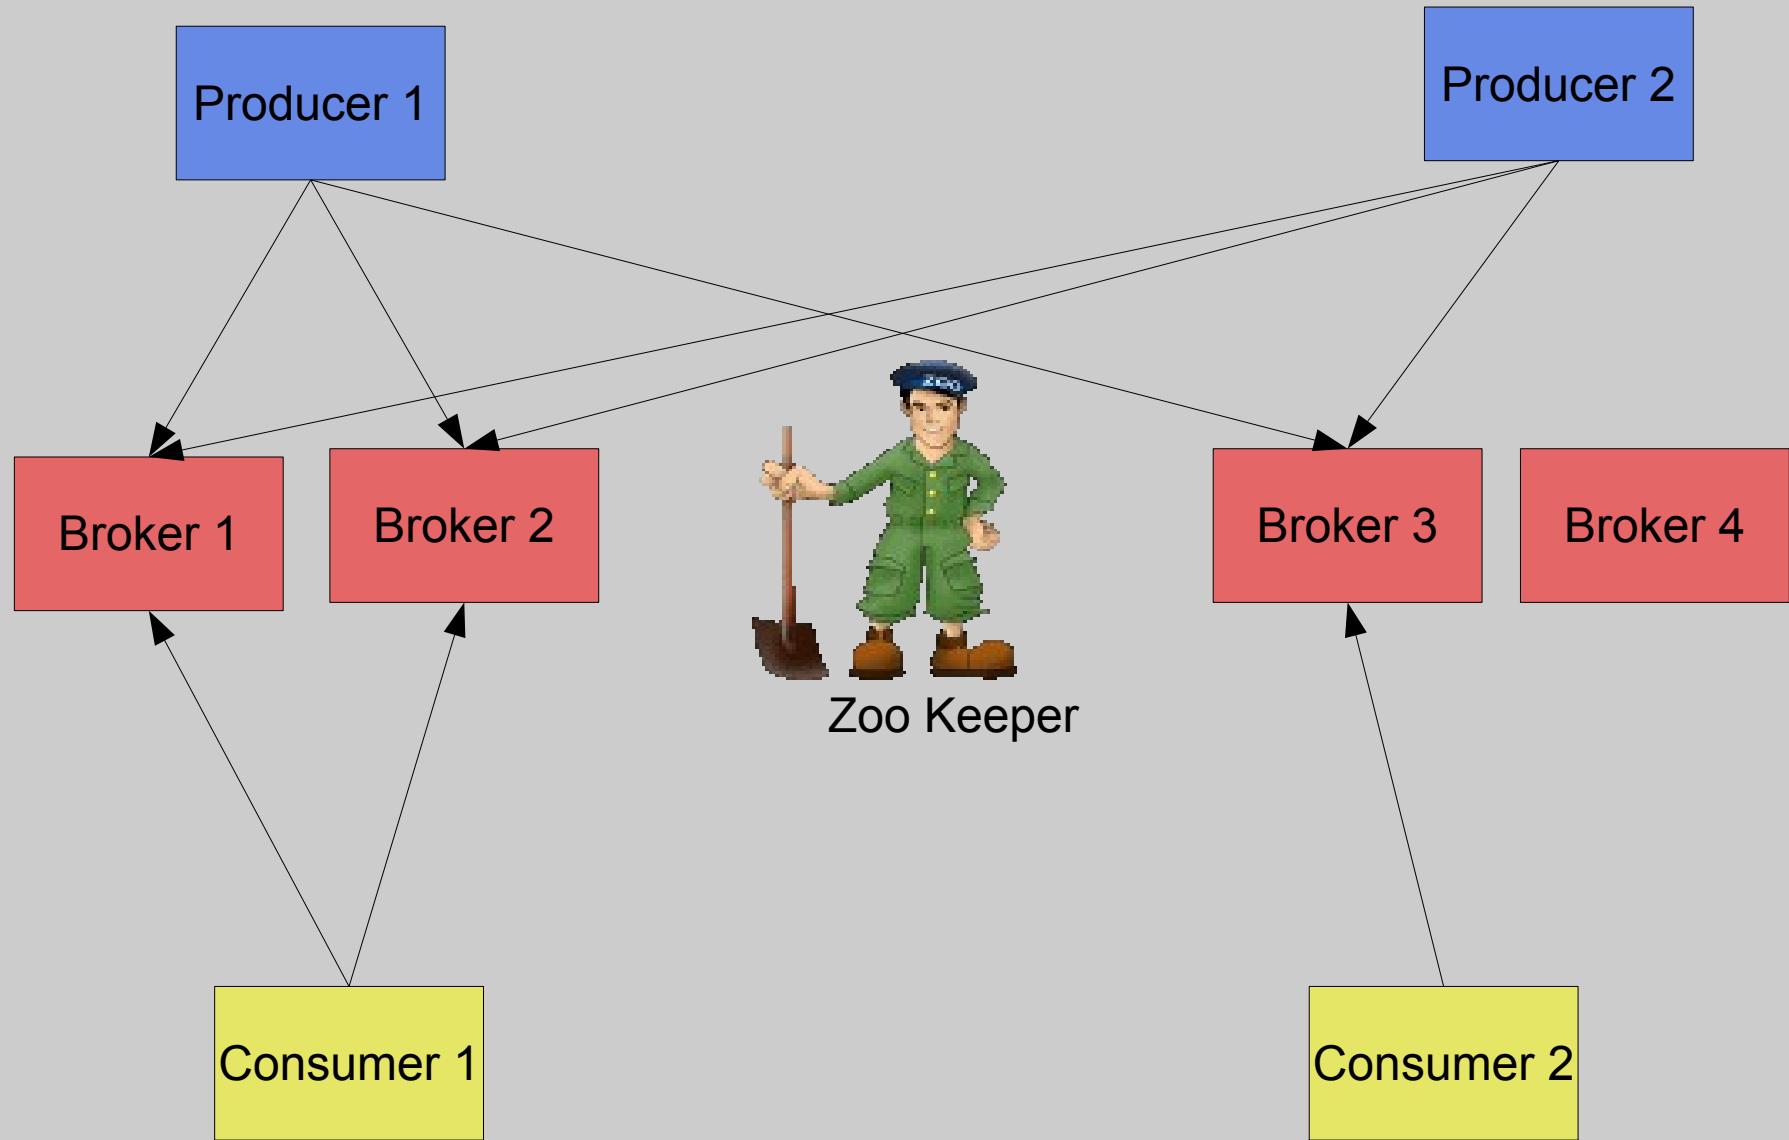


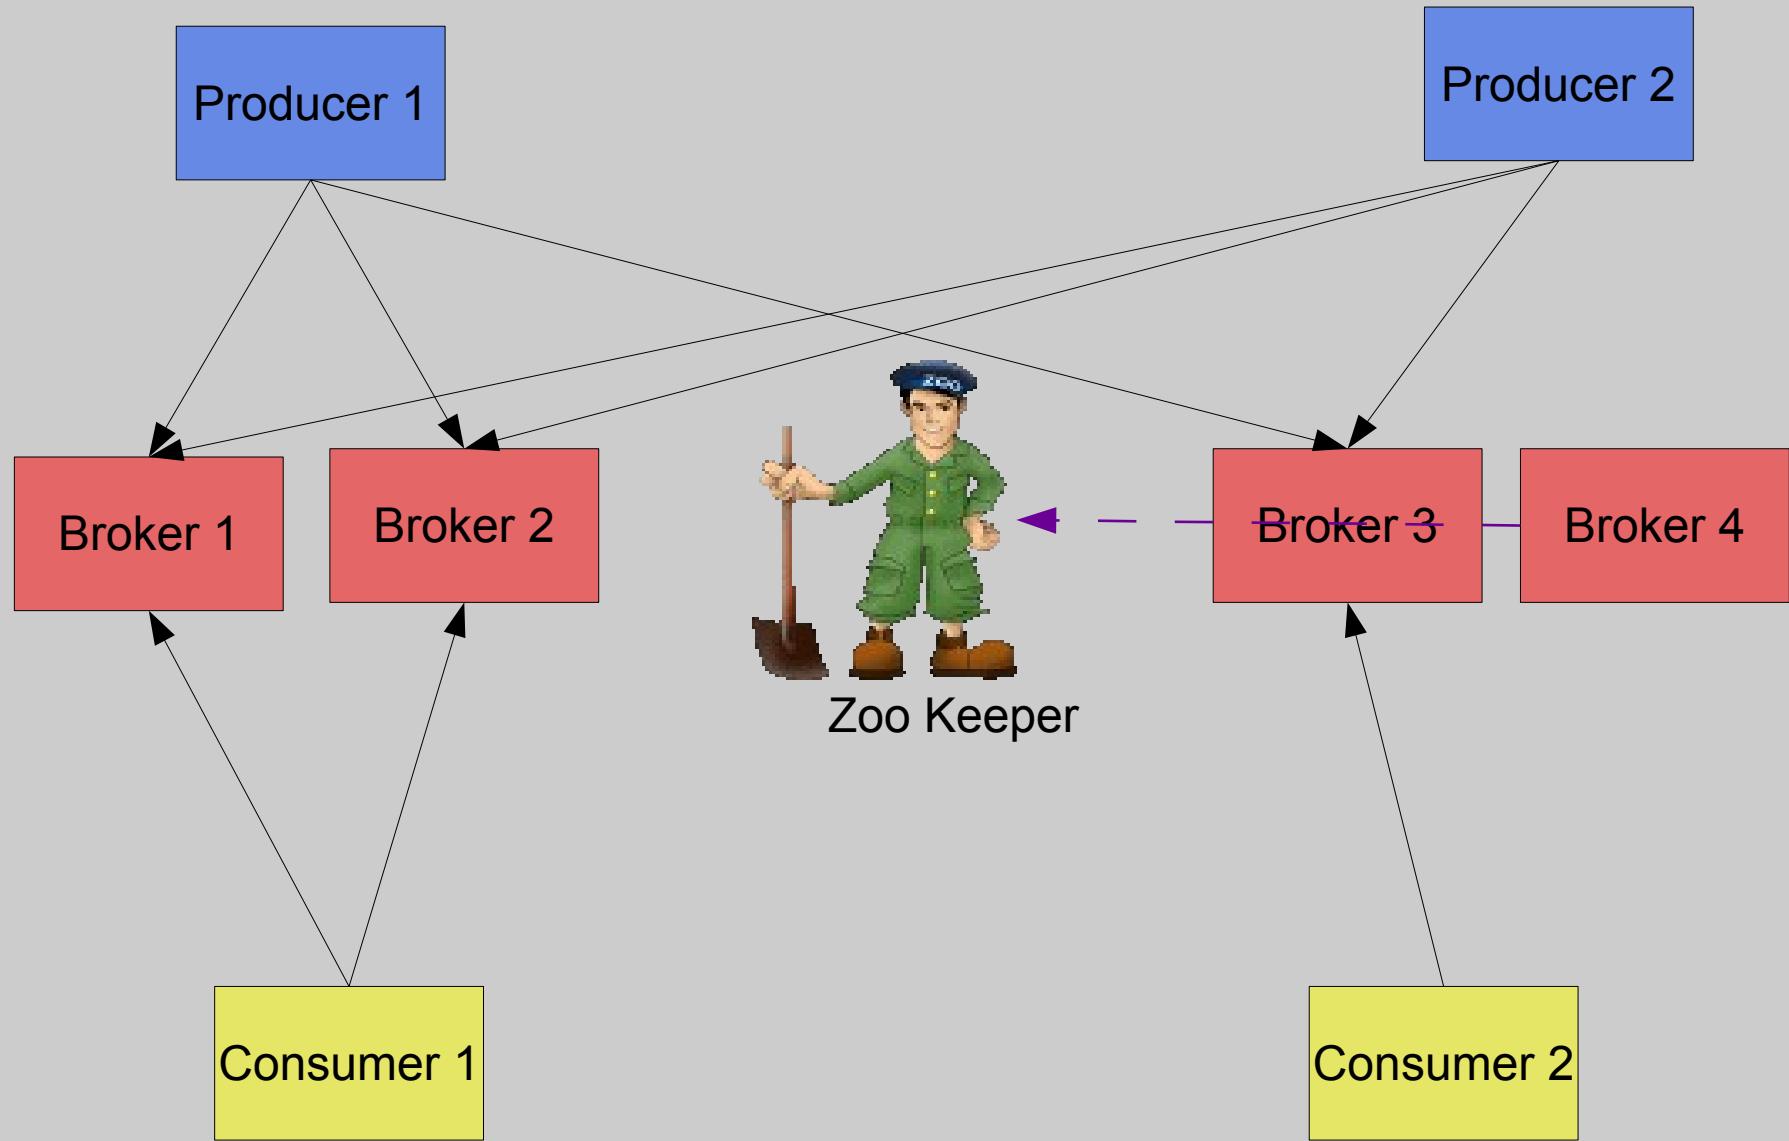
A partition is Consumed by Exactly One Group's Consumer

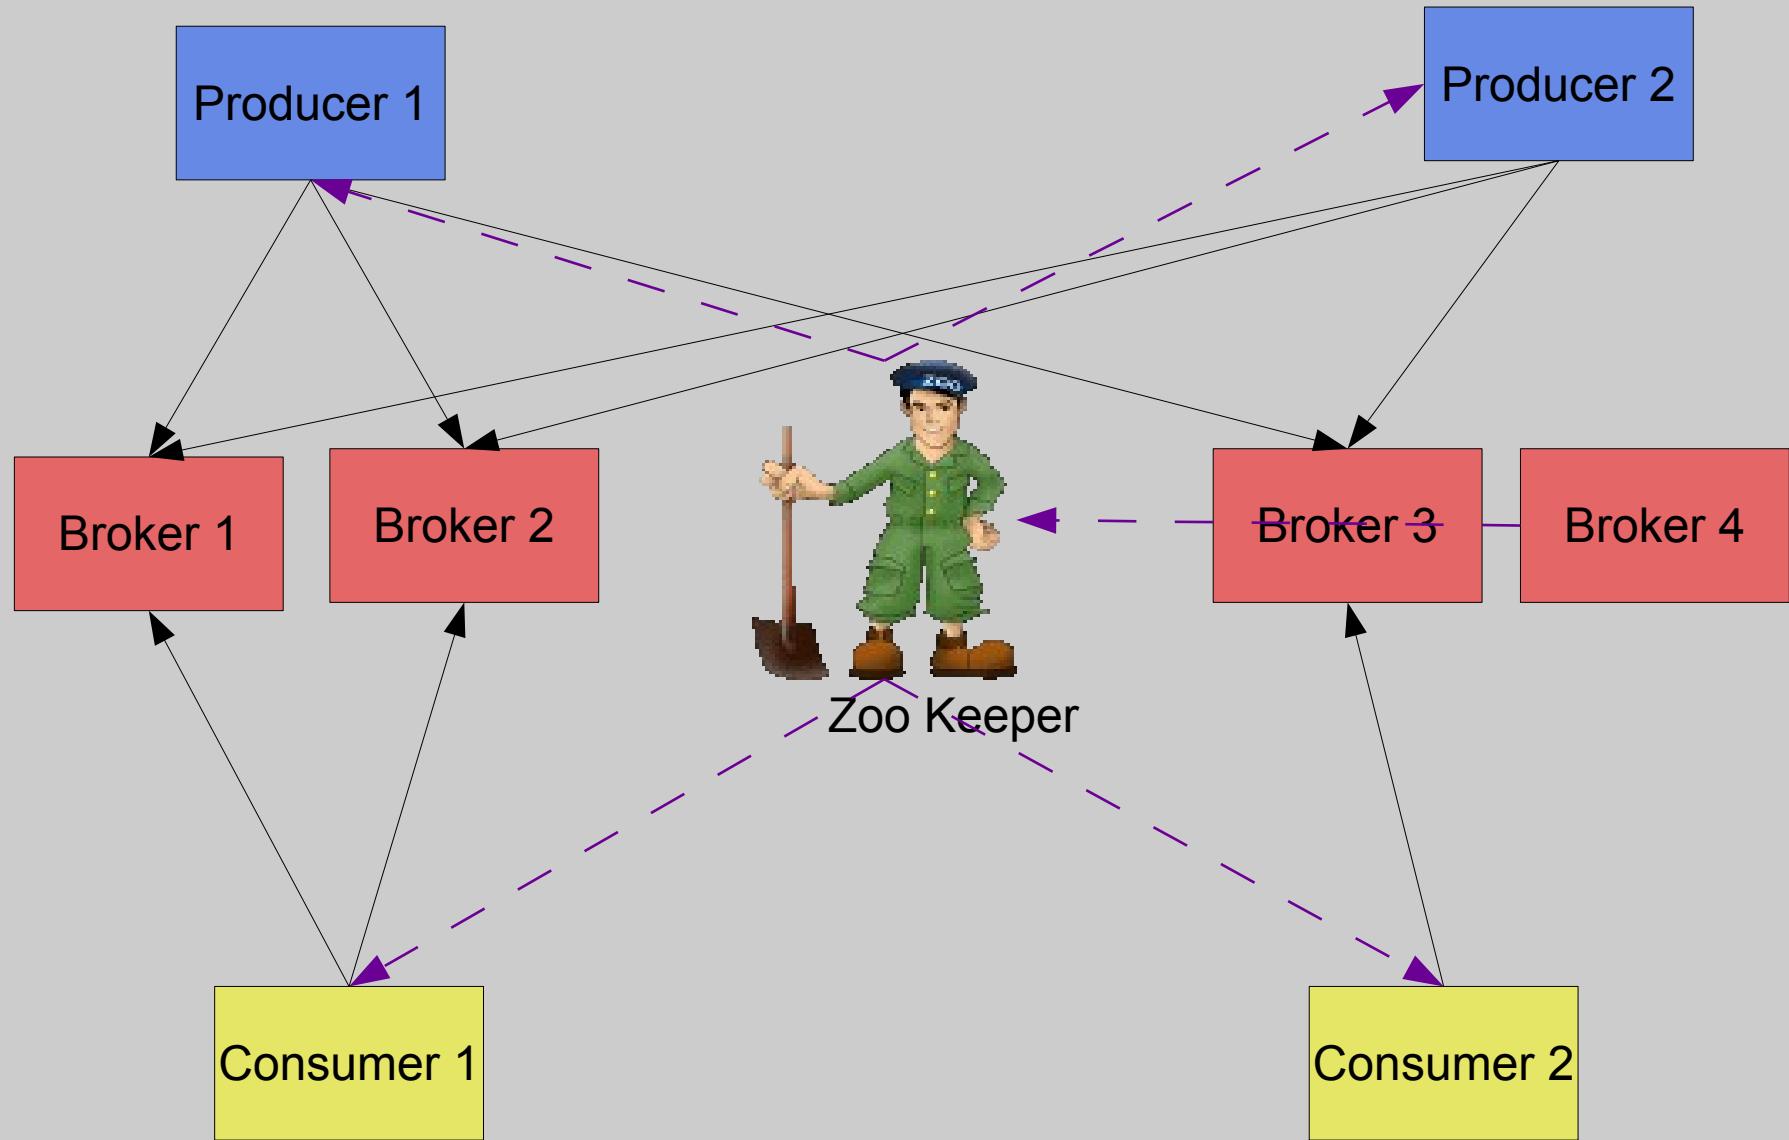


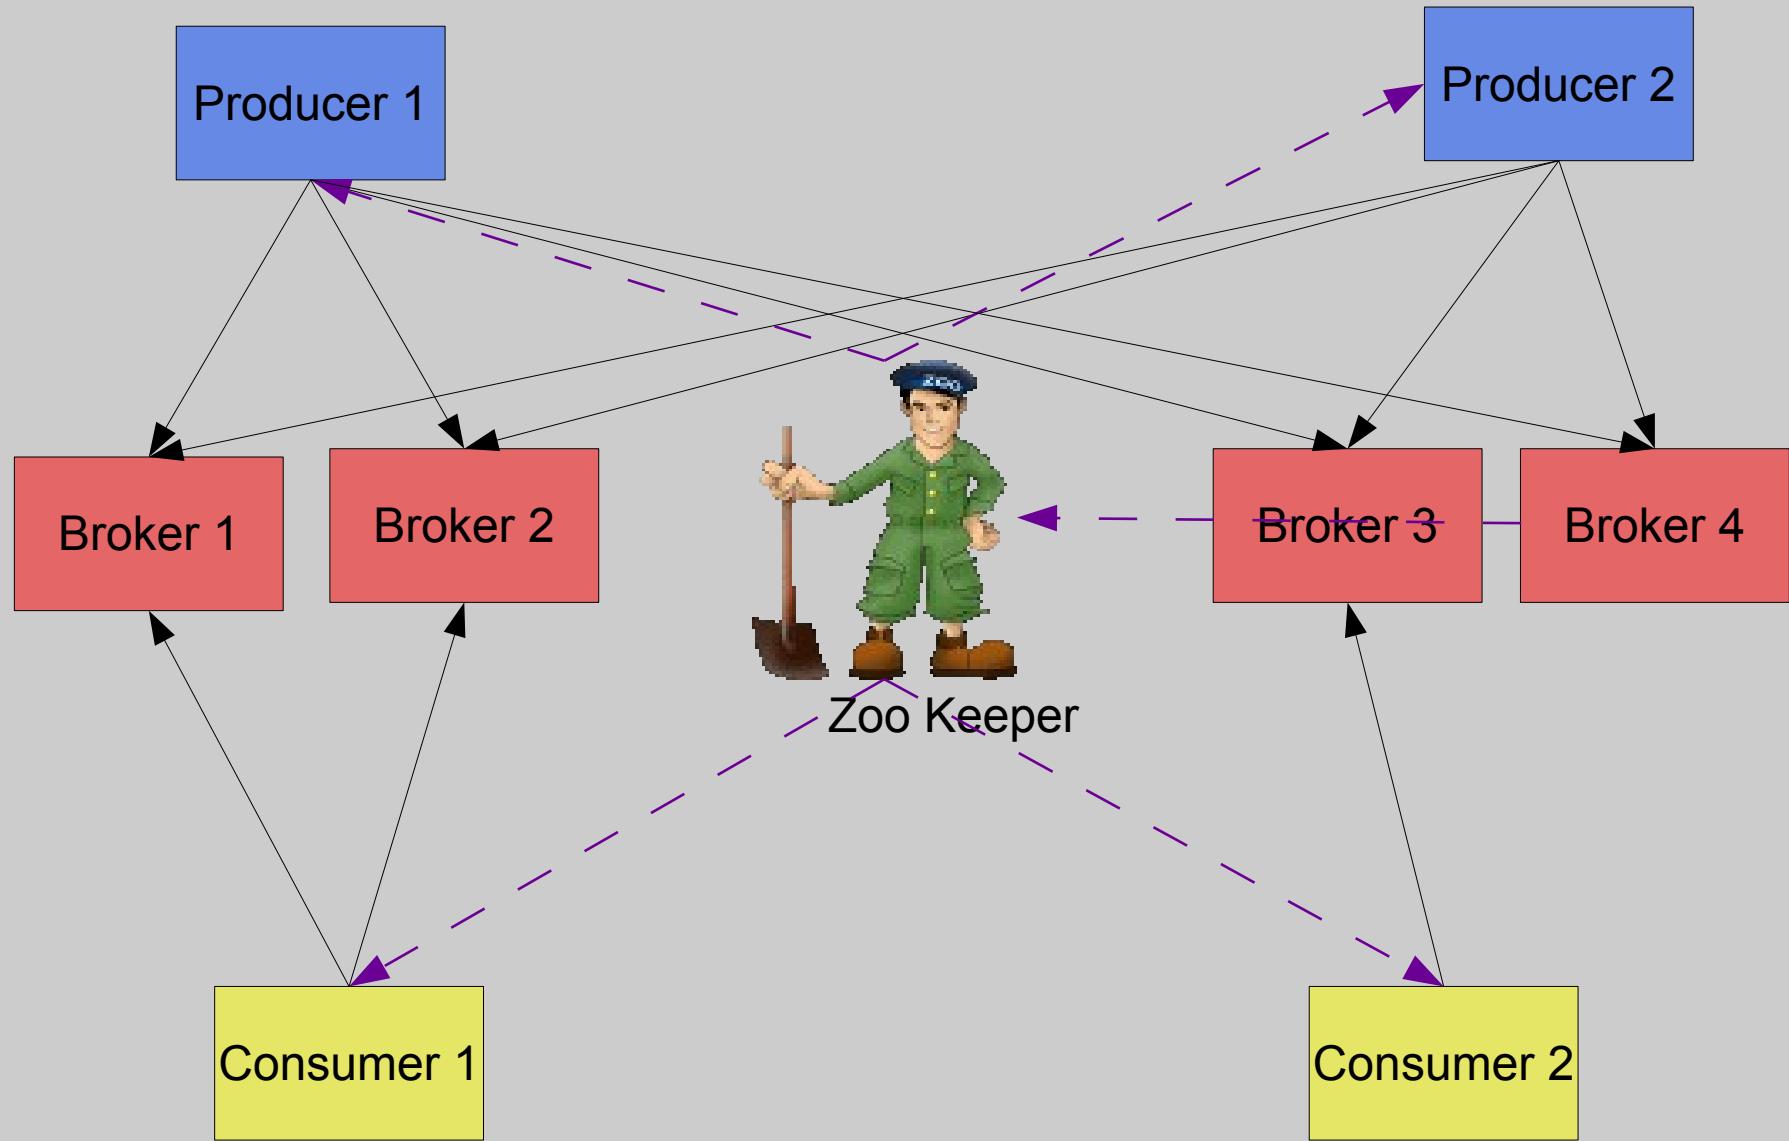
Distributed & Fault-Tolerant

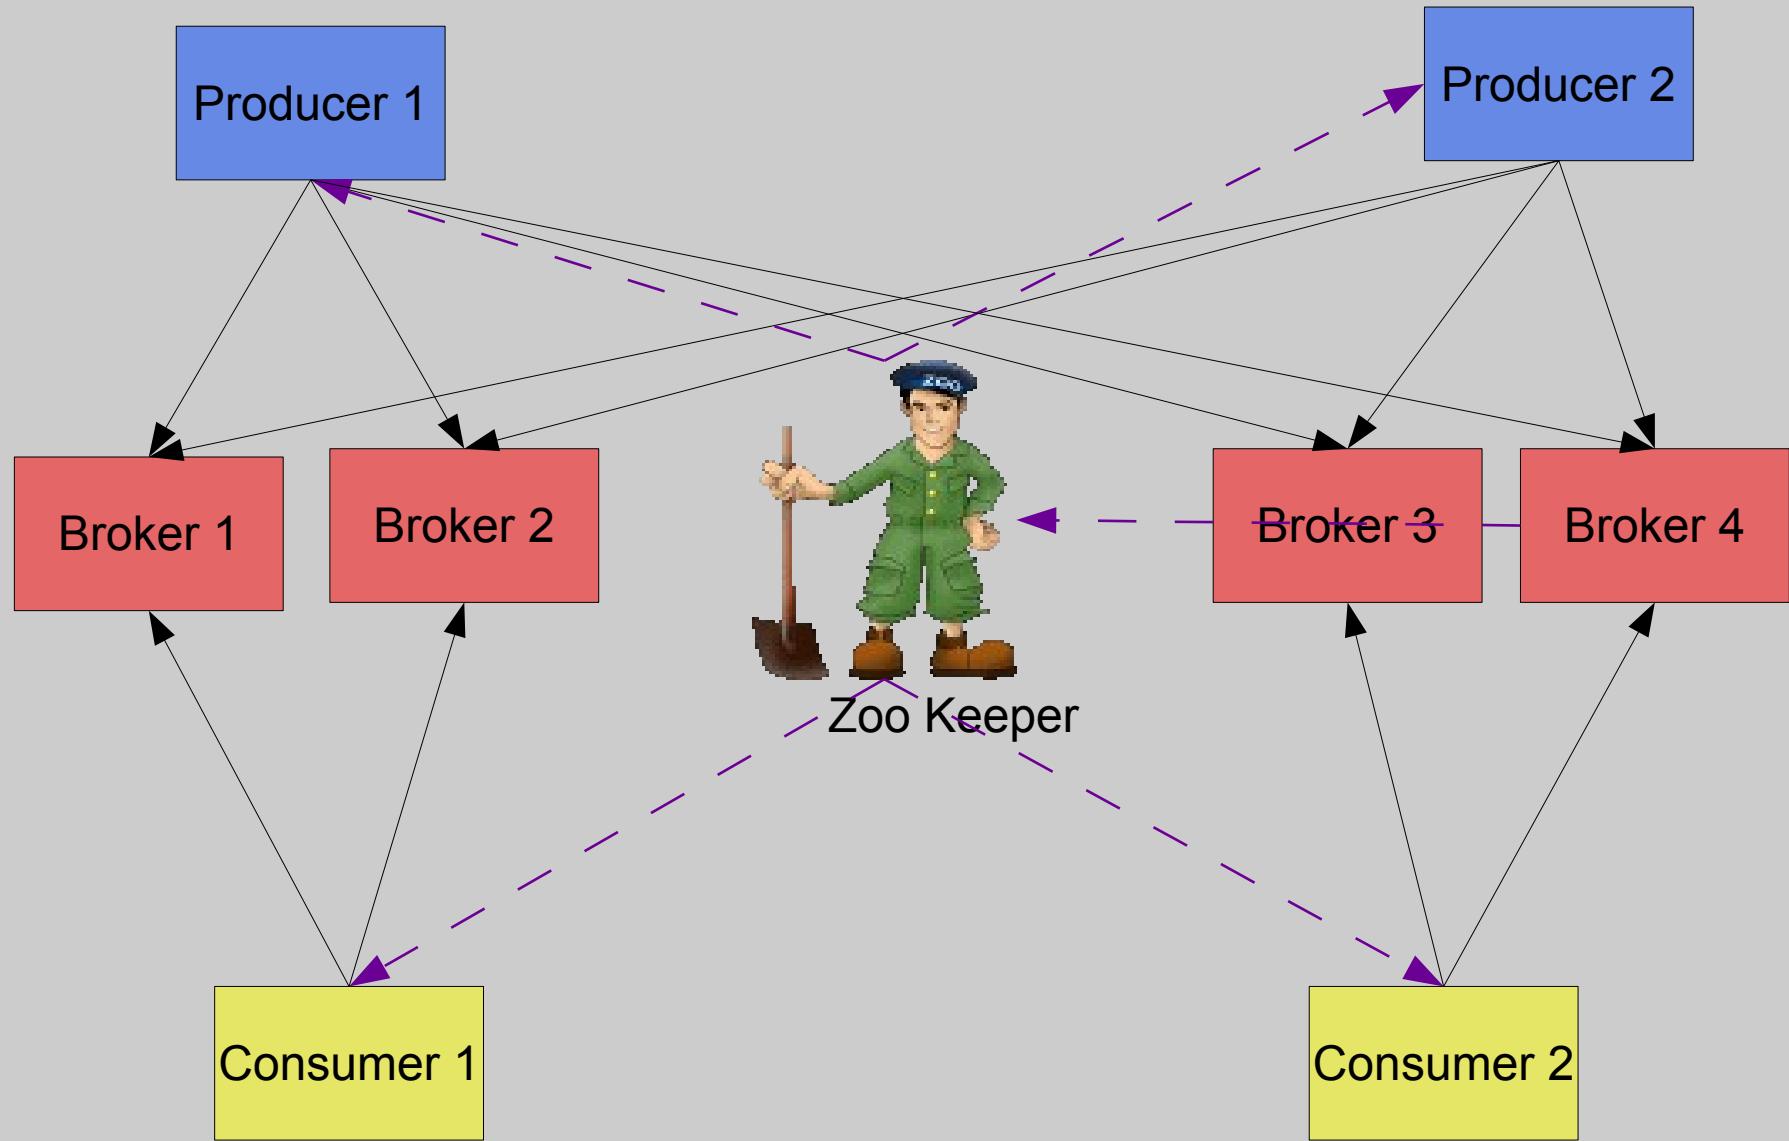


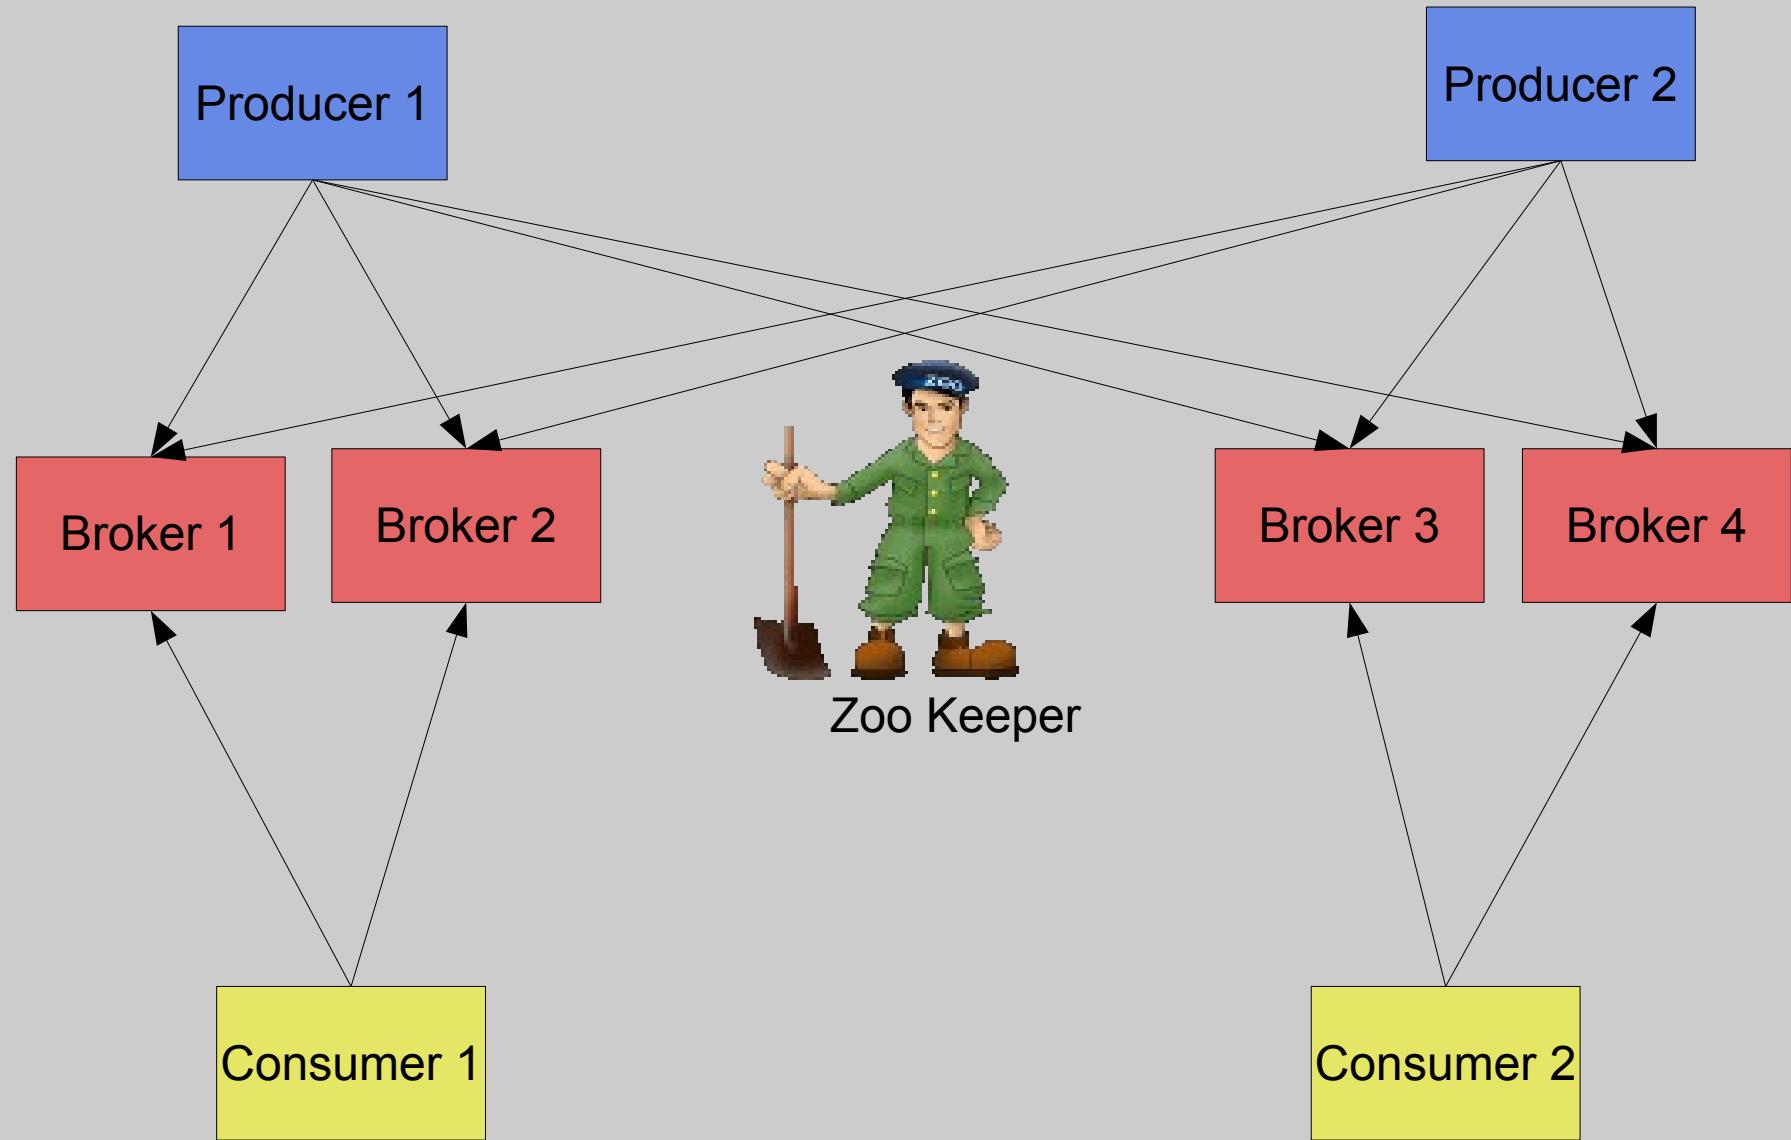


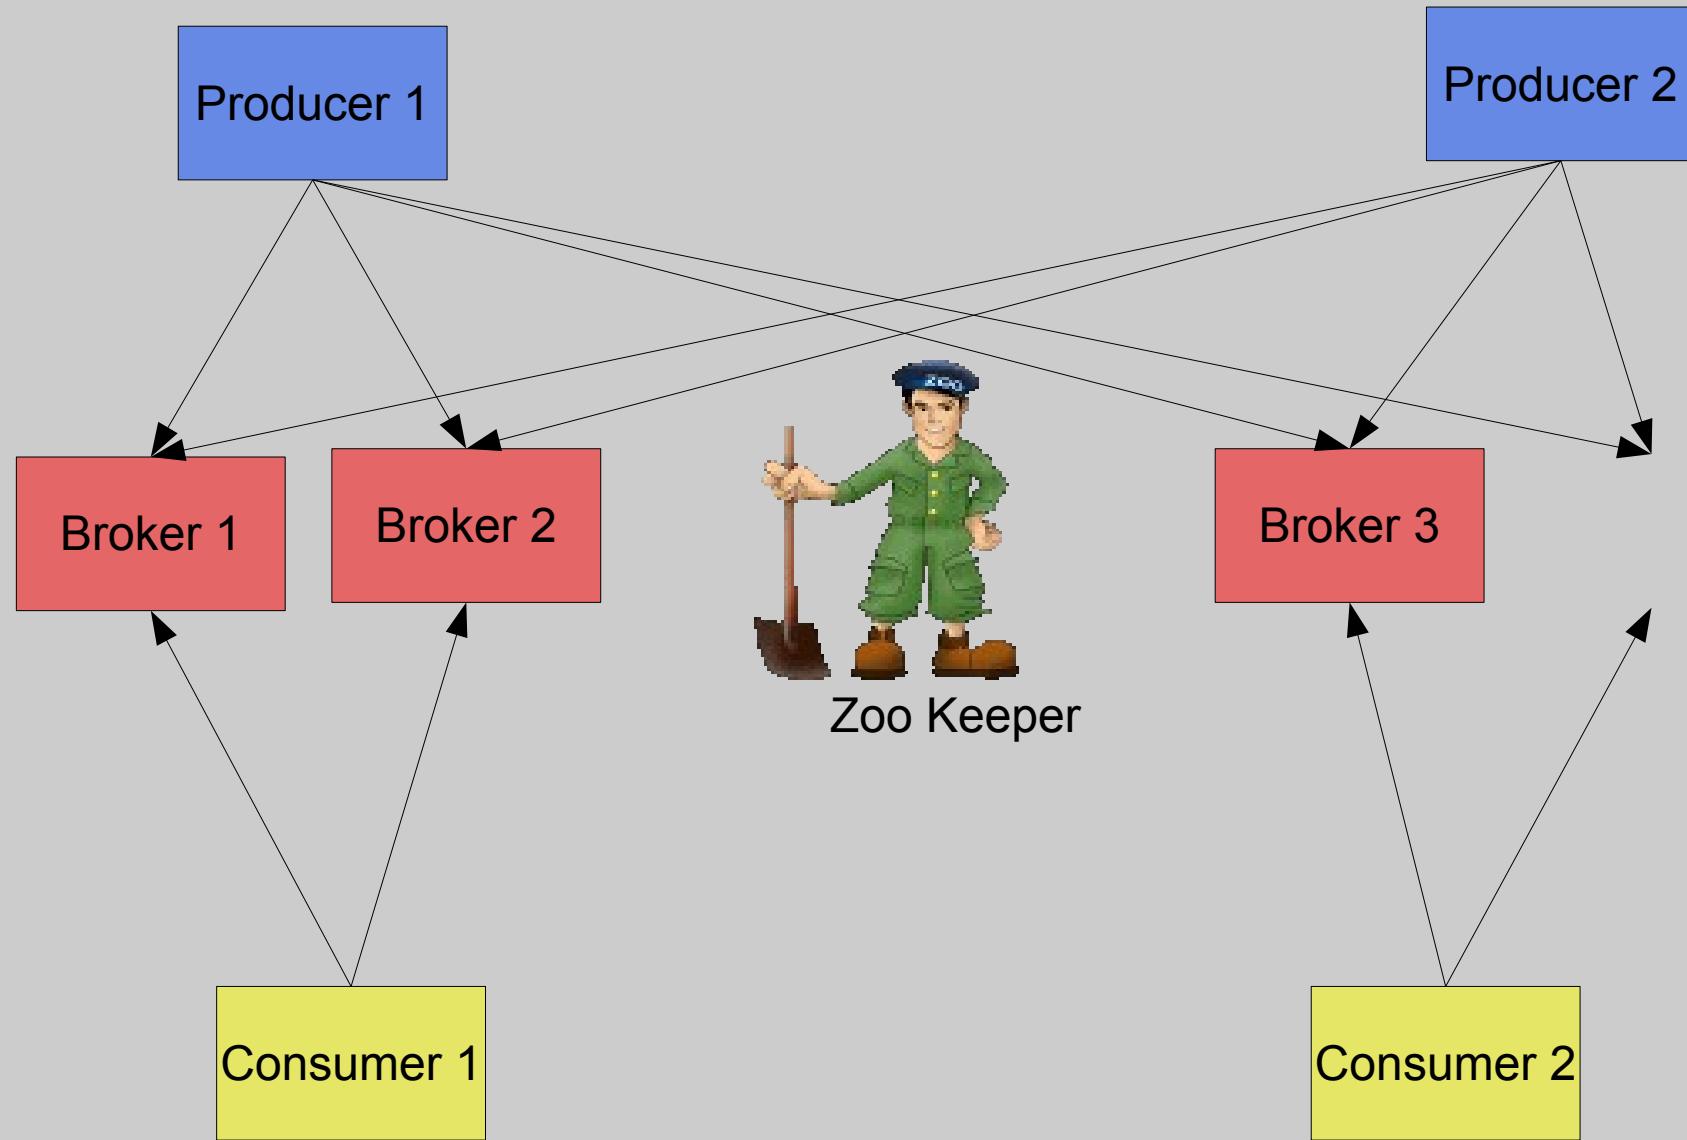


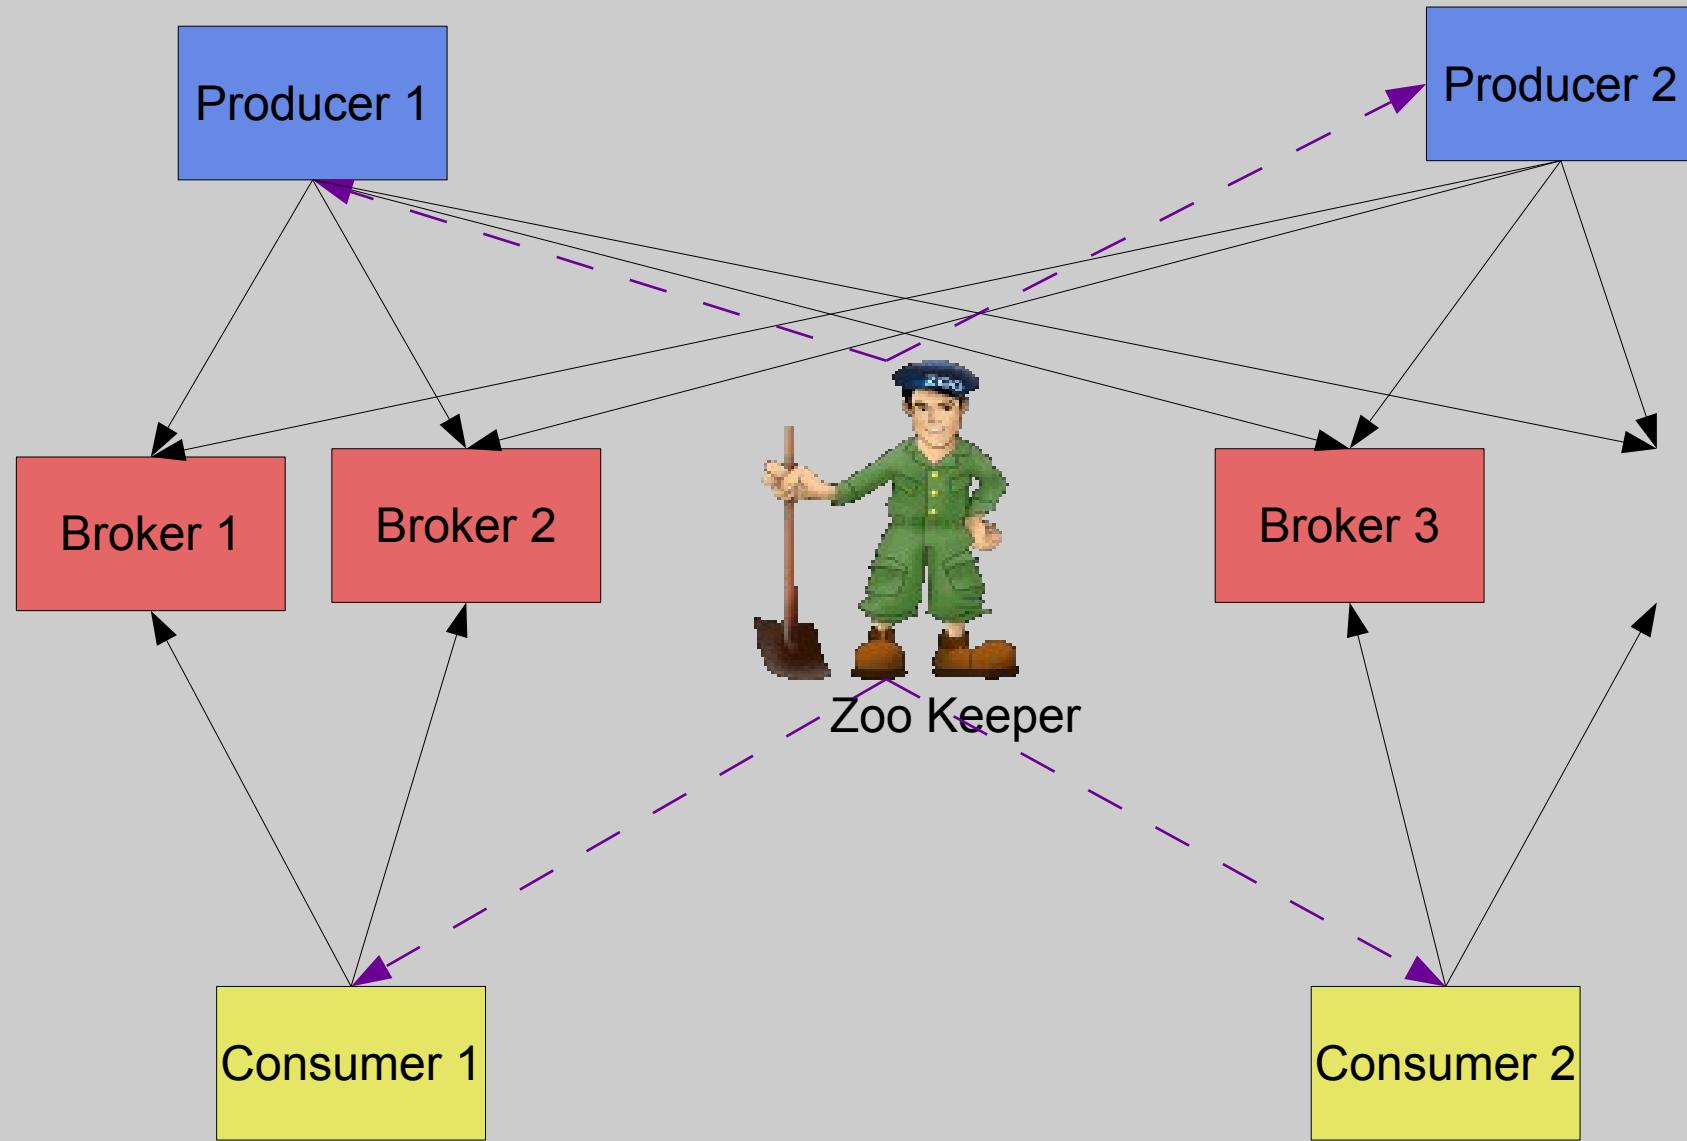


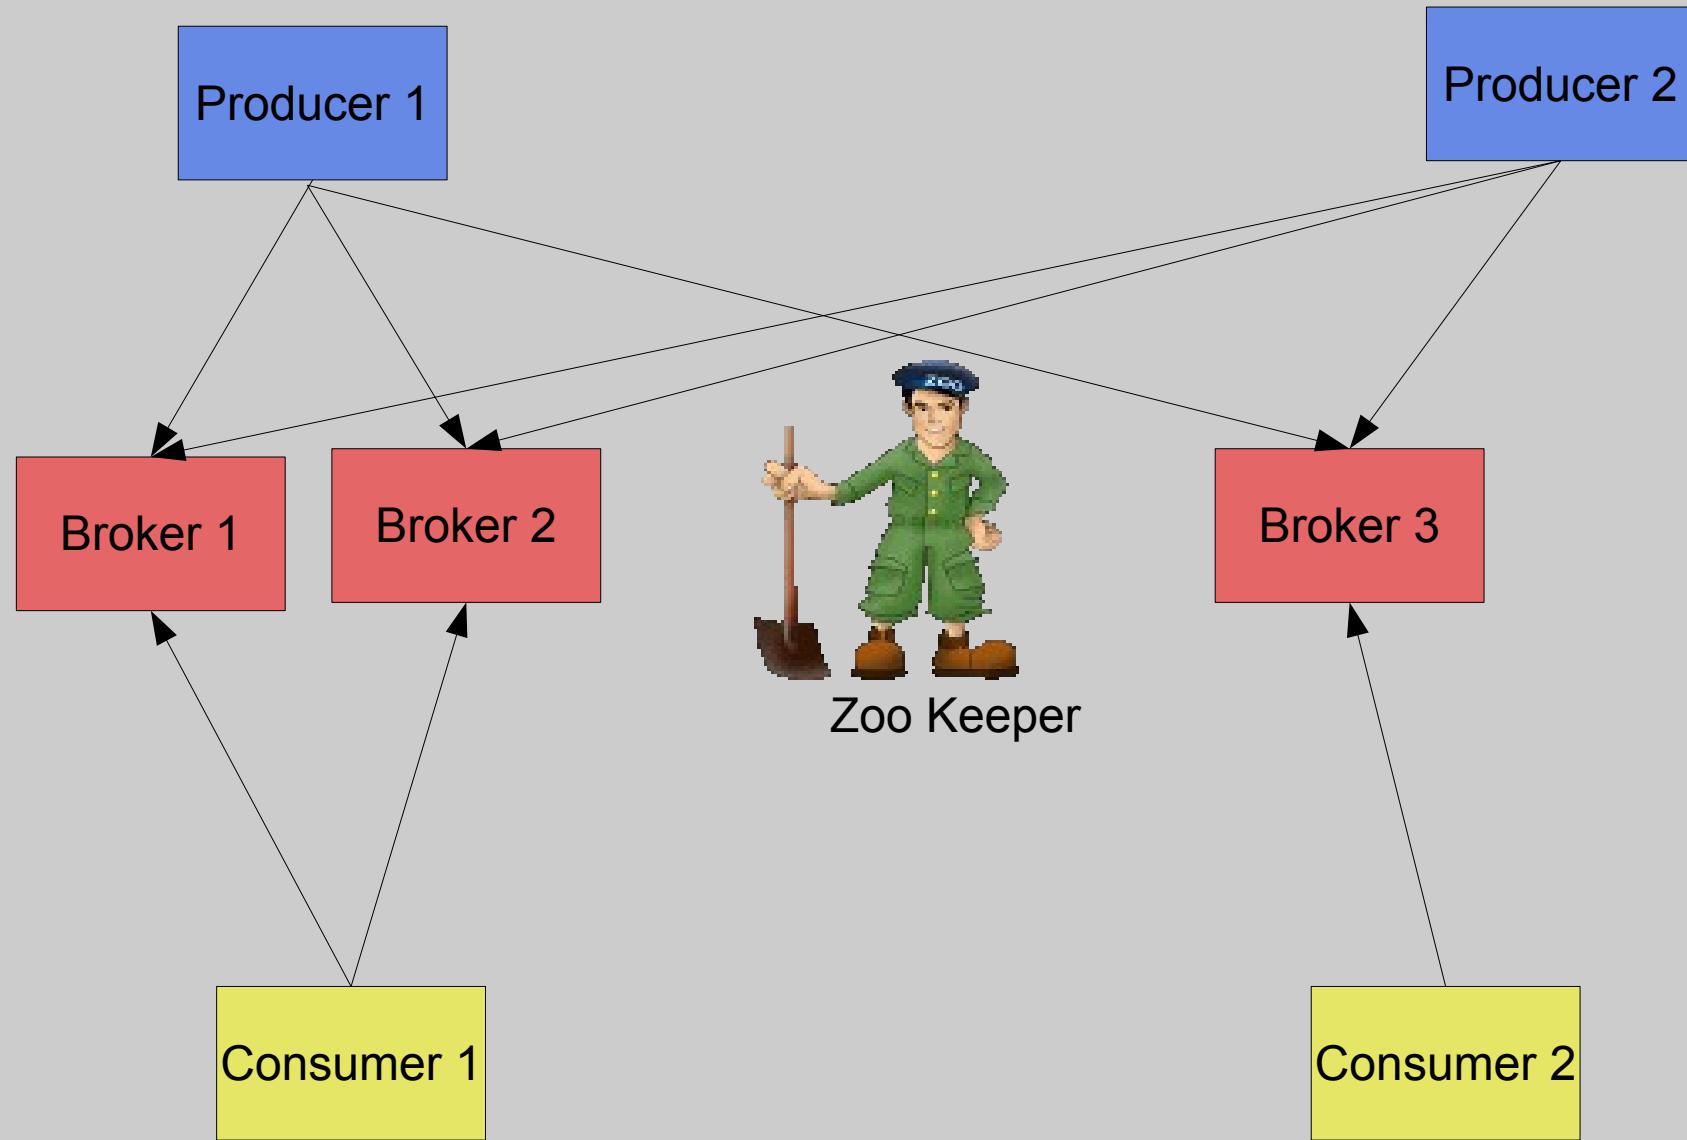


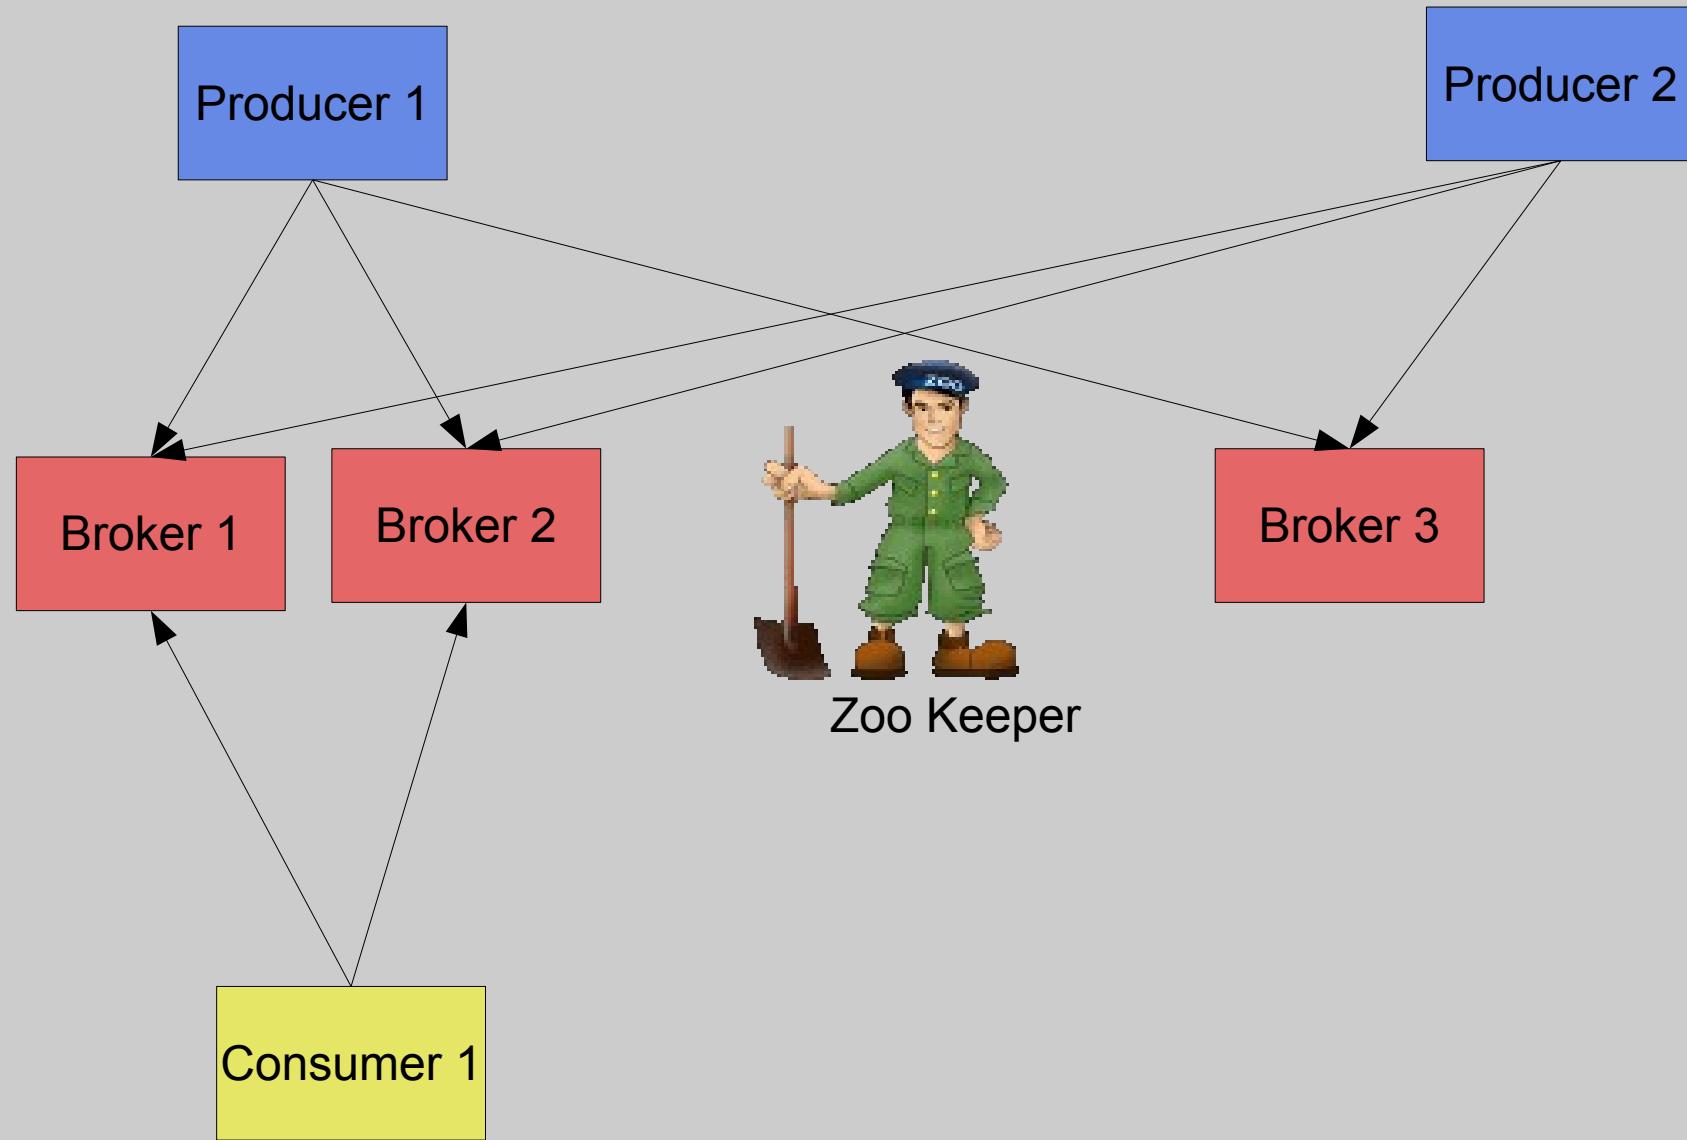


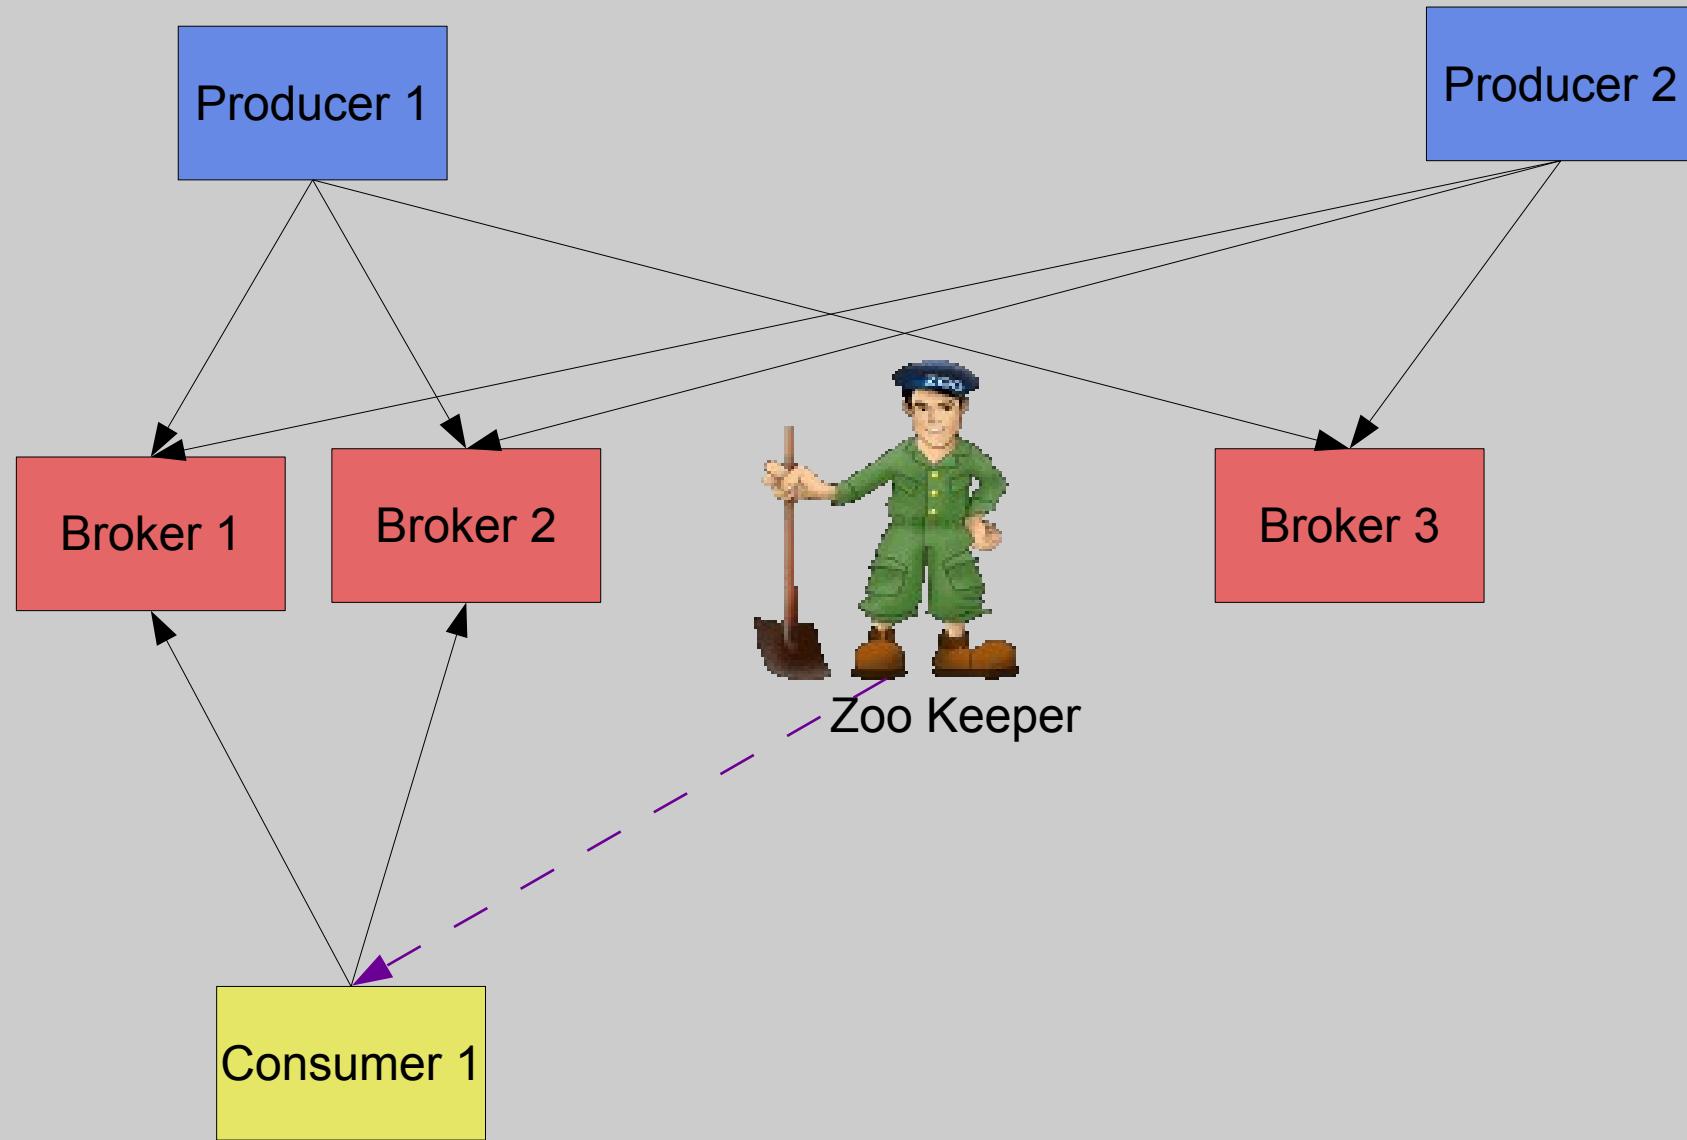


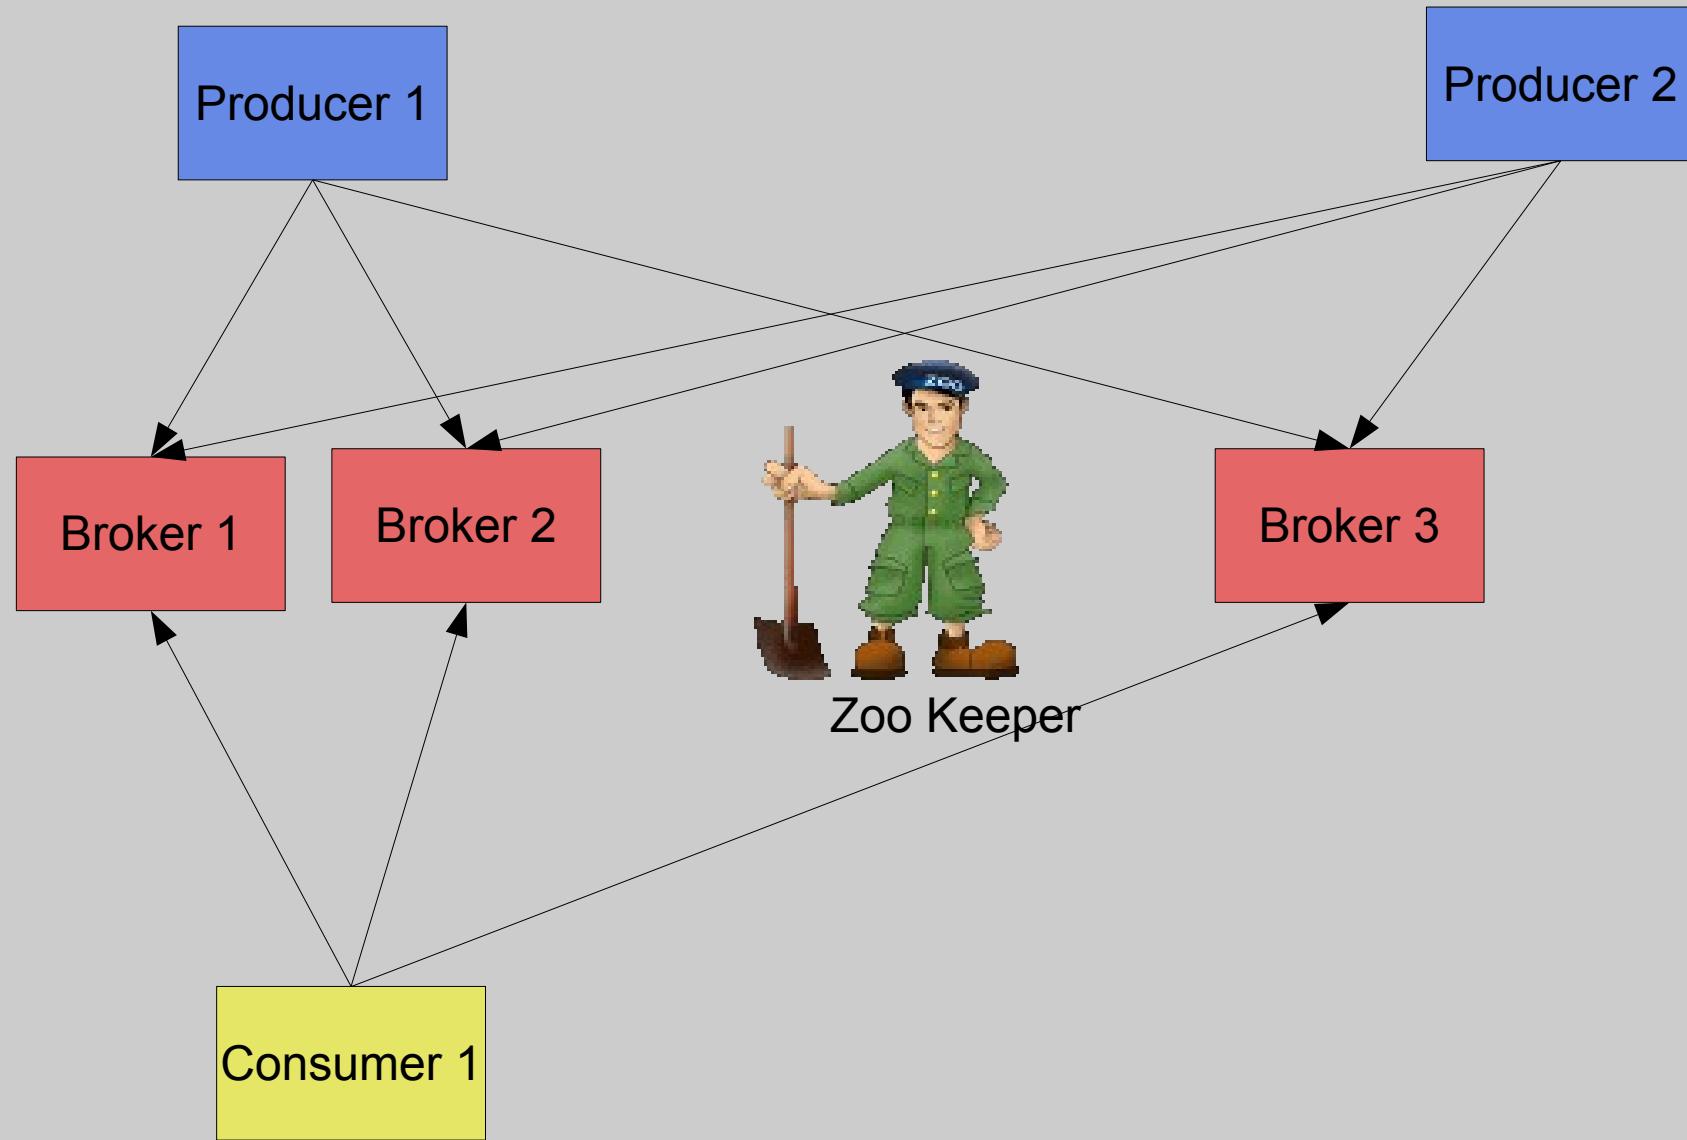










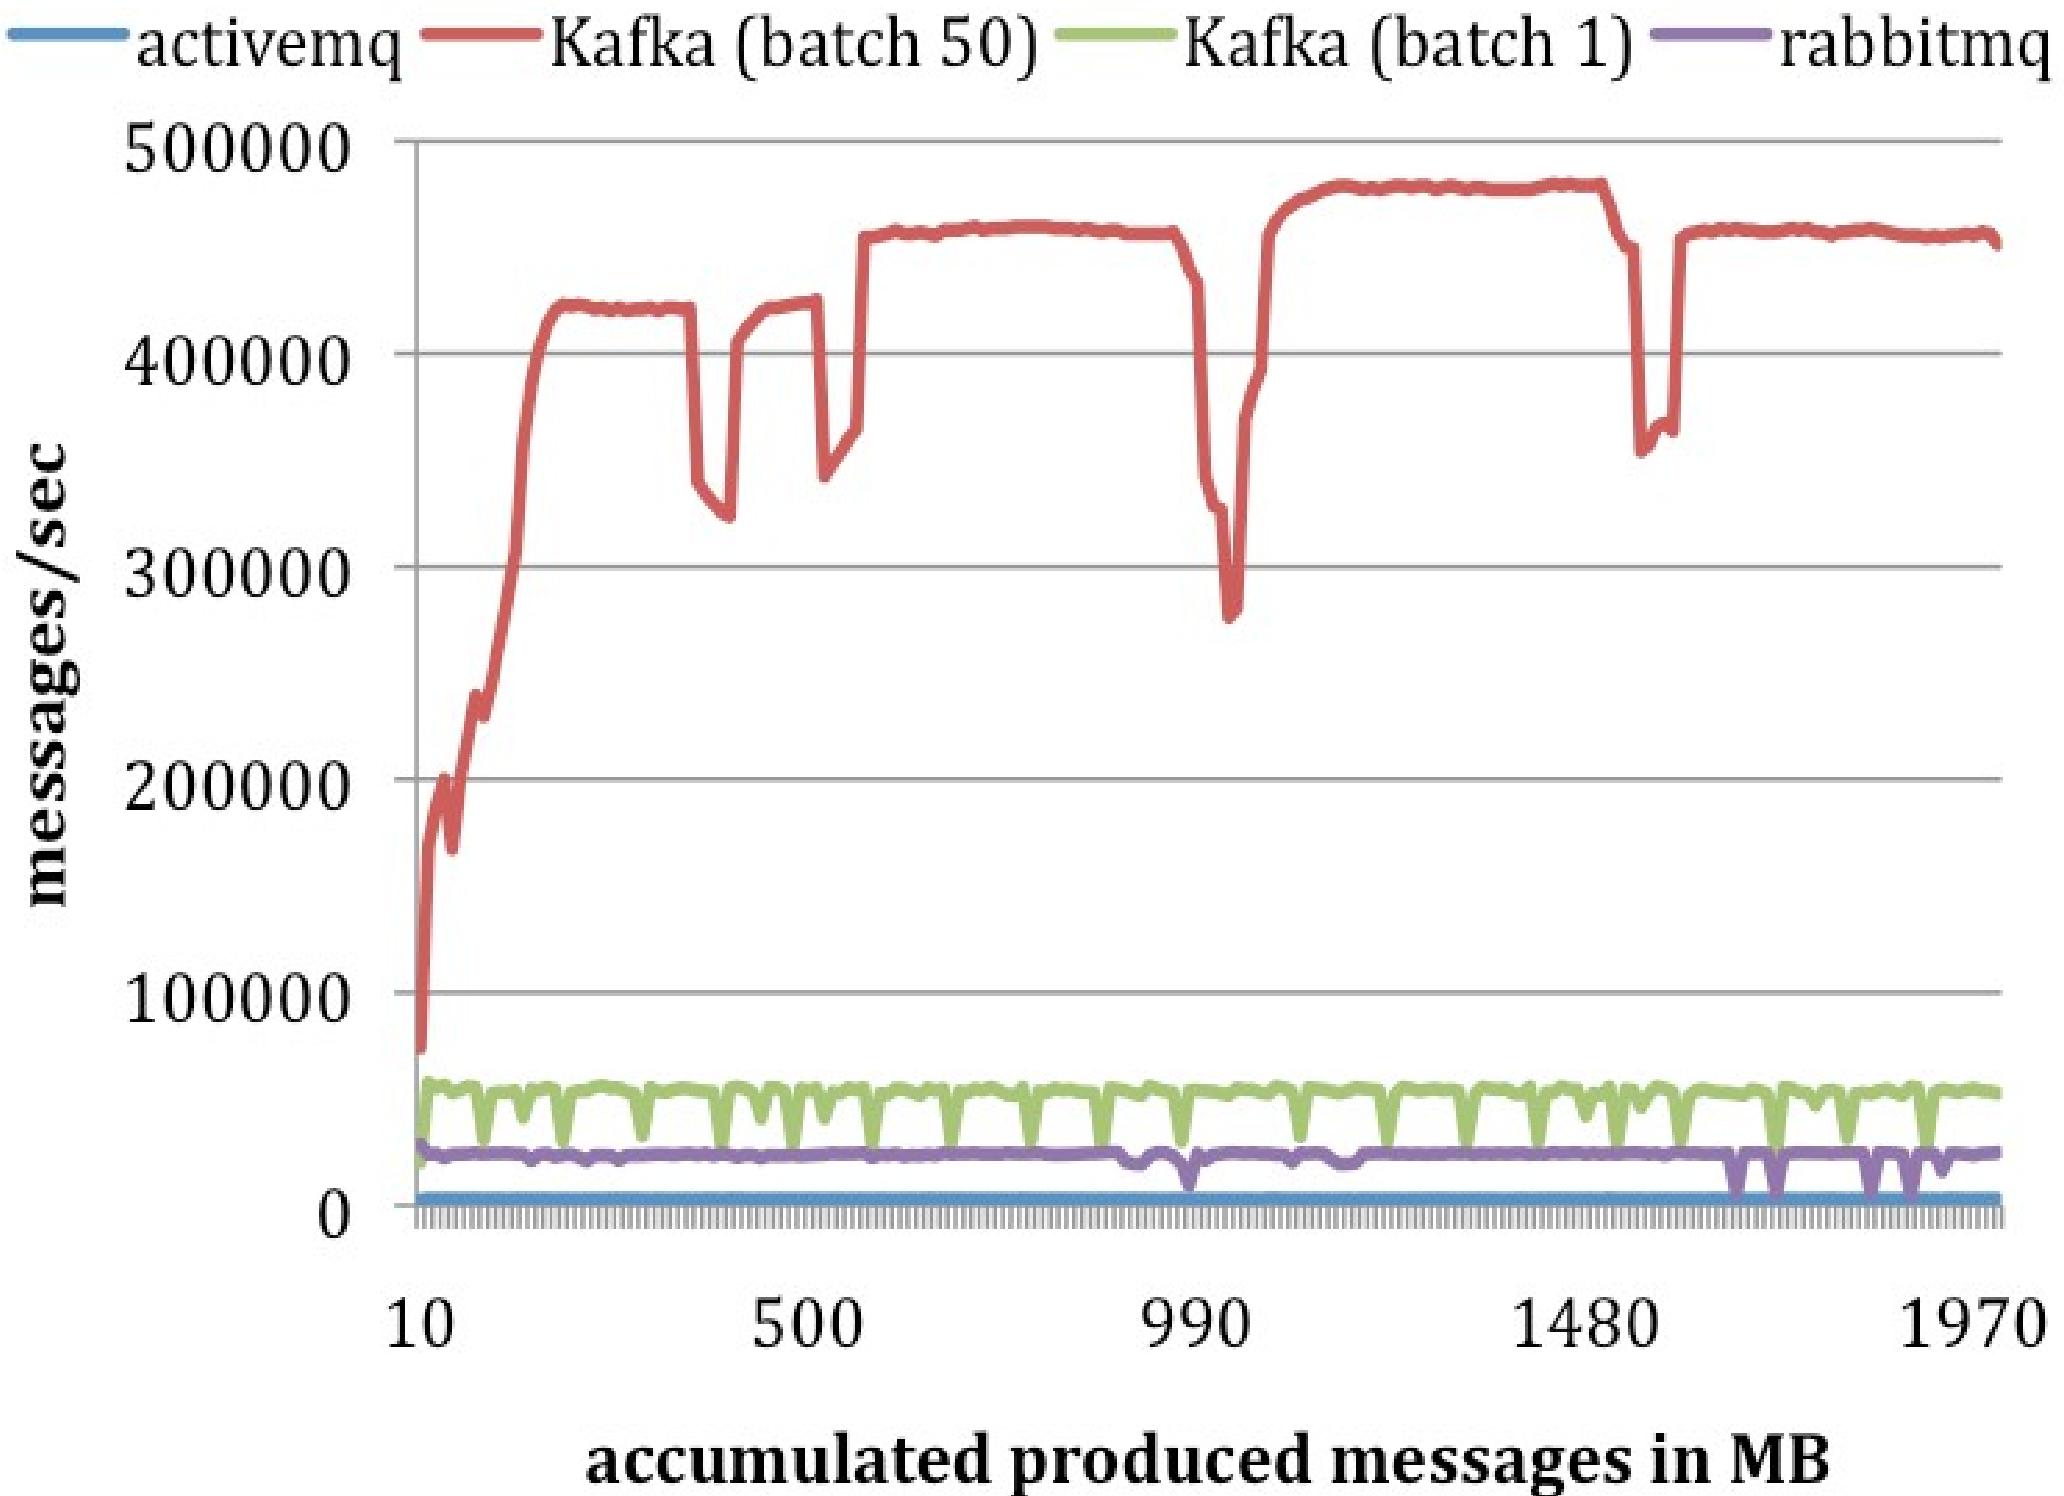


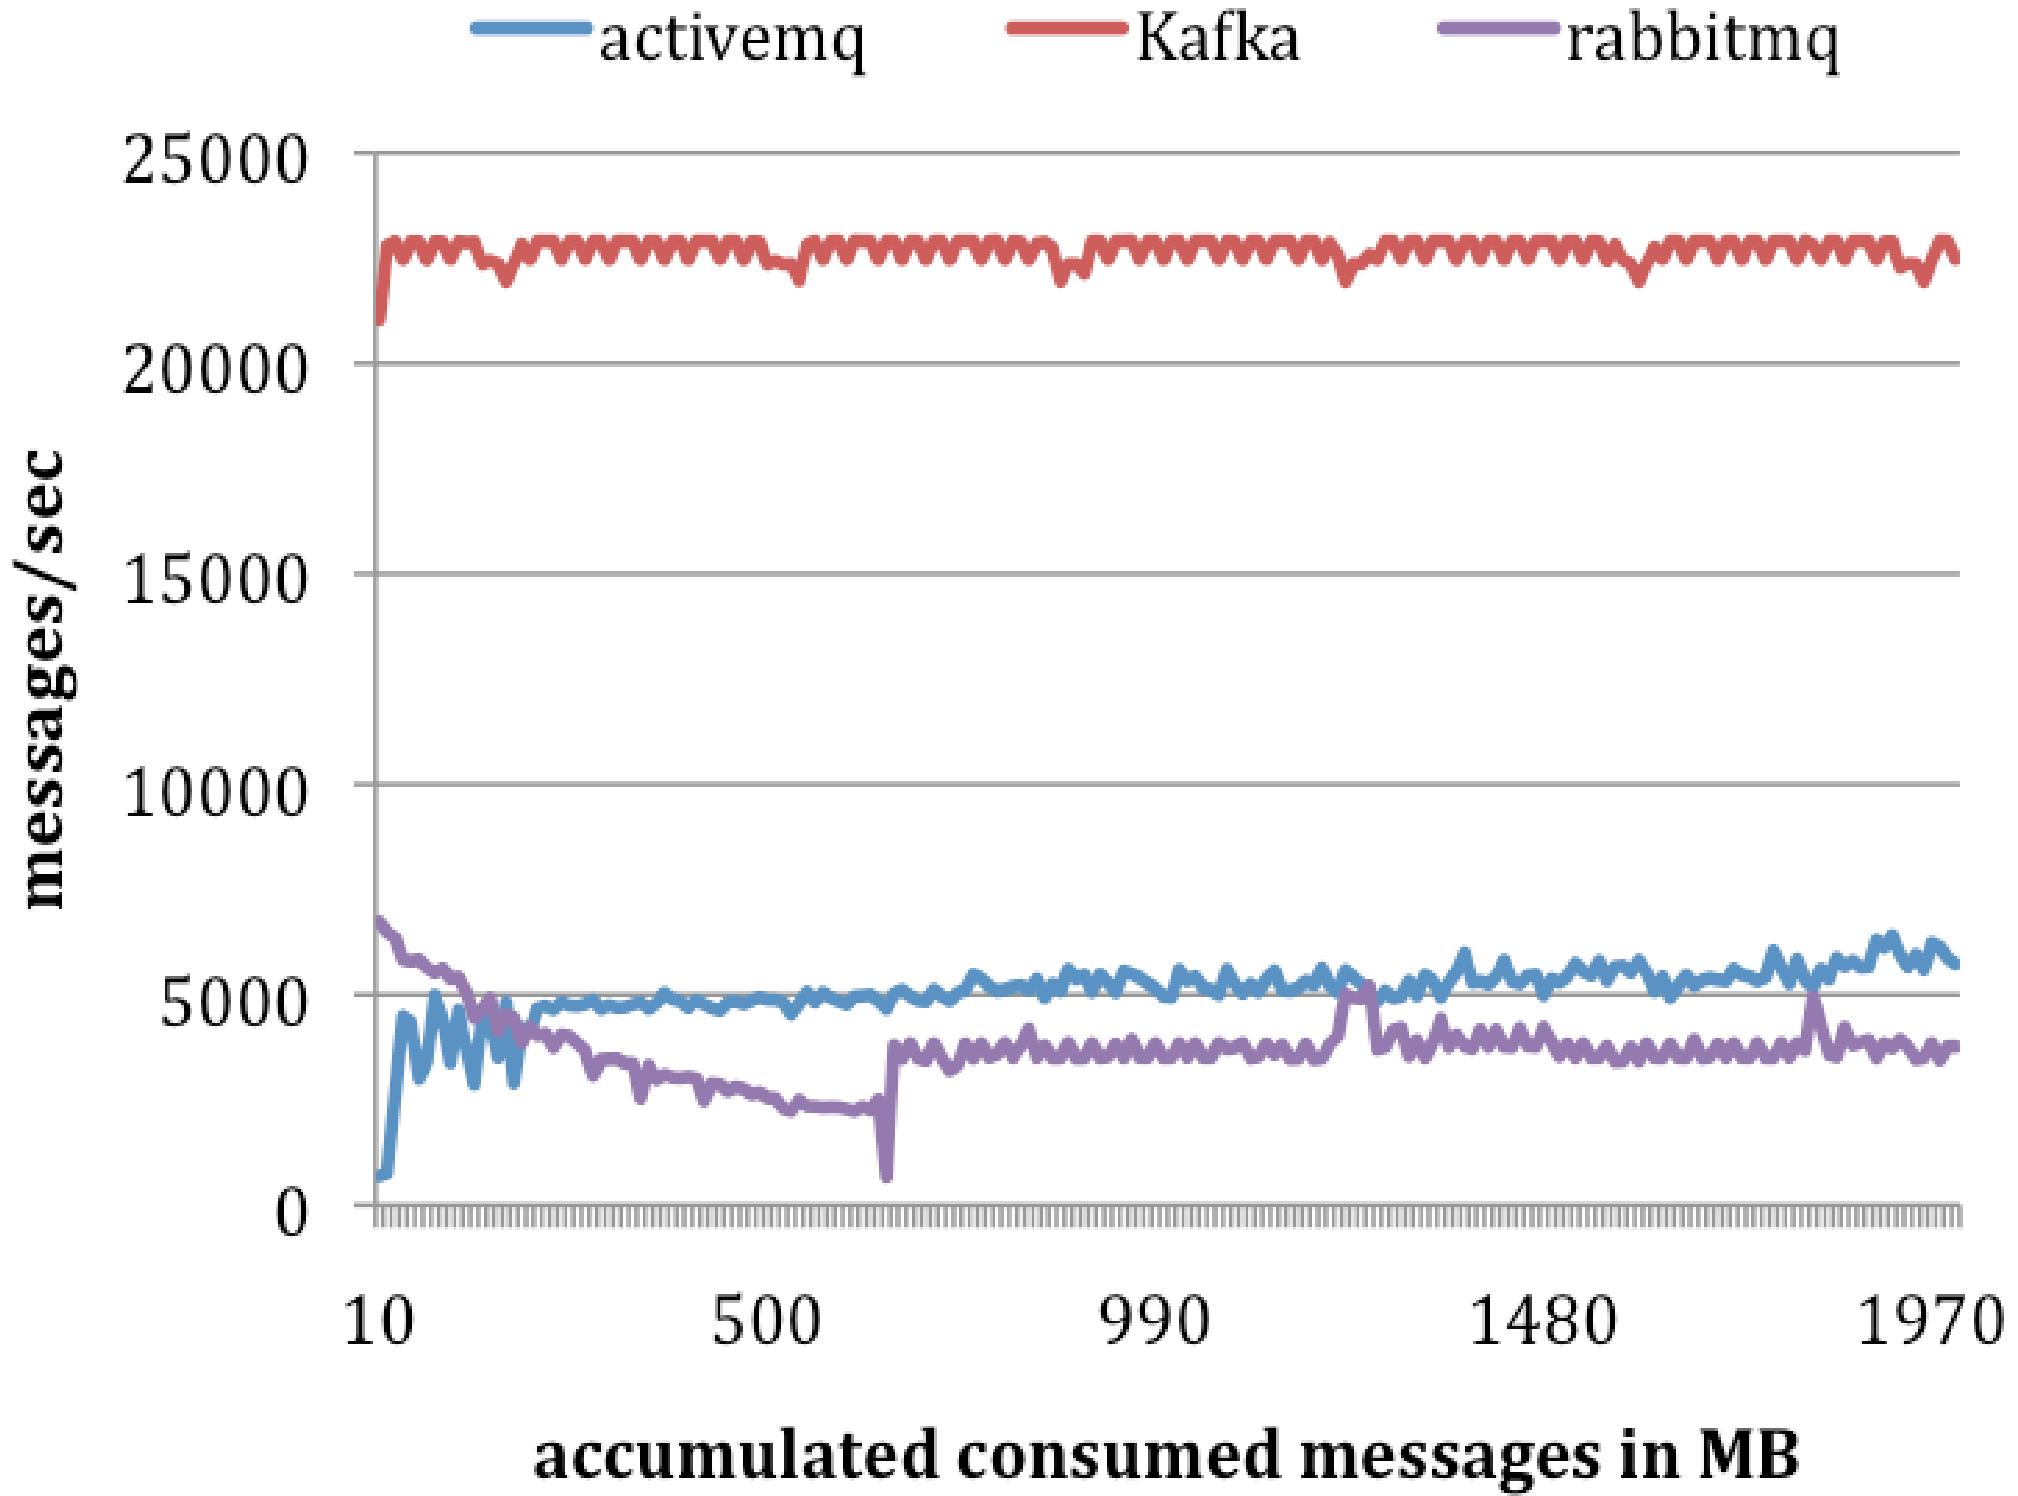
Performance Benchmark

1 Broker

1 Producer

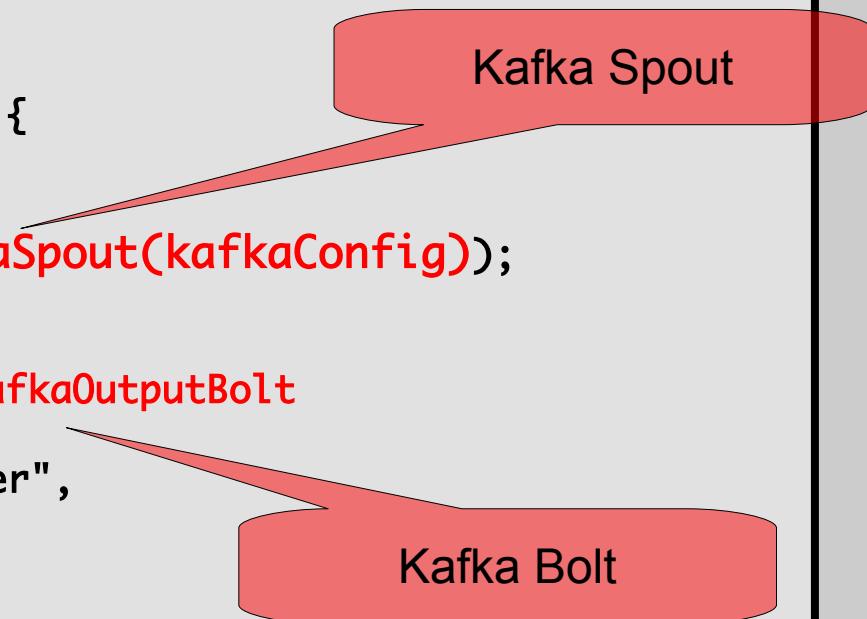
1 Consumer





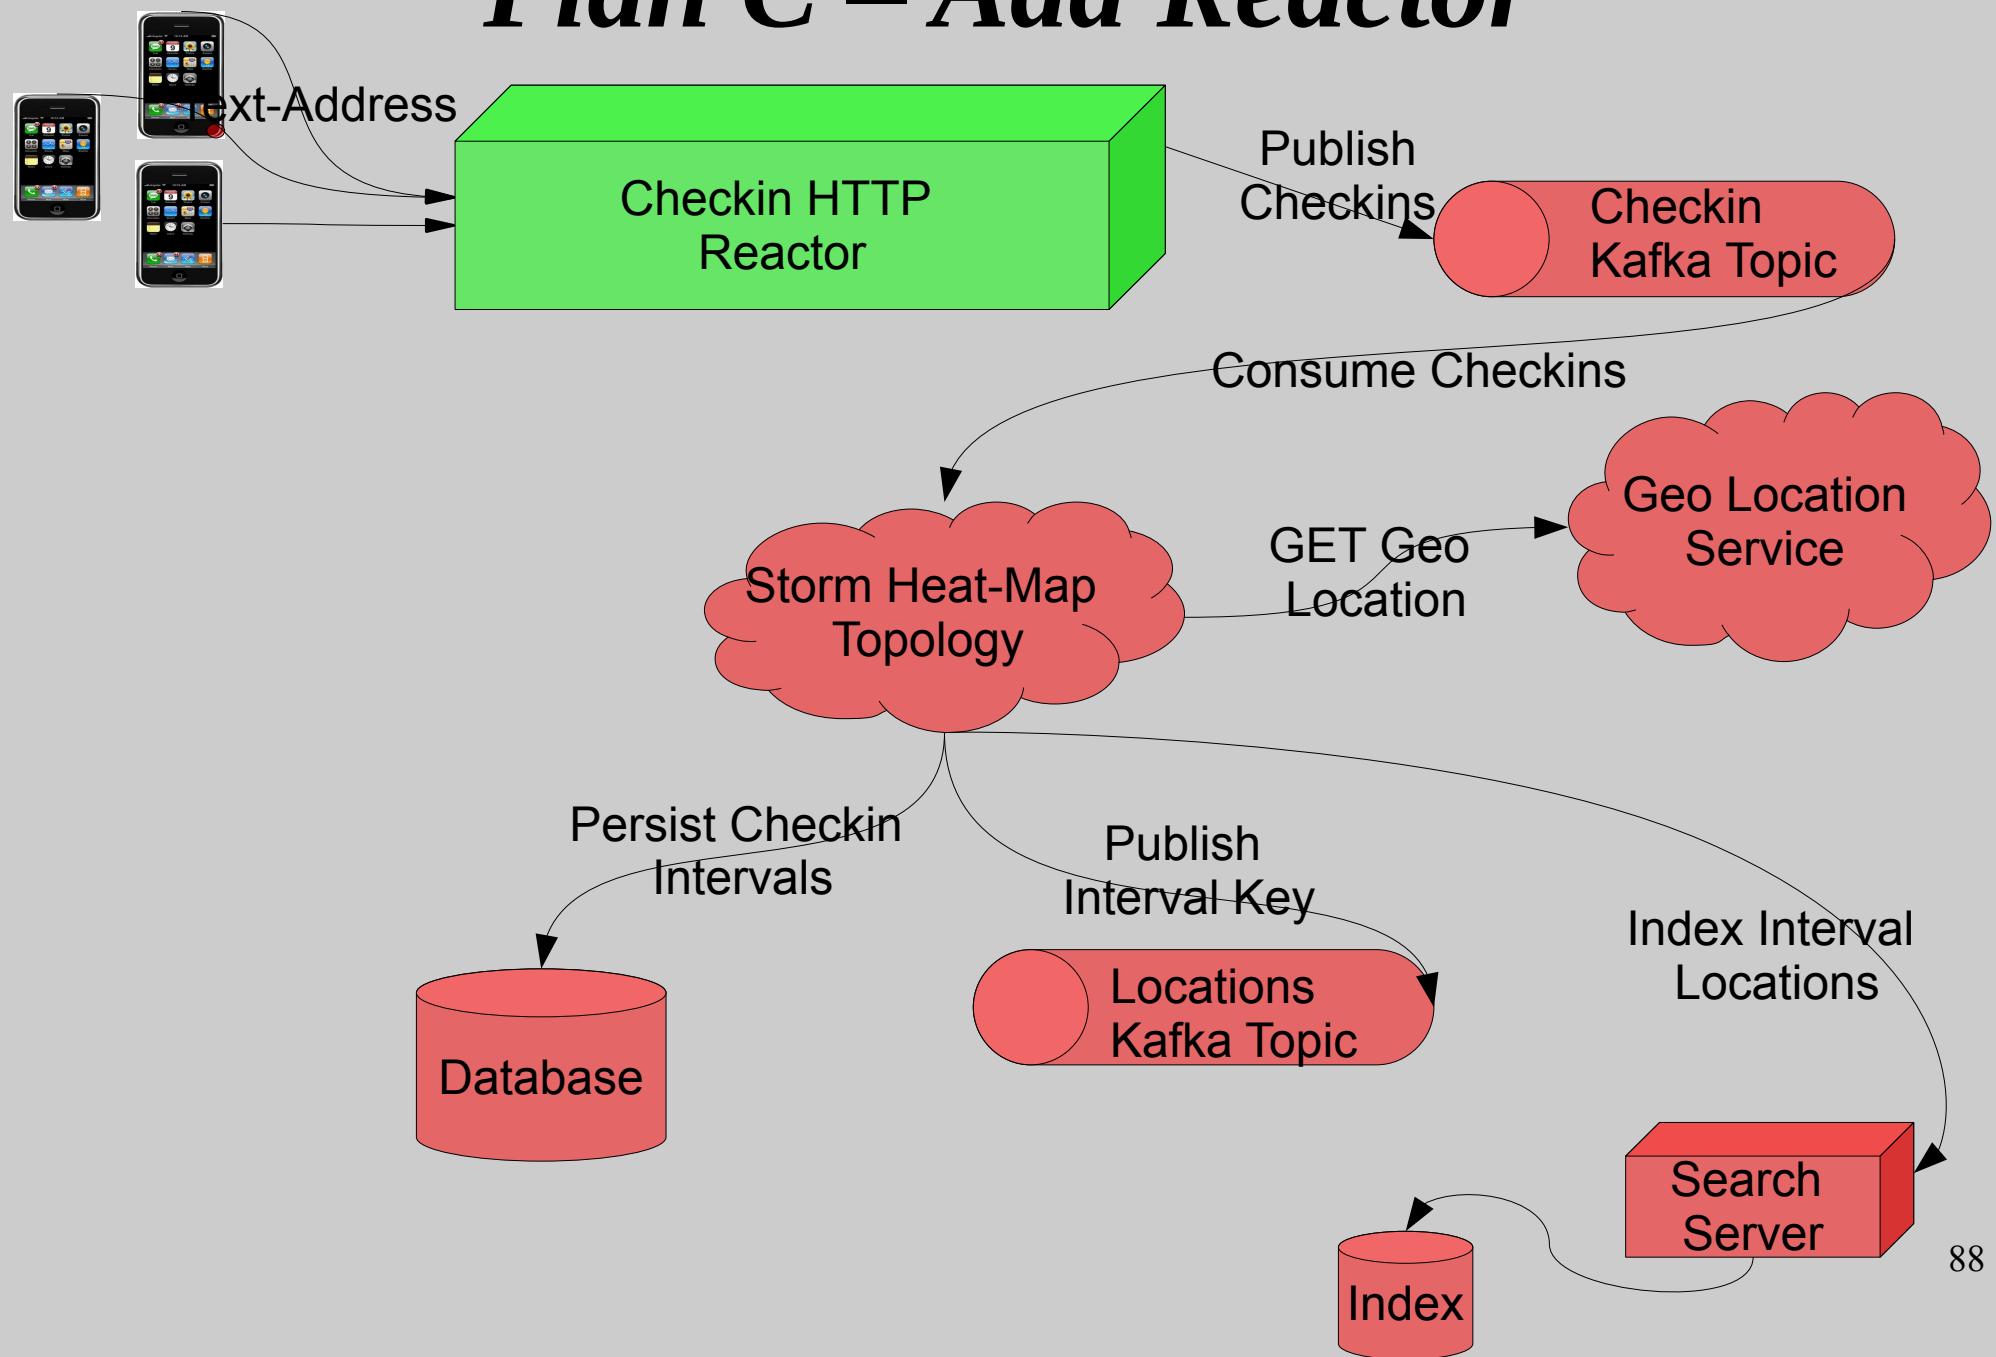
Add Kafka to our Topology

```
public class LocalTopologyRunner {  
    ...  
  
    private static TopologyBuilder buildTopology() {  
        ...  
        builder.setSpout("checkins", new KafkaSpout(kafkaConfig));  
  
        ...  
        builder.setBolt("kafkaProducer", new KafkaOutputBolt  
            ("localhost:9092",  
             "kafka.serializer.StringEncoder",  
             "locations-topic"))  
            .shuffleGrouping("persistor");  
  
        return builder;  
    }  
}
```

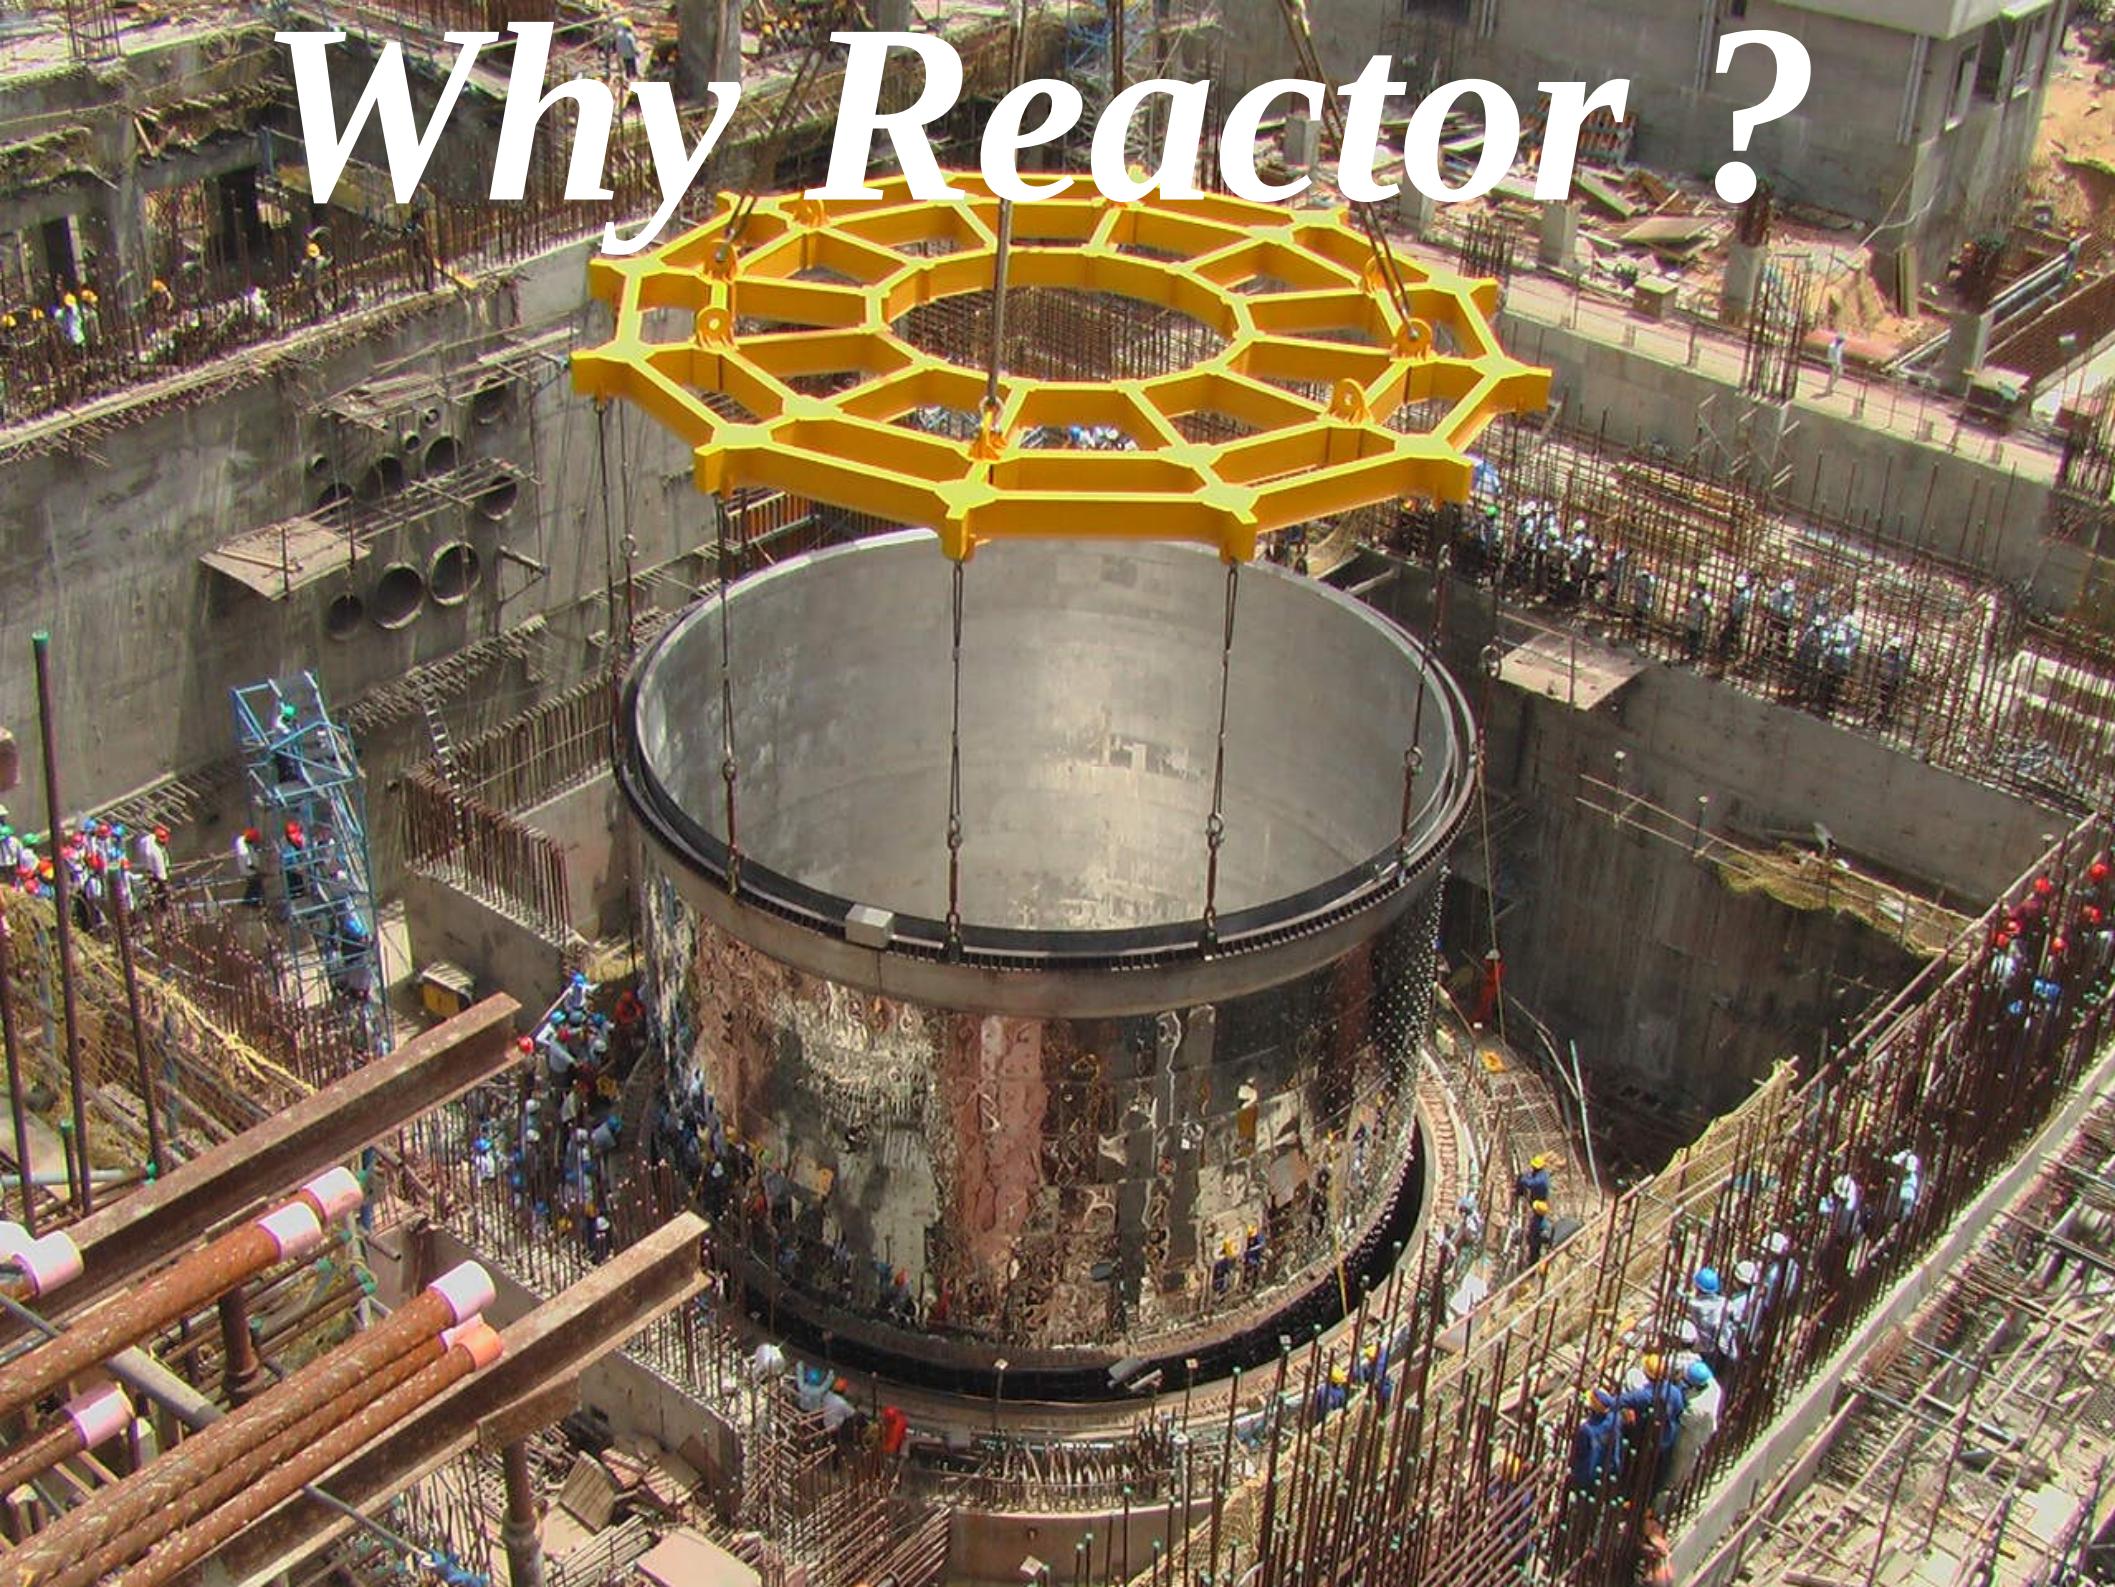


The code snippet shows the configuration of a topology. It includes a call to `builder.setSpout` which creates a `KafkaSpout` component. This component is highlighted with a red oval labeled "Kafka Spout". Below it, there is another red oval labeled "Kafka Bolt" with a line pointing to the `KafkaOutputBolt` component in the code.

Plan C – Add Reactor



Why Reactor ?

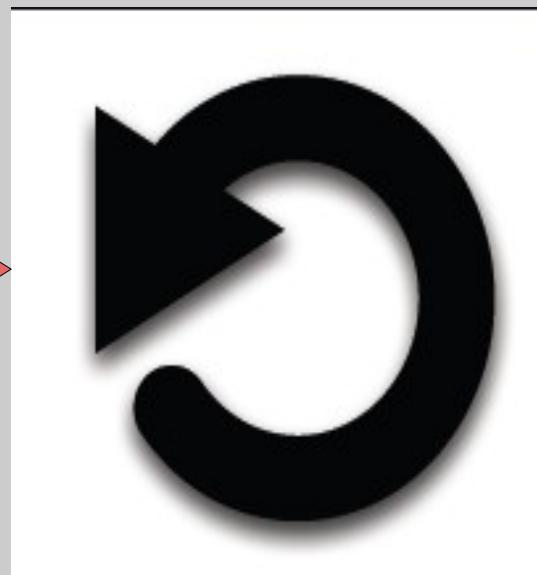


C10K Problem

2008:
Thread Per Request/Response

Reactor Pattern Paradigm

Application
registers
handlers



...events
trigger
handlers

Reactor Pattern – Key Points

- Single thread / single event loop
- EVERYTHING runs on it
- You MUST NOT block the event loop
- Many Implementations (partial list):
 - Node.js (JavaScript), EventMachine (Ruby), Twisted (Python)... and Vert.X

Reactor Pattern Problems

- Some work is naturally blocking:
 - Intensive data crunching
 - 3rd-party blocking API's (e.g. JDBC)
- Pure reactor (e.g. Node.js) is not a good fit for this kind of work!



Vert.x is a lightweight, high performance application platform for the JVM that's designed for modern mobile, web, and enterprise applications.

Polyglot

Write your application components in Java, JavaScript, CoffeeScript, Ruby, Python or Groovy...

... or mix and match several programming languages in a single app.

Simplicity

...without being simplistic.

Simple, powerful, APIs enable you to write non-blocking network enabled applications with ease. No complex configuration or boilerplate required.

Scalability

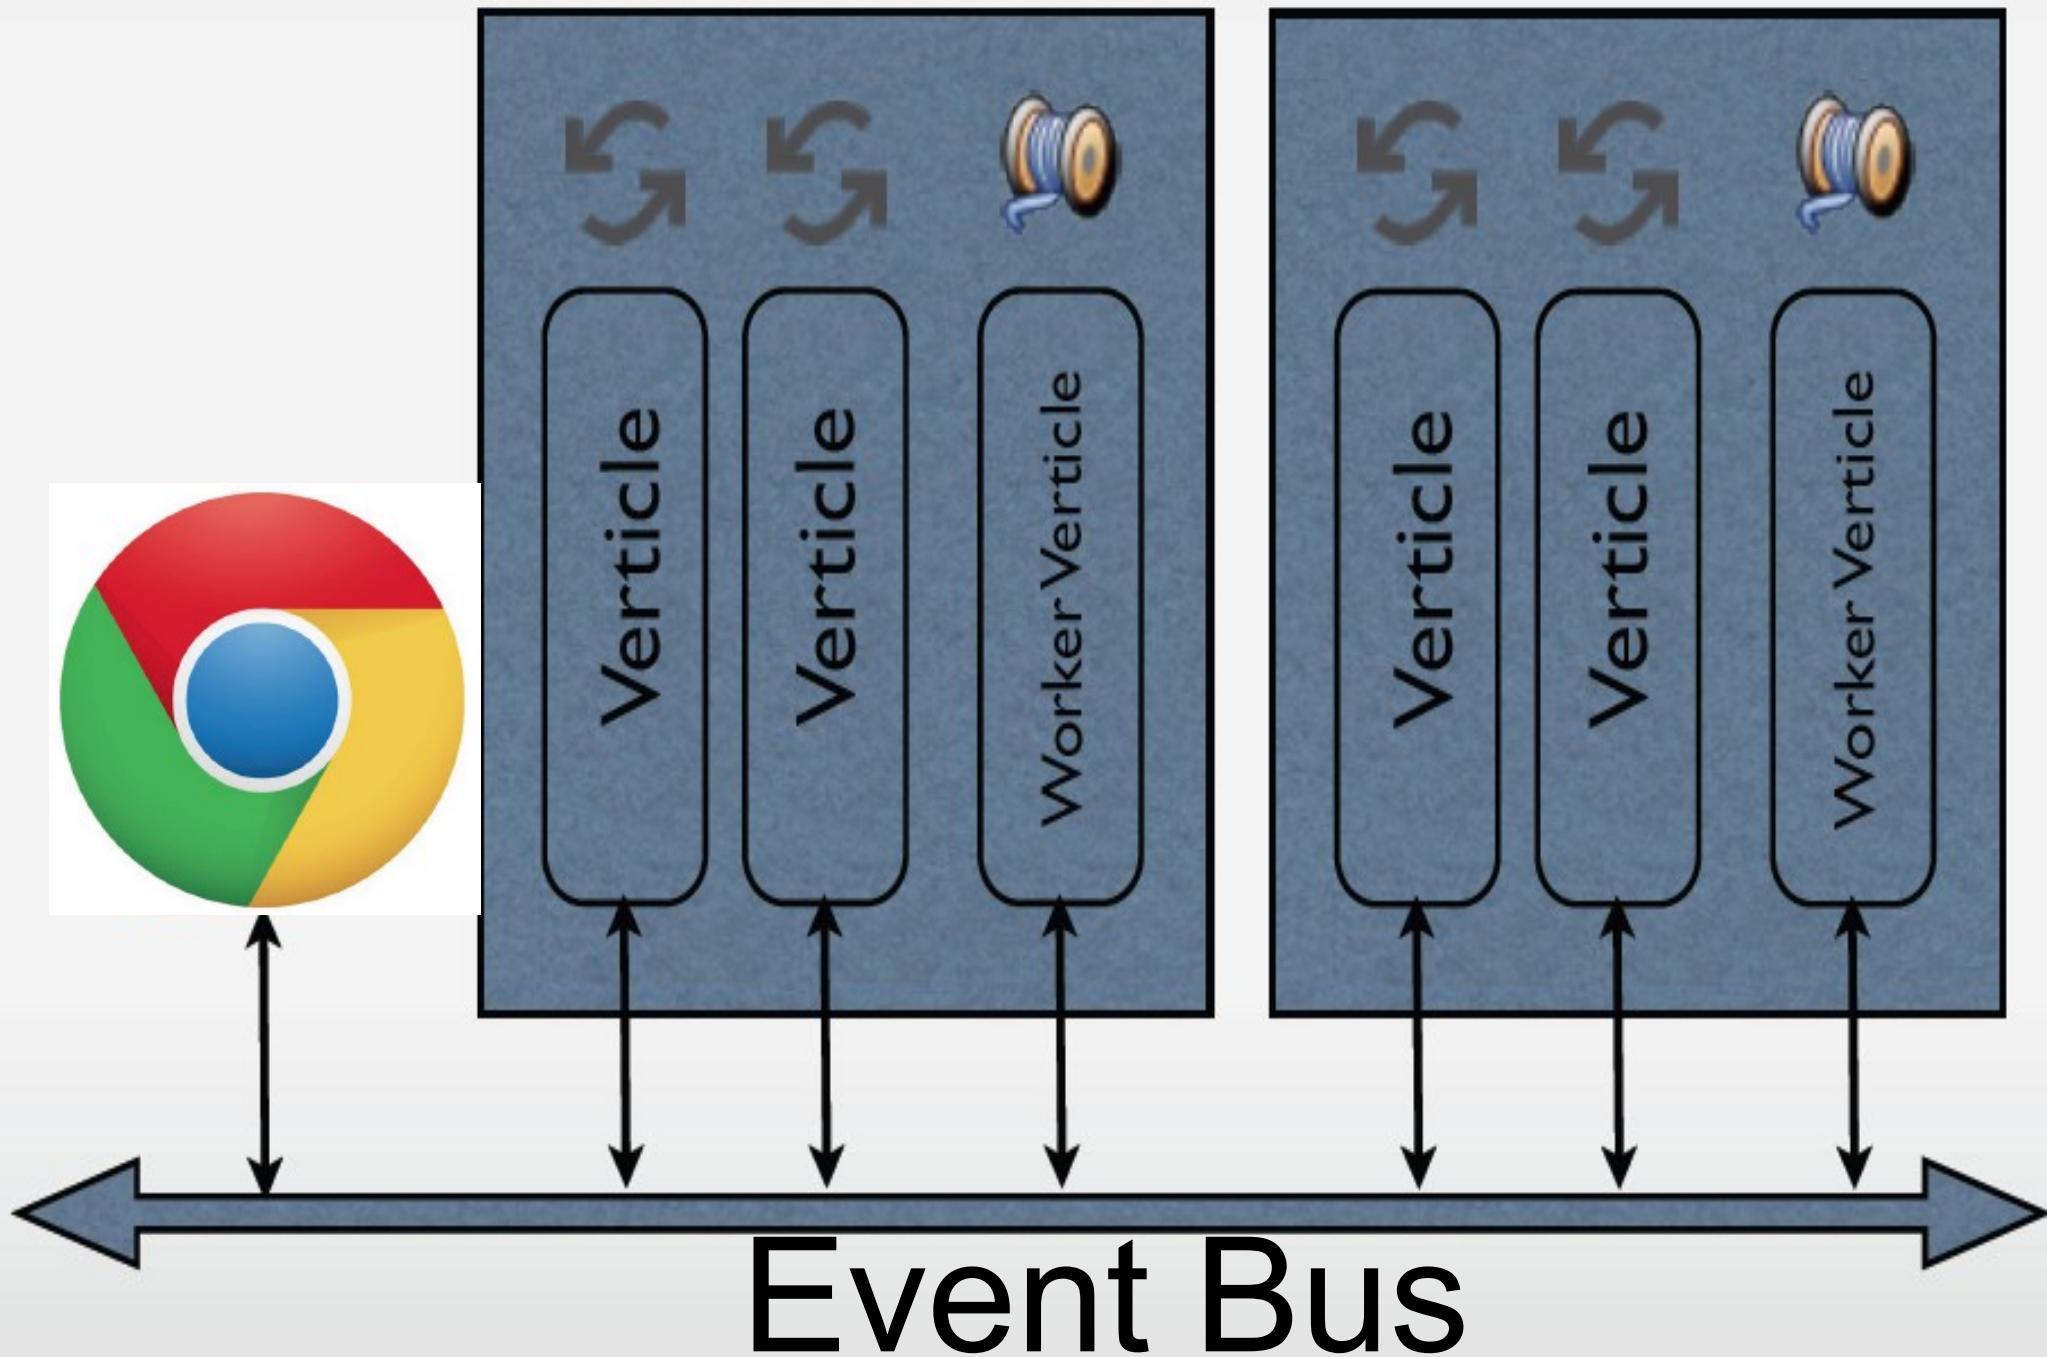
Scales using messaging passing to efficiently utilise your server cores.

Uses non blocking I/O to serve many connections with minimal threads.

Concurrency

Simple actor-like concurrency model frees you from the pain of traditional multi-threaded programming.

Vert.X Architecture



Vert.X Goodies

- Growing Module
- Repository
- web server
- Persistors (Mongo, JDBC, ...)
- Work queue
- Authentication
- Manager
- Session manager
- Socket.IO
- TCP/SSL servers/clients
- HTTP/HTTPS servers/clients
- WebSockets support
- SockJS support
- Timers
- Buffers
- Streams and Pumps
- Routing
- Asynchronous File I/O

Node.JS vs Vert.X



Node.js vs Vert.X

- Node.js
 - JavaScript Only
 - Inherently Single Threaded
 - No help much with IPC
 - All code **MUST** be in Event loop
- Vert.X
 - Polyglot (JavaScript, Java, Ruby, Python...)
 - Leverages JVM multi-threading
 - Nervous Event Bus
 - Blocking work can be done off the event loop

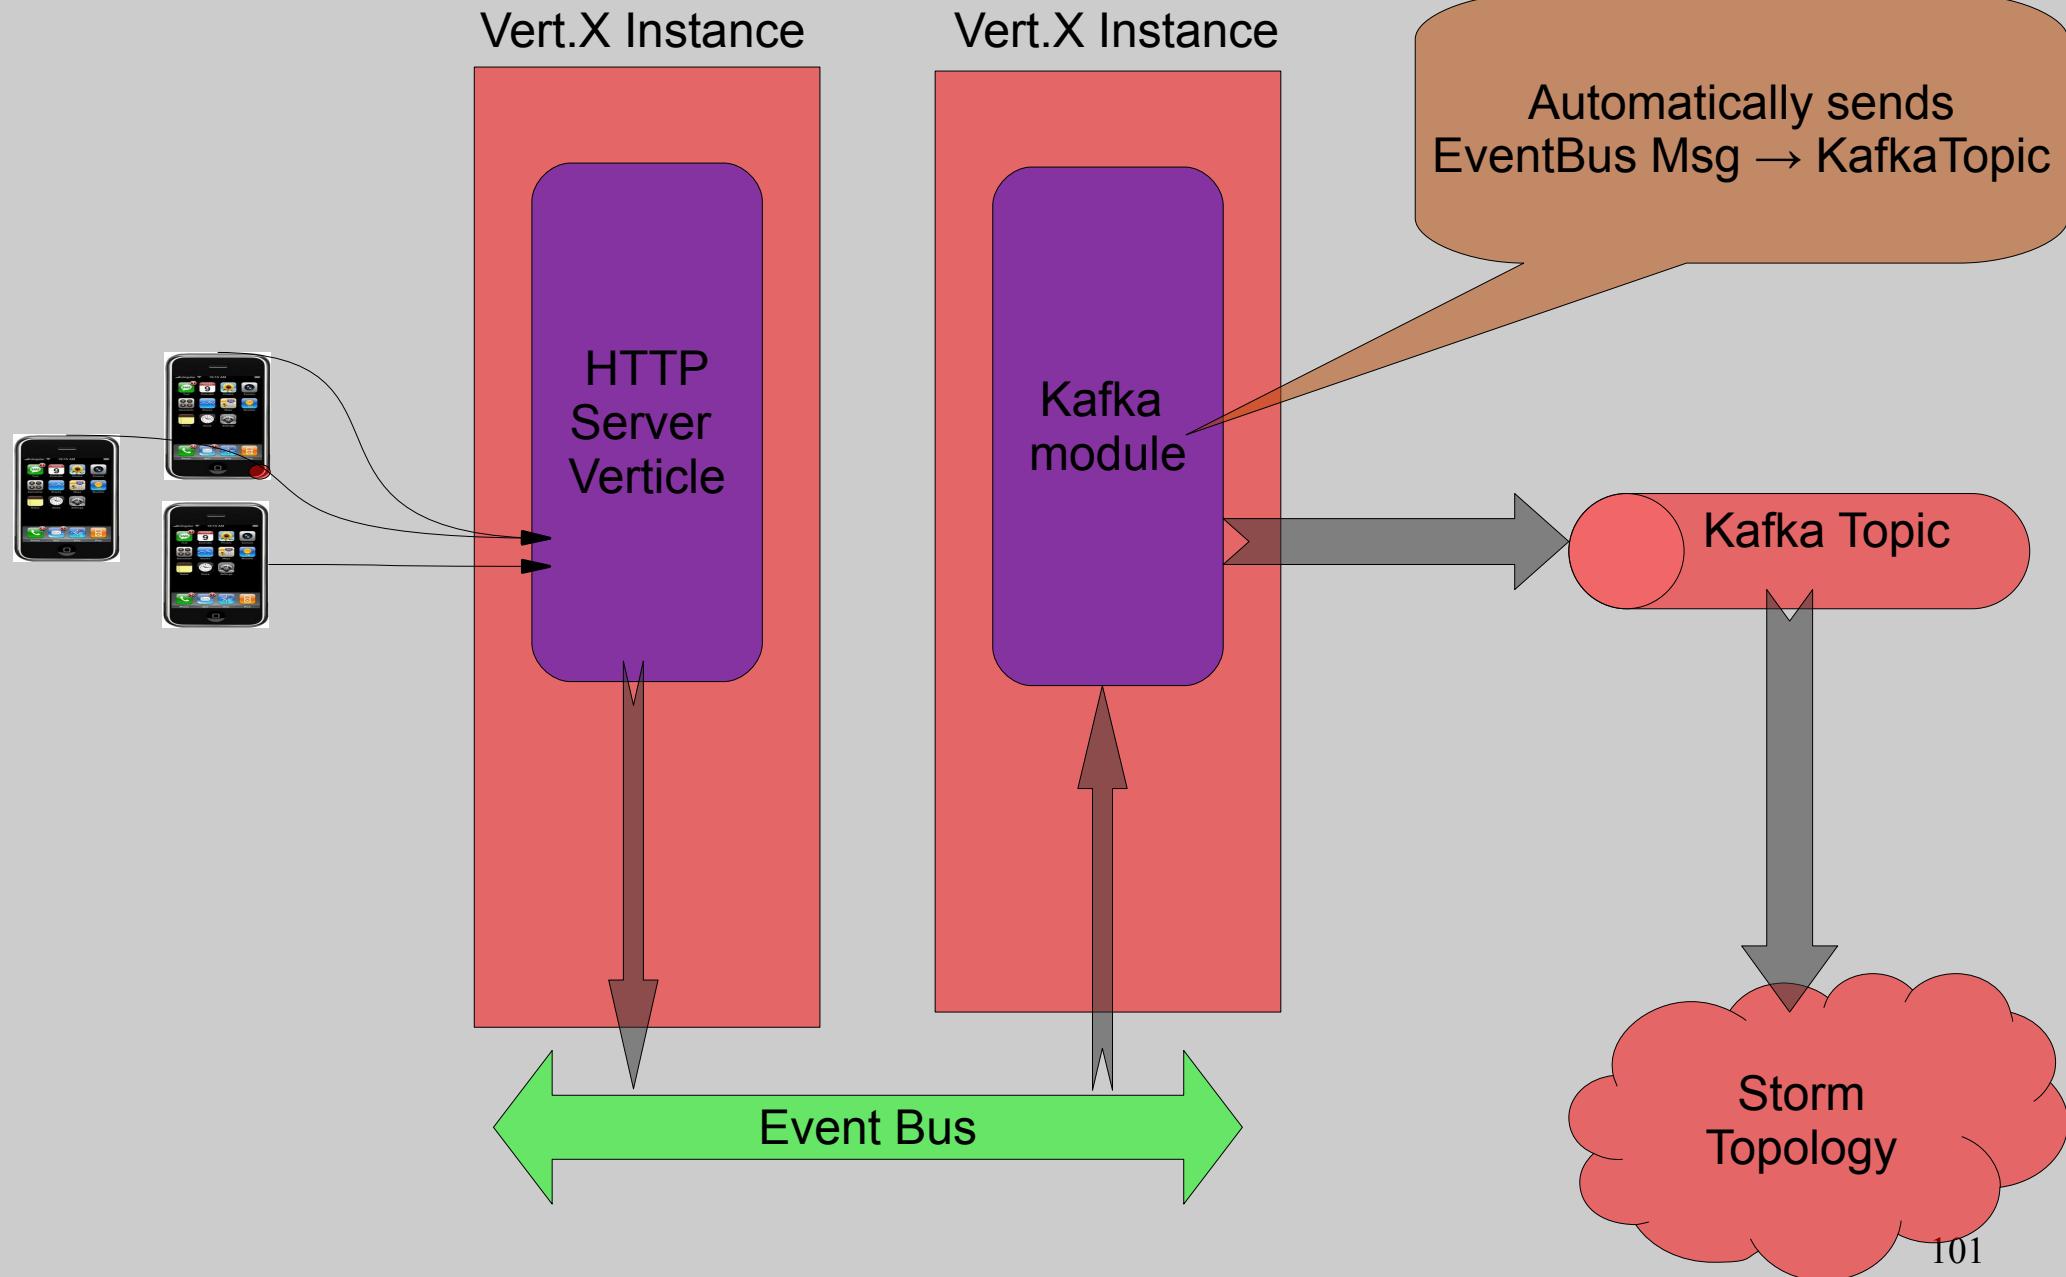
Node.js vs Vert.X Benchmark



<http://vertxproject.wordpress.com/2012/05/09/vert-x-vs-node-js-simple-http-benchmarks/>

*AMD Phenom II X6 (6 core), 8GB
RAM, Ubuntu 11.04*

HeatMap Reactor Architecture



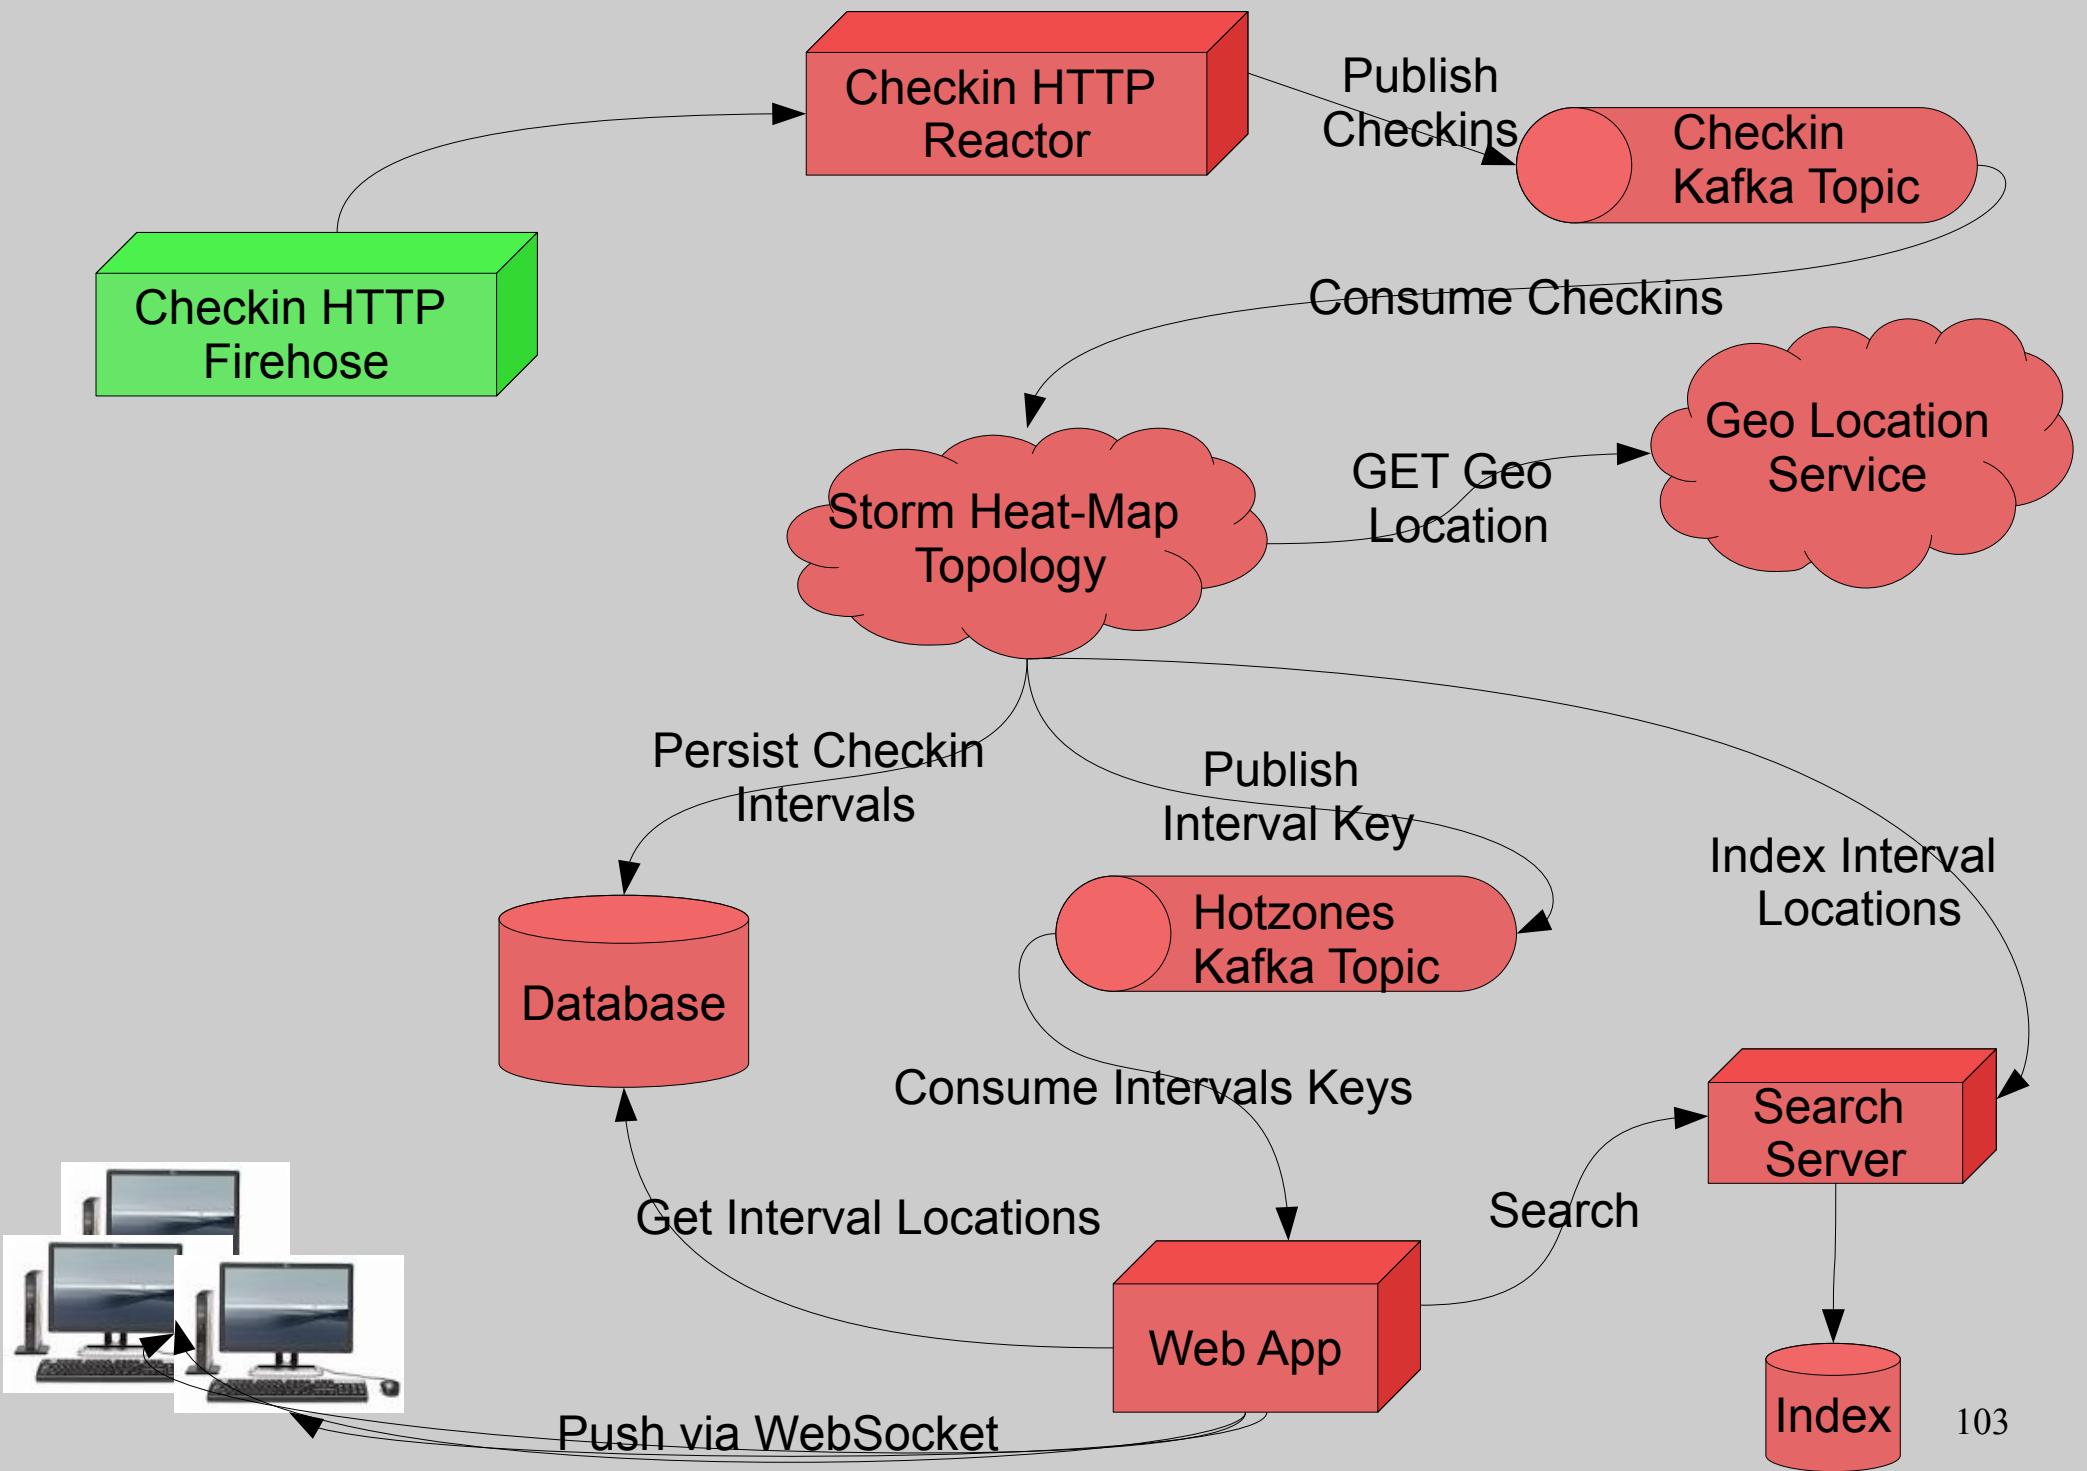
Heat-Map Server – Only 6 LOC !

```
var vertx = require('vertx');
var container = require('vertx/container');
var console = require('vertx/console');
var config = container.config;

vertx.createHttpServer().requestHandler(function(request) {
    request.dataHandler(function(buffer) {
        vertx.eventBus.send(config.address, {payload:buffer.toString()});
    });
    request.response.end();
}).listen(config.httpServerPort, config.httpServerHost);

console.log("HTTP CheckinsReactor started on port "+config.httpServerPort);
```

Send checkin
to Vert.X EventBus



Demo



Summary

When You go out to Salsa Club



- Good Music
- Crowded

More Conclusions..

- Storm – Great for real-time BigData processing.
Complementary for Hadoop batch jobs.
- Kafka – Great messaging for logs/events data, been served as a good “source” for Storm spout
- Vert.X – Worth trial and check as an alternative for reactor.

Thanks

